



EKS Networking 101

AWS Professional Services
2023.11



Agenda

- Networking foundation
- EKS Control Plane Network
- EKS Data Plane Network (Amazon VPC CNI)
- Ingress Controller

Topics not covered in this session

- Kubernetes Basics
- Linux IPtables
- Mesh Network (App Mesh, Istio)
- Amazon VPC Lattice

Networking foundation



Linux Networking Primitives for Kubernetes



Node

The basis of Kubernetes
Workload



Cgroups

Linux kernel feature



Namespace

Process isolation



Network Namespace

Creates a copy of the
network stack



Veth

Virtual Ethernet



Bridge Interface

Provide connectivity



Pause Container

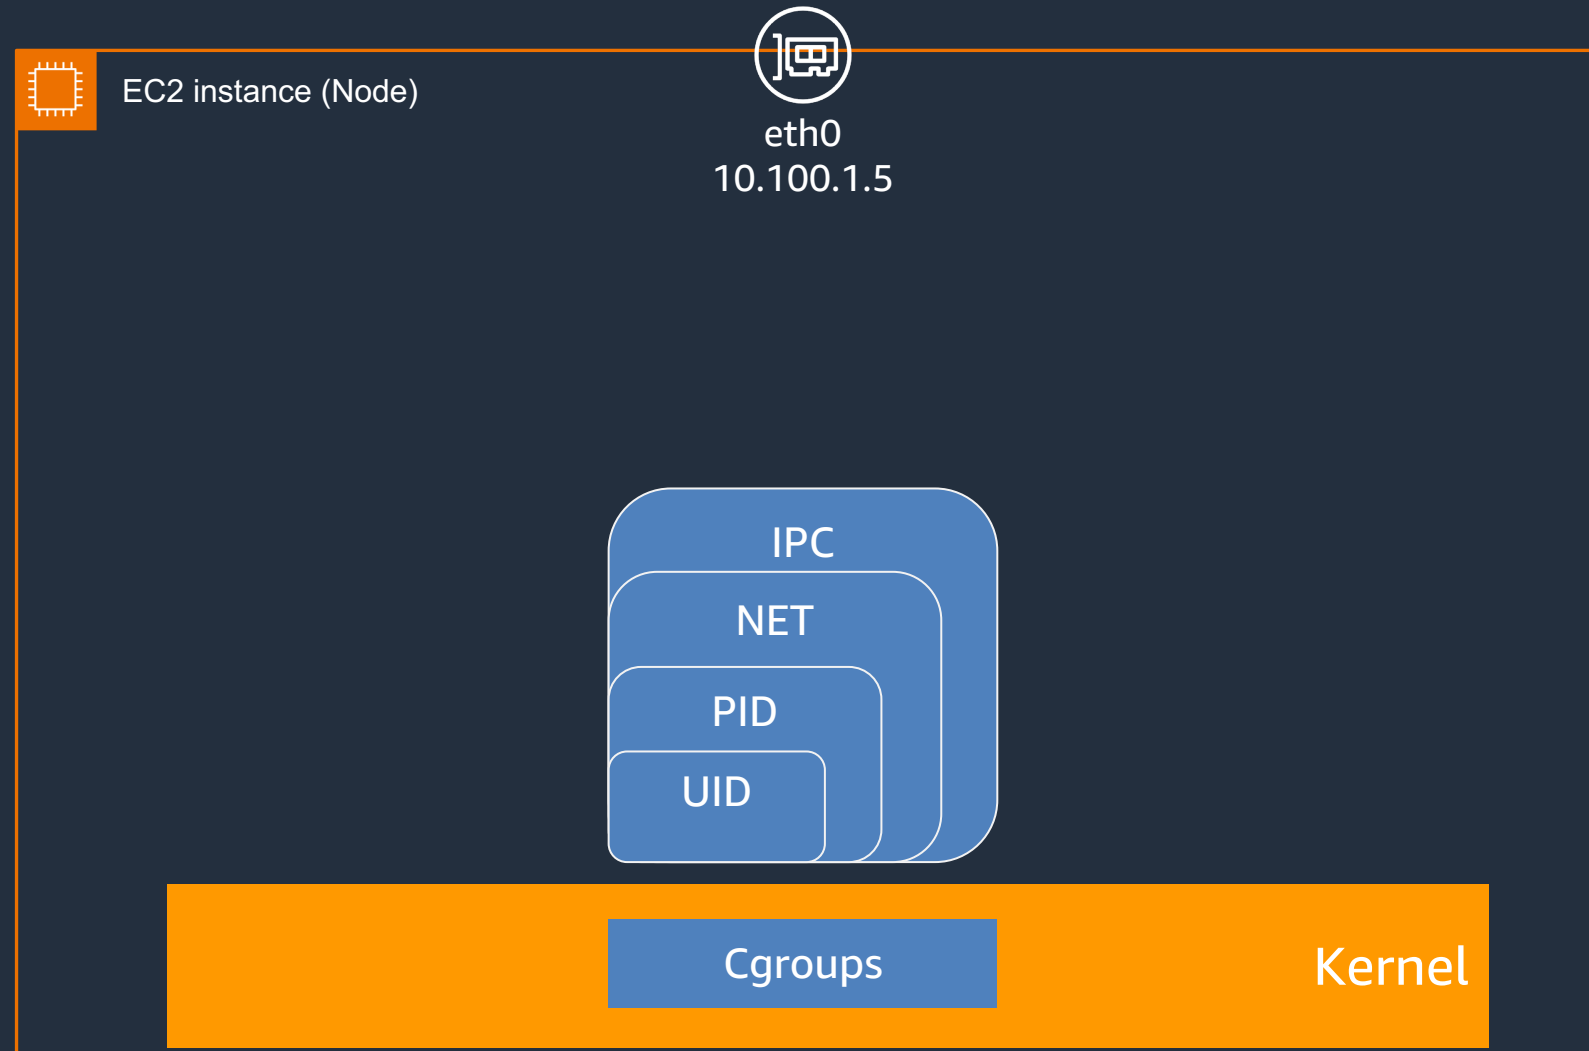
Manages the namespace
for each pod



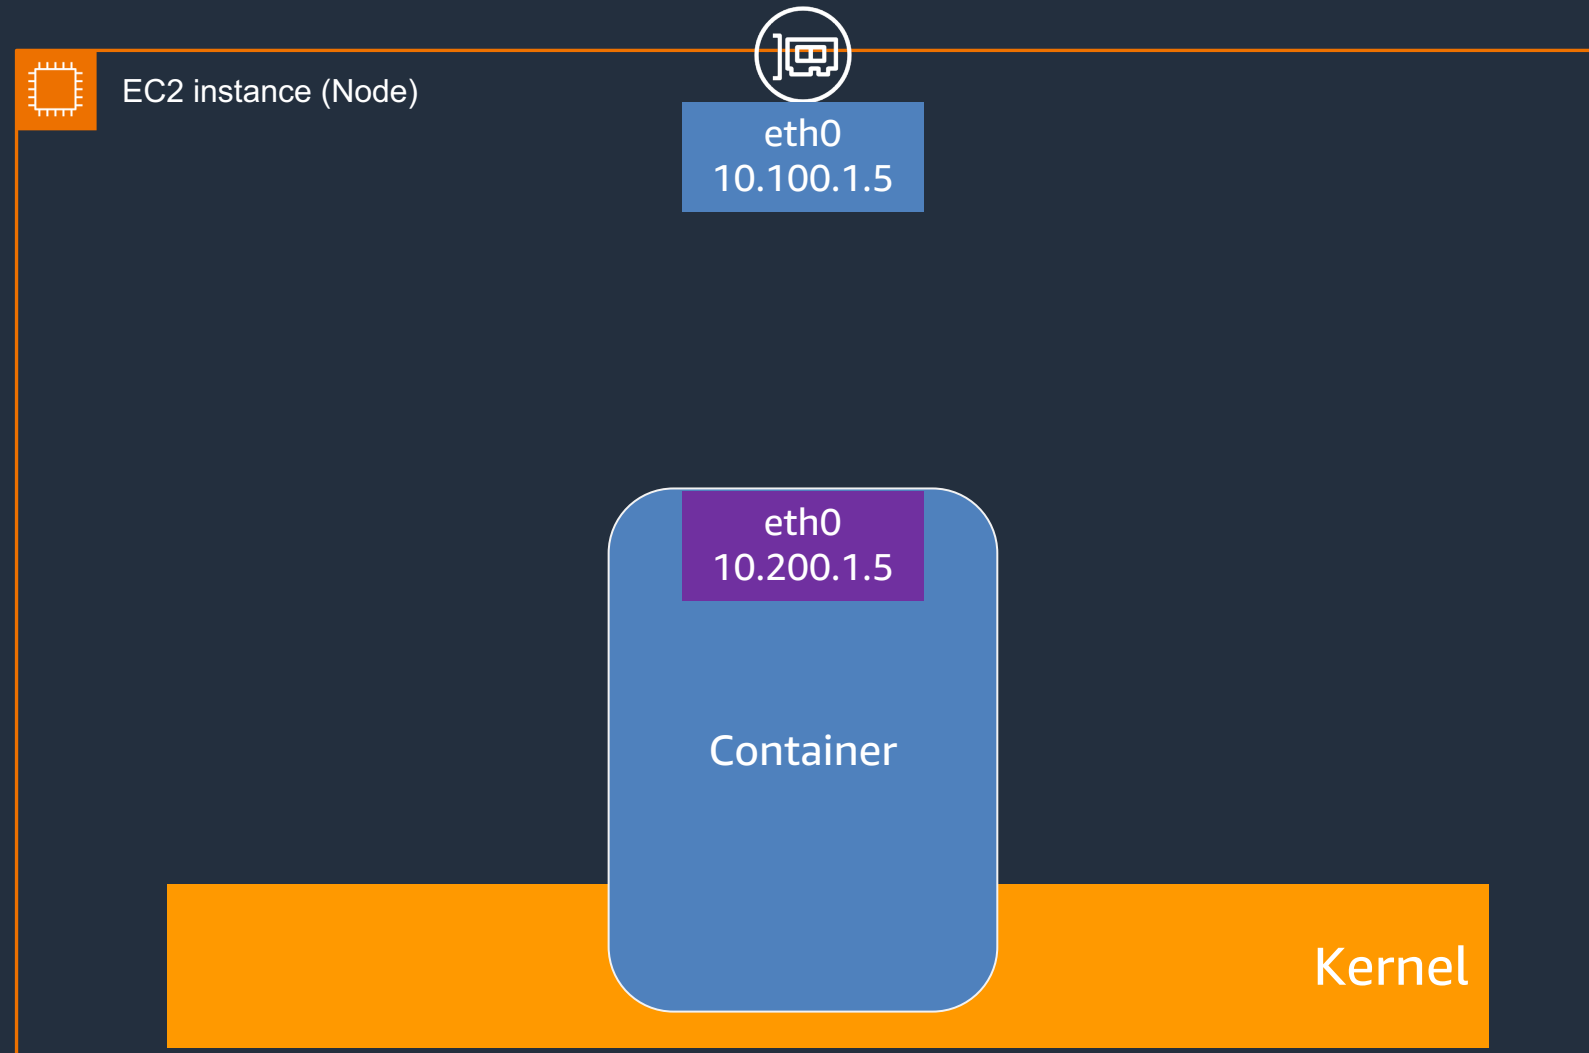
Kubenet

CNI Plugin

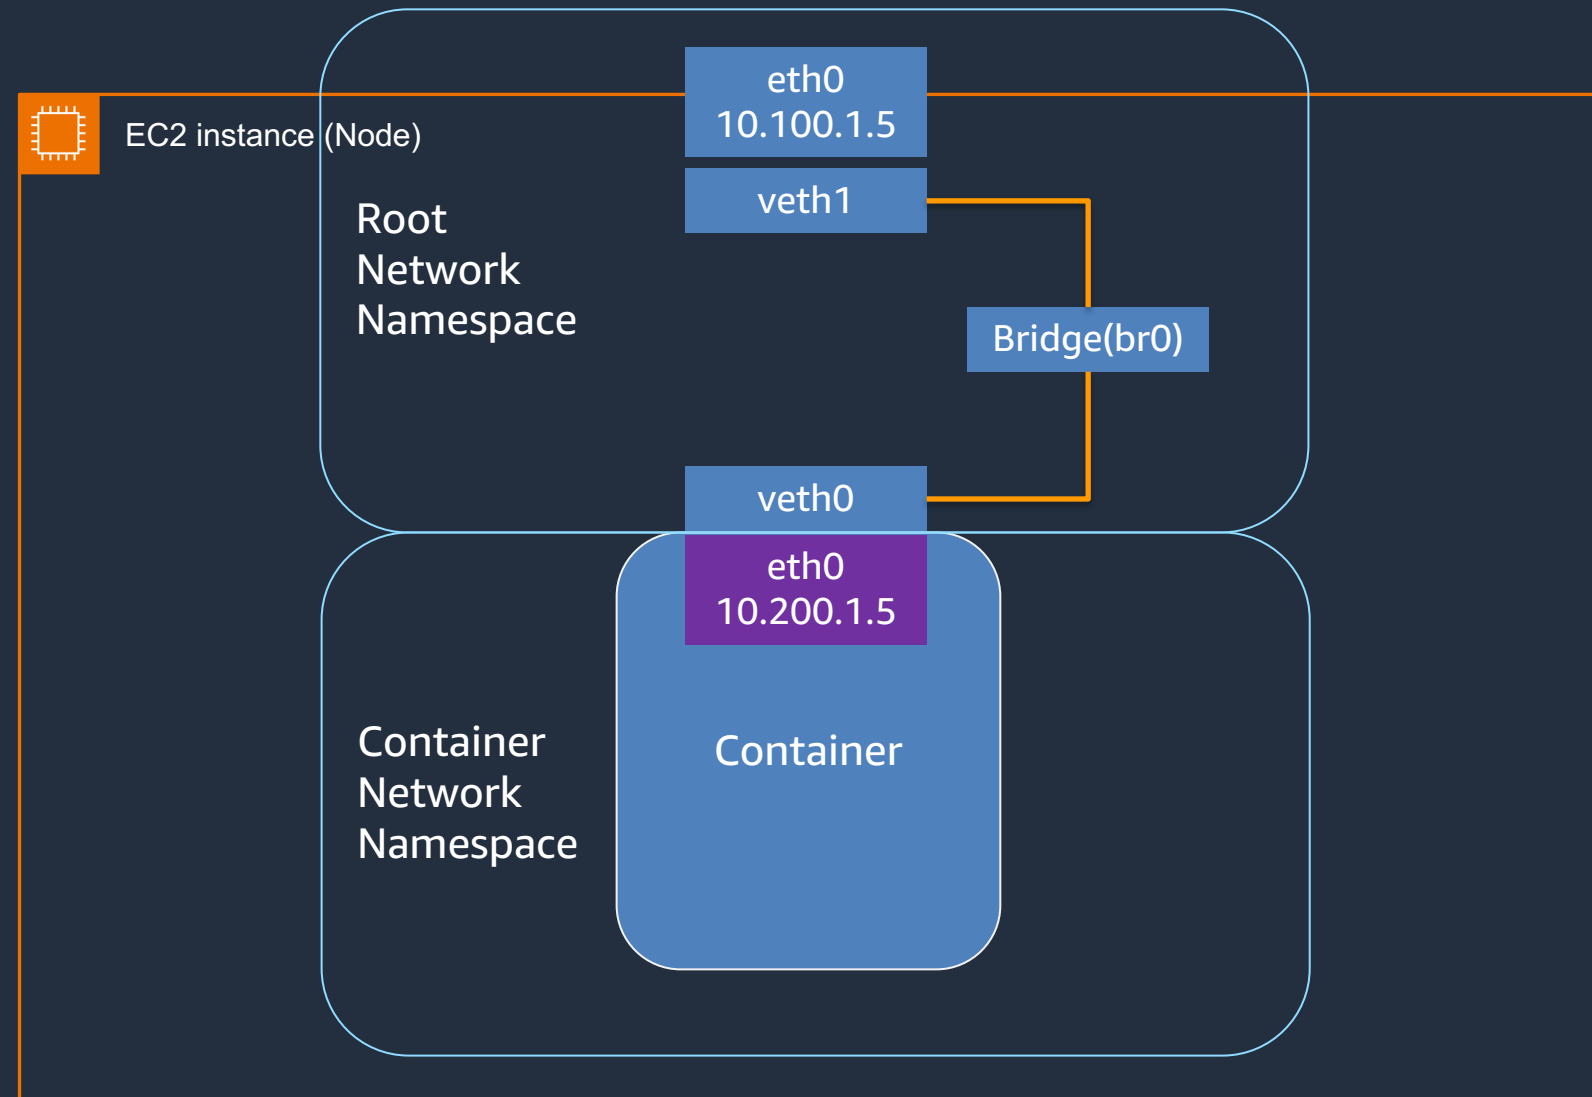
Linux Networking Primitives for Kubernetes



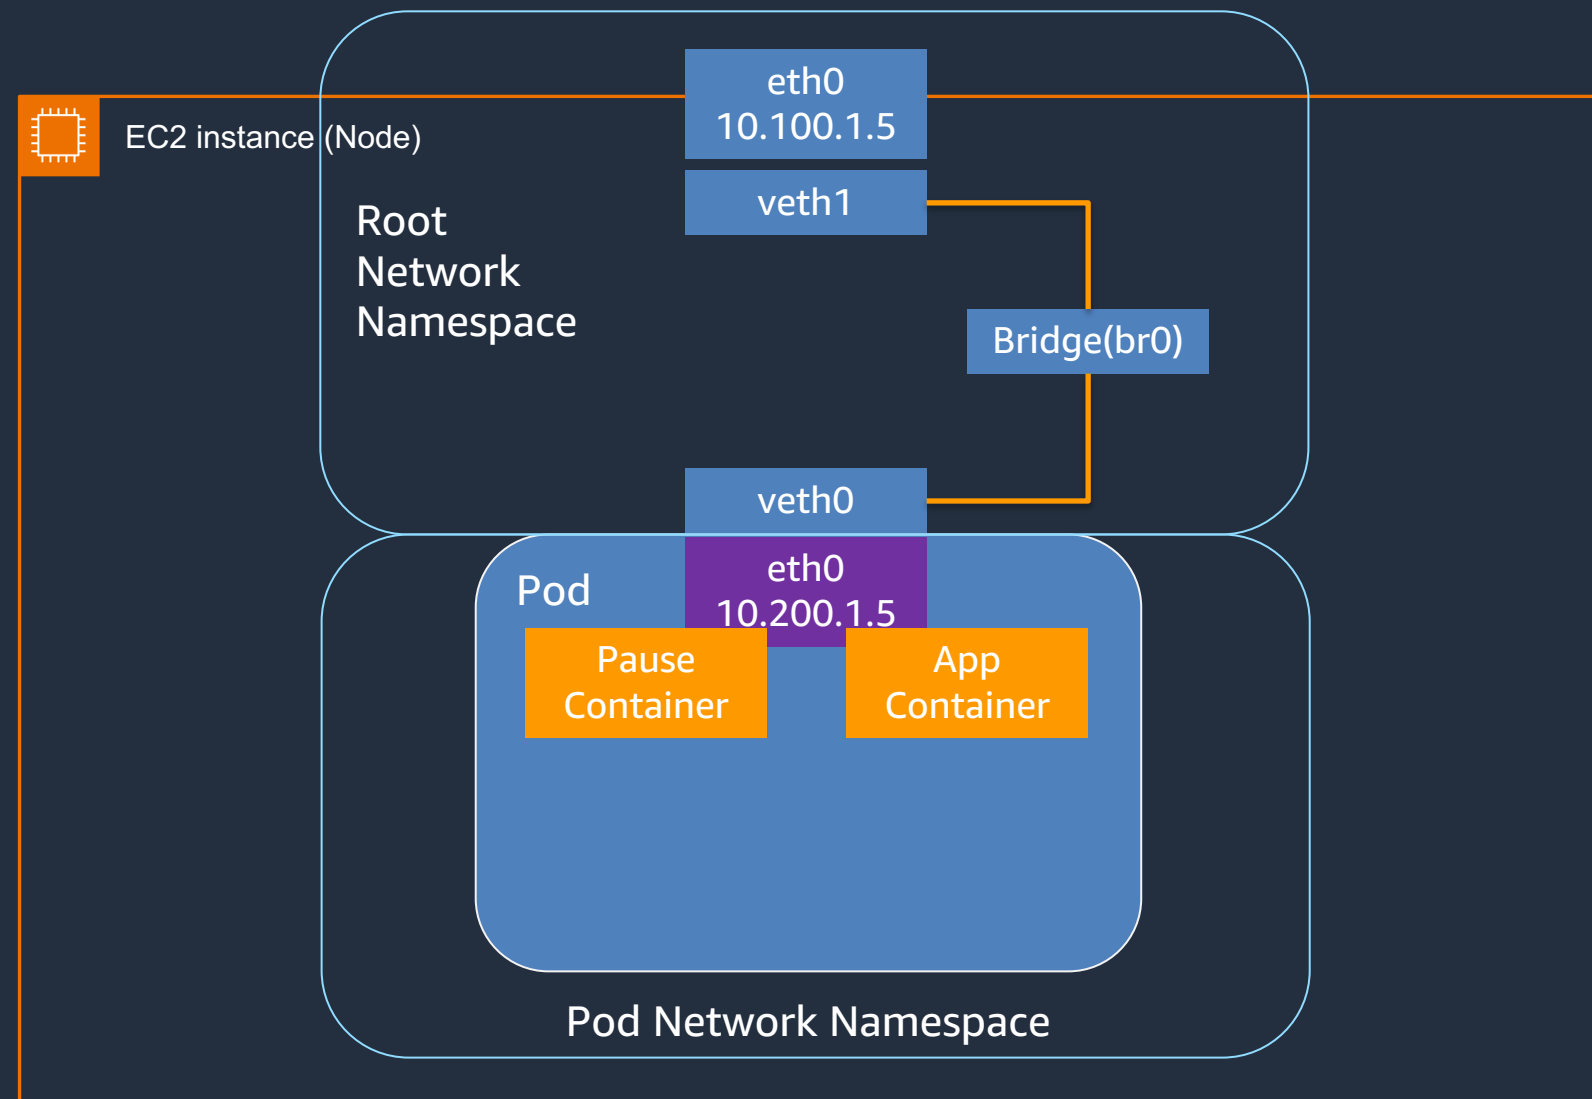
Linux Networking Primitives for Kubernetes



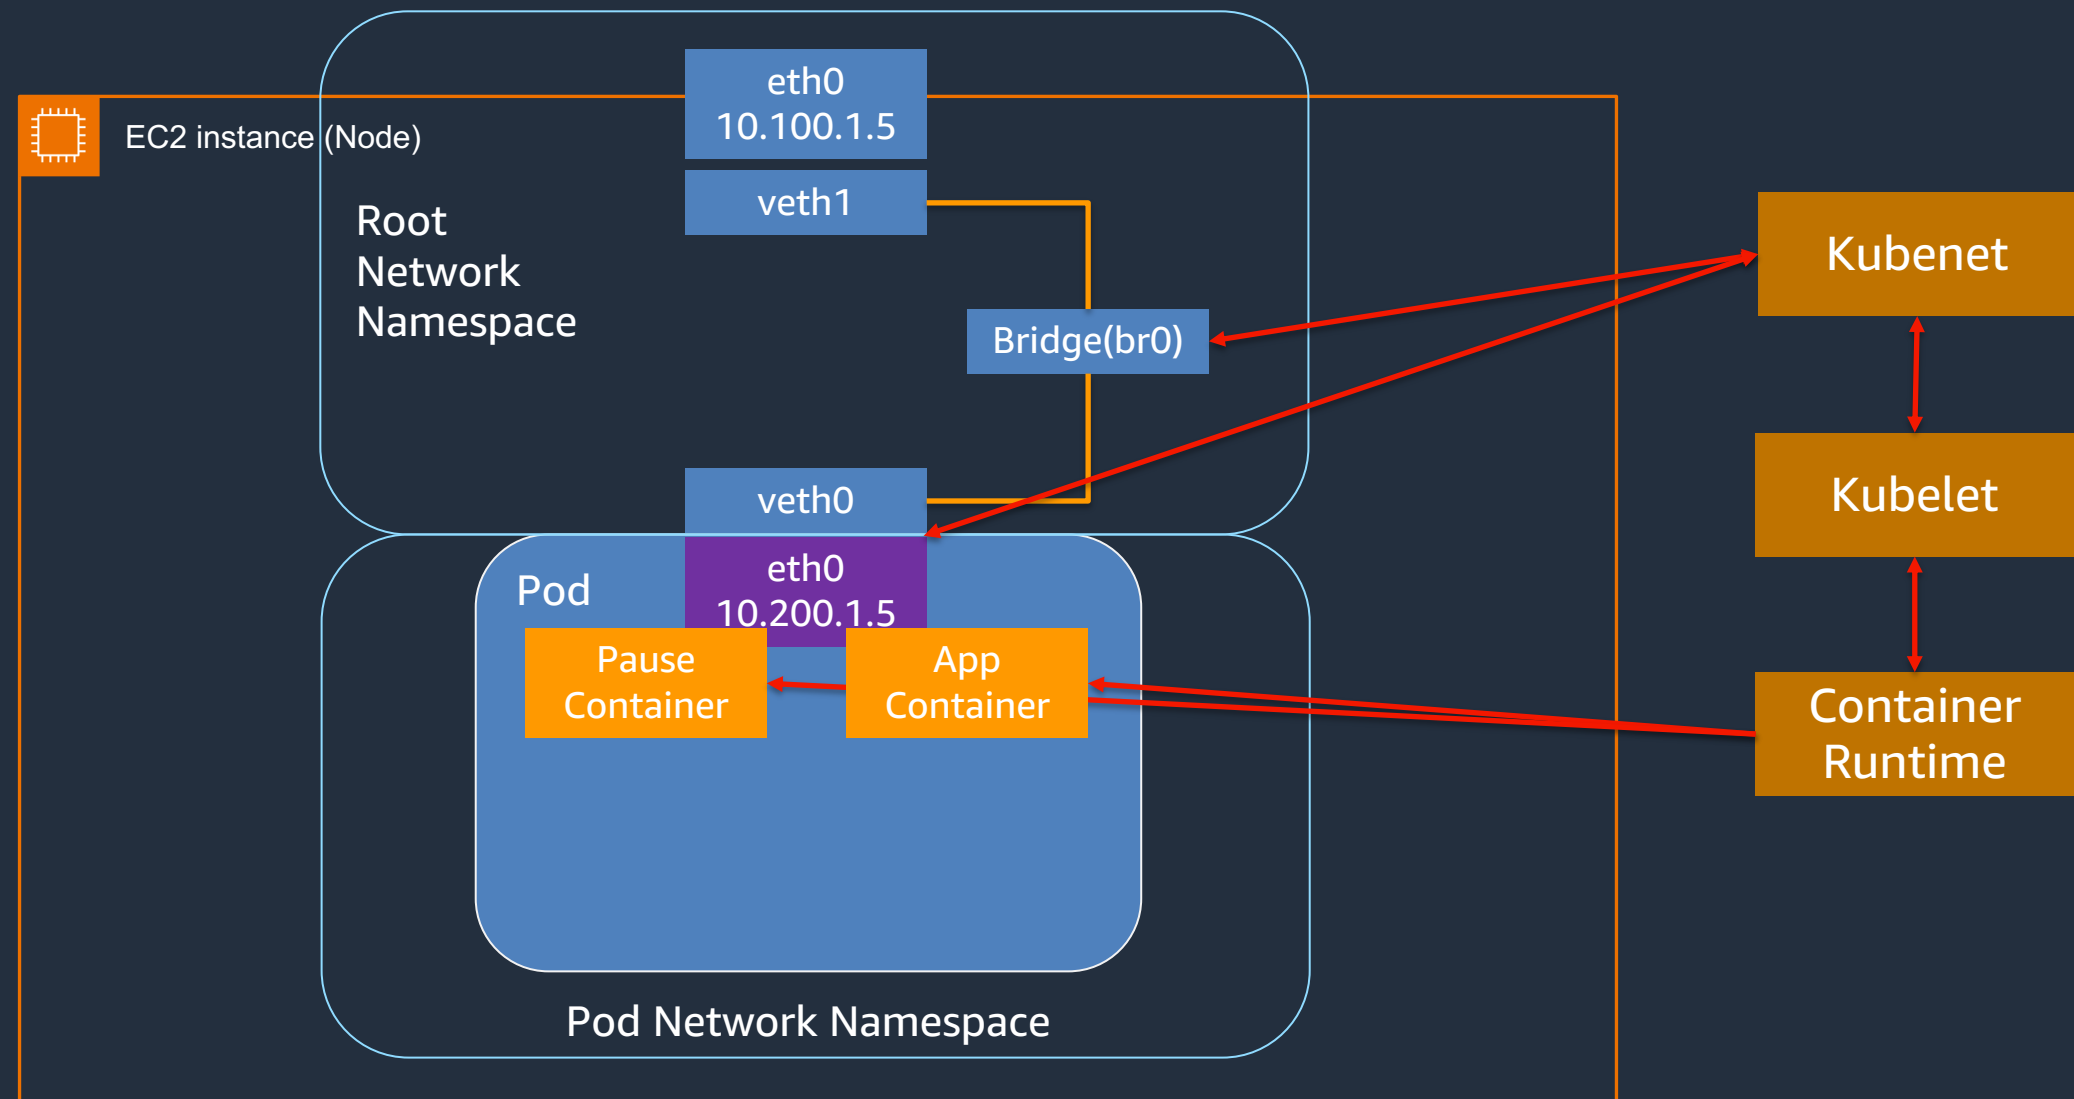
Linux Networking Primitives for Kubernetes



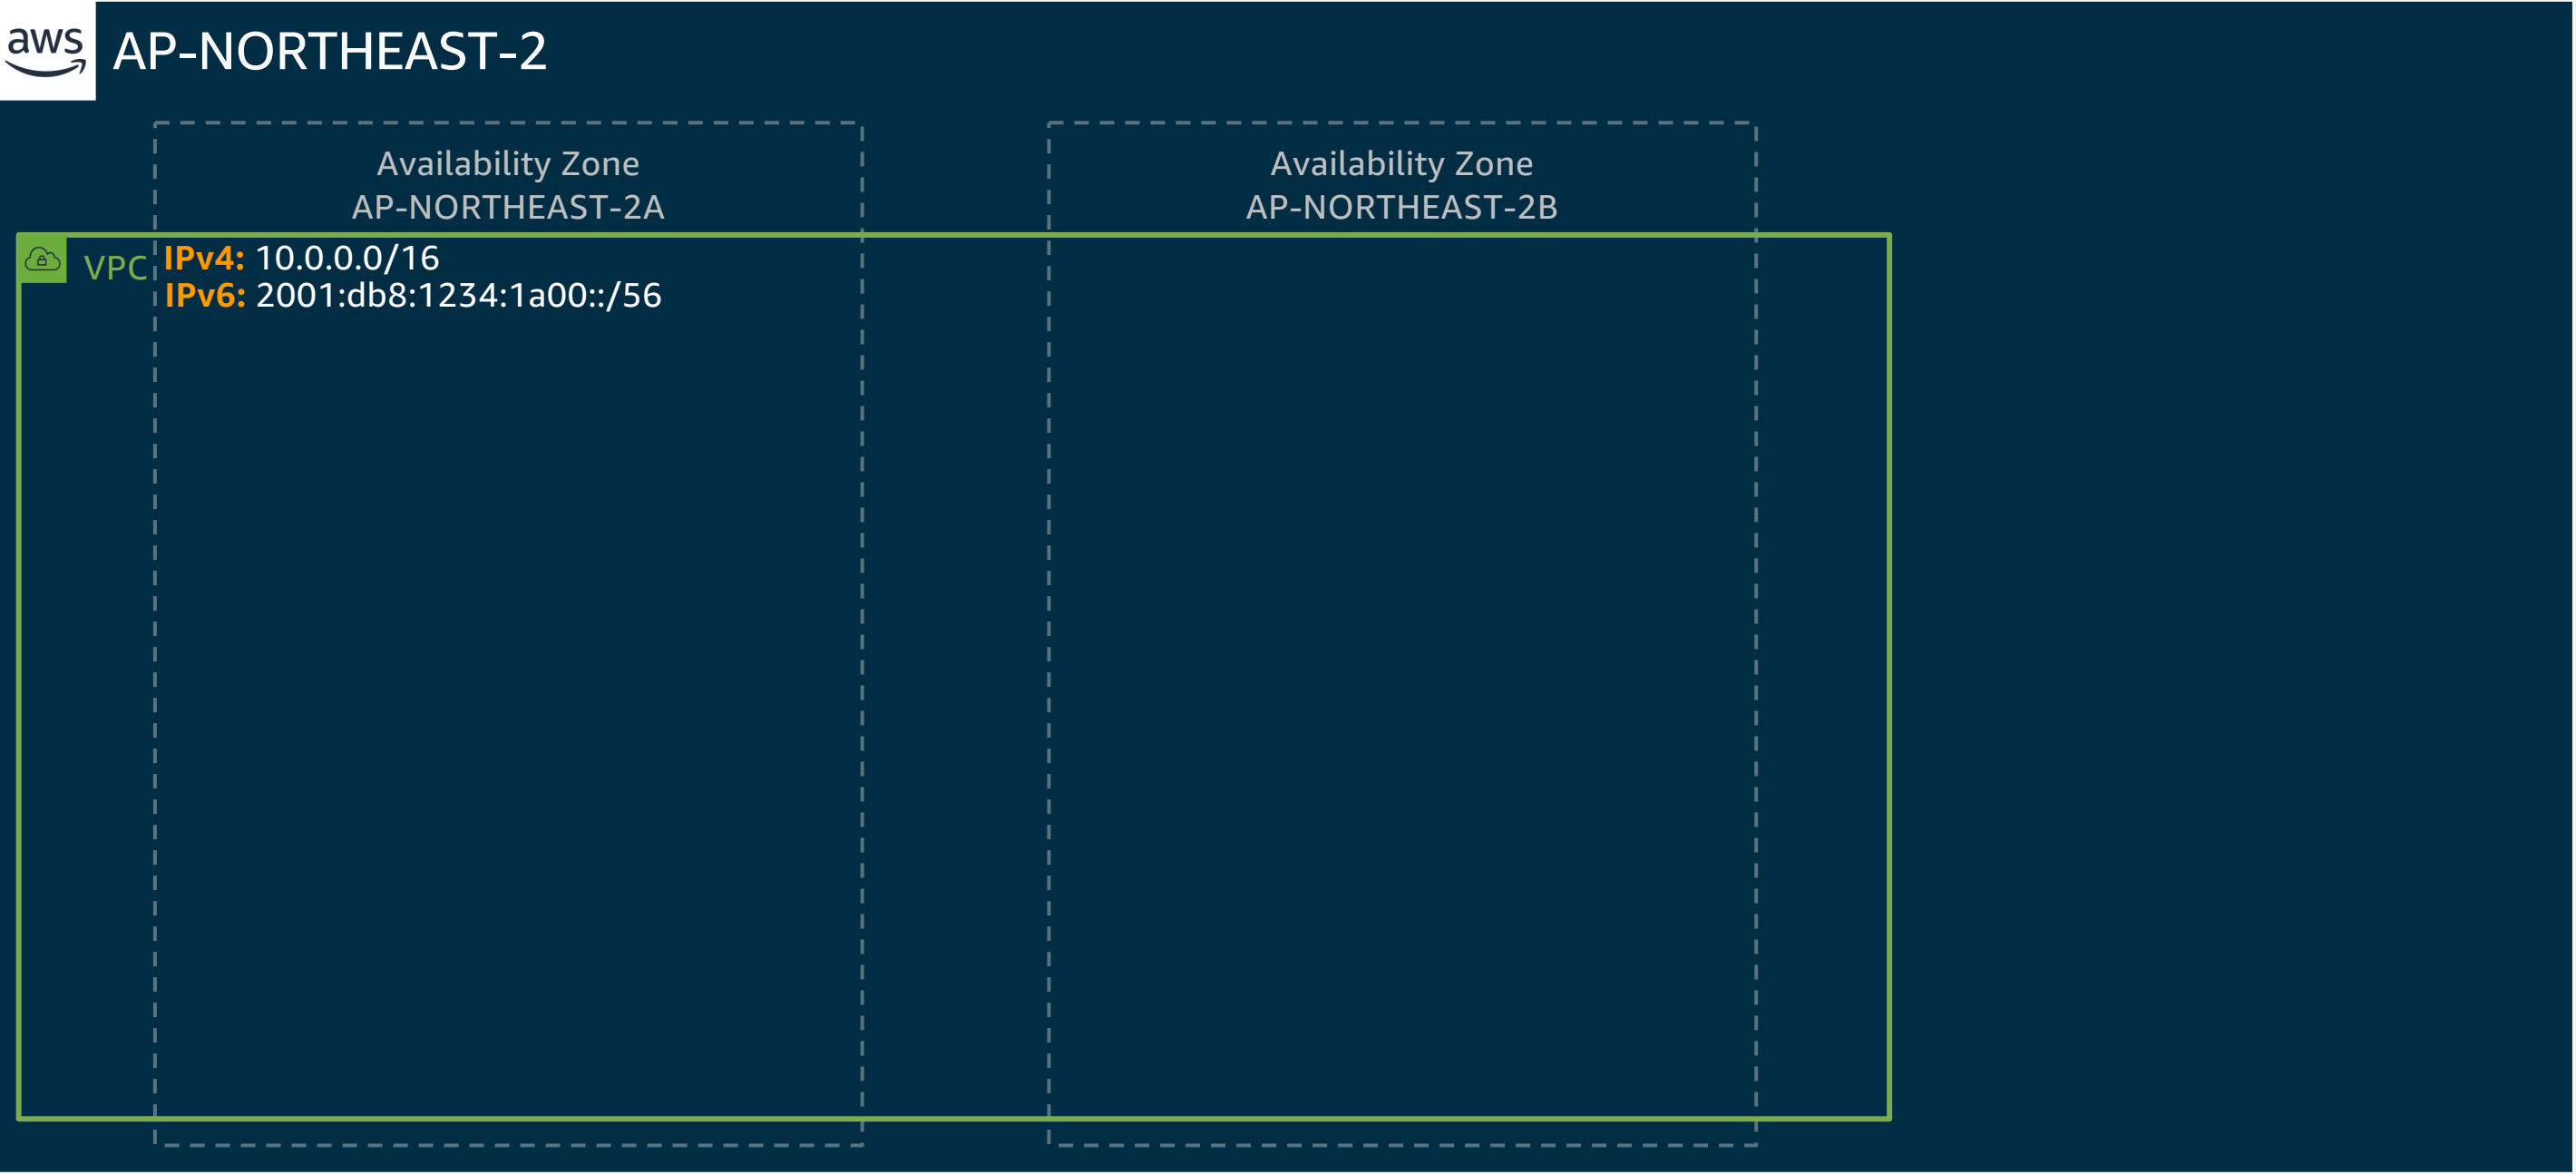
Linux Networking Primitives for Kubernetes



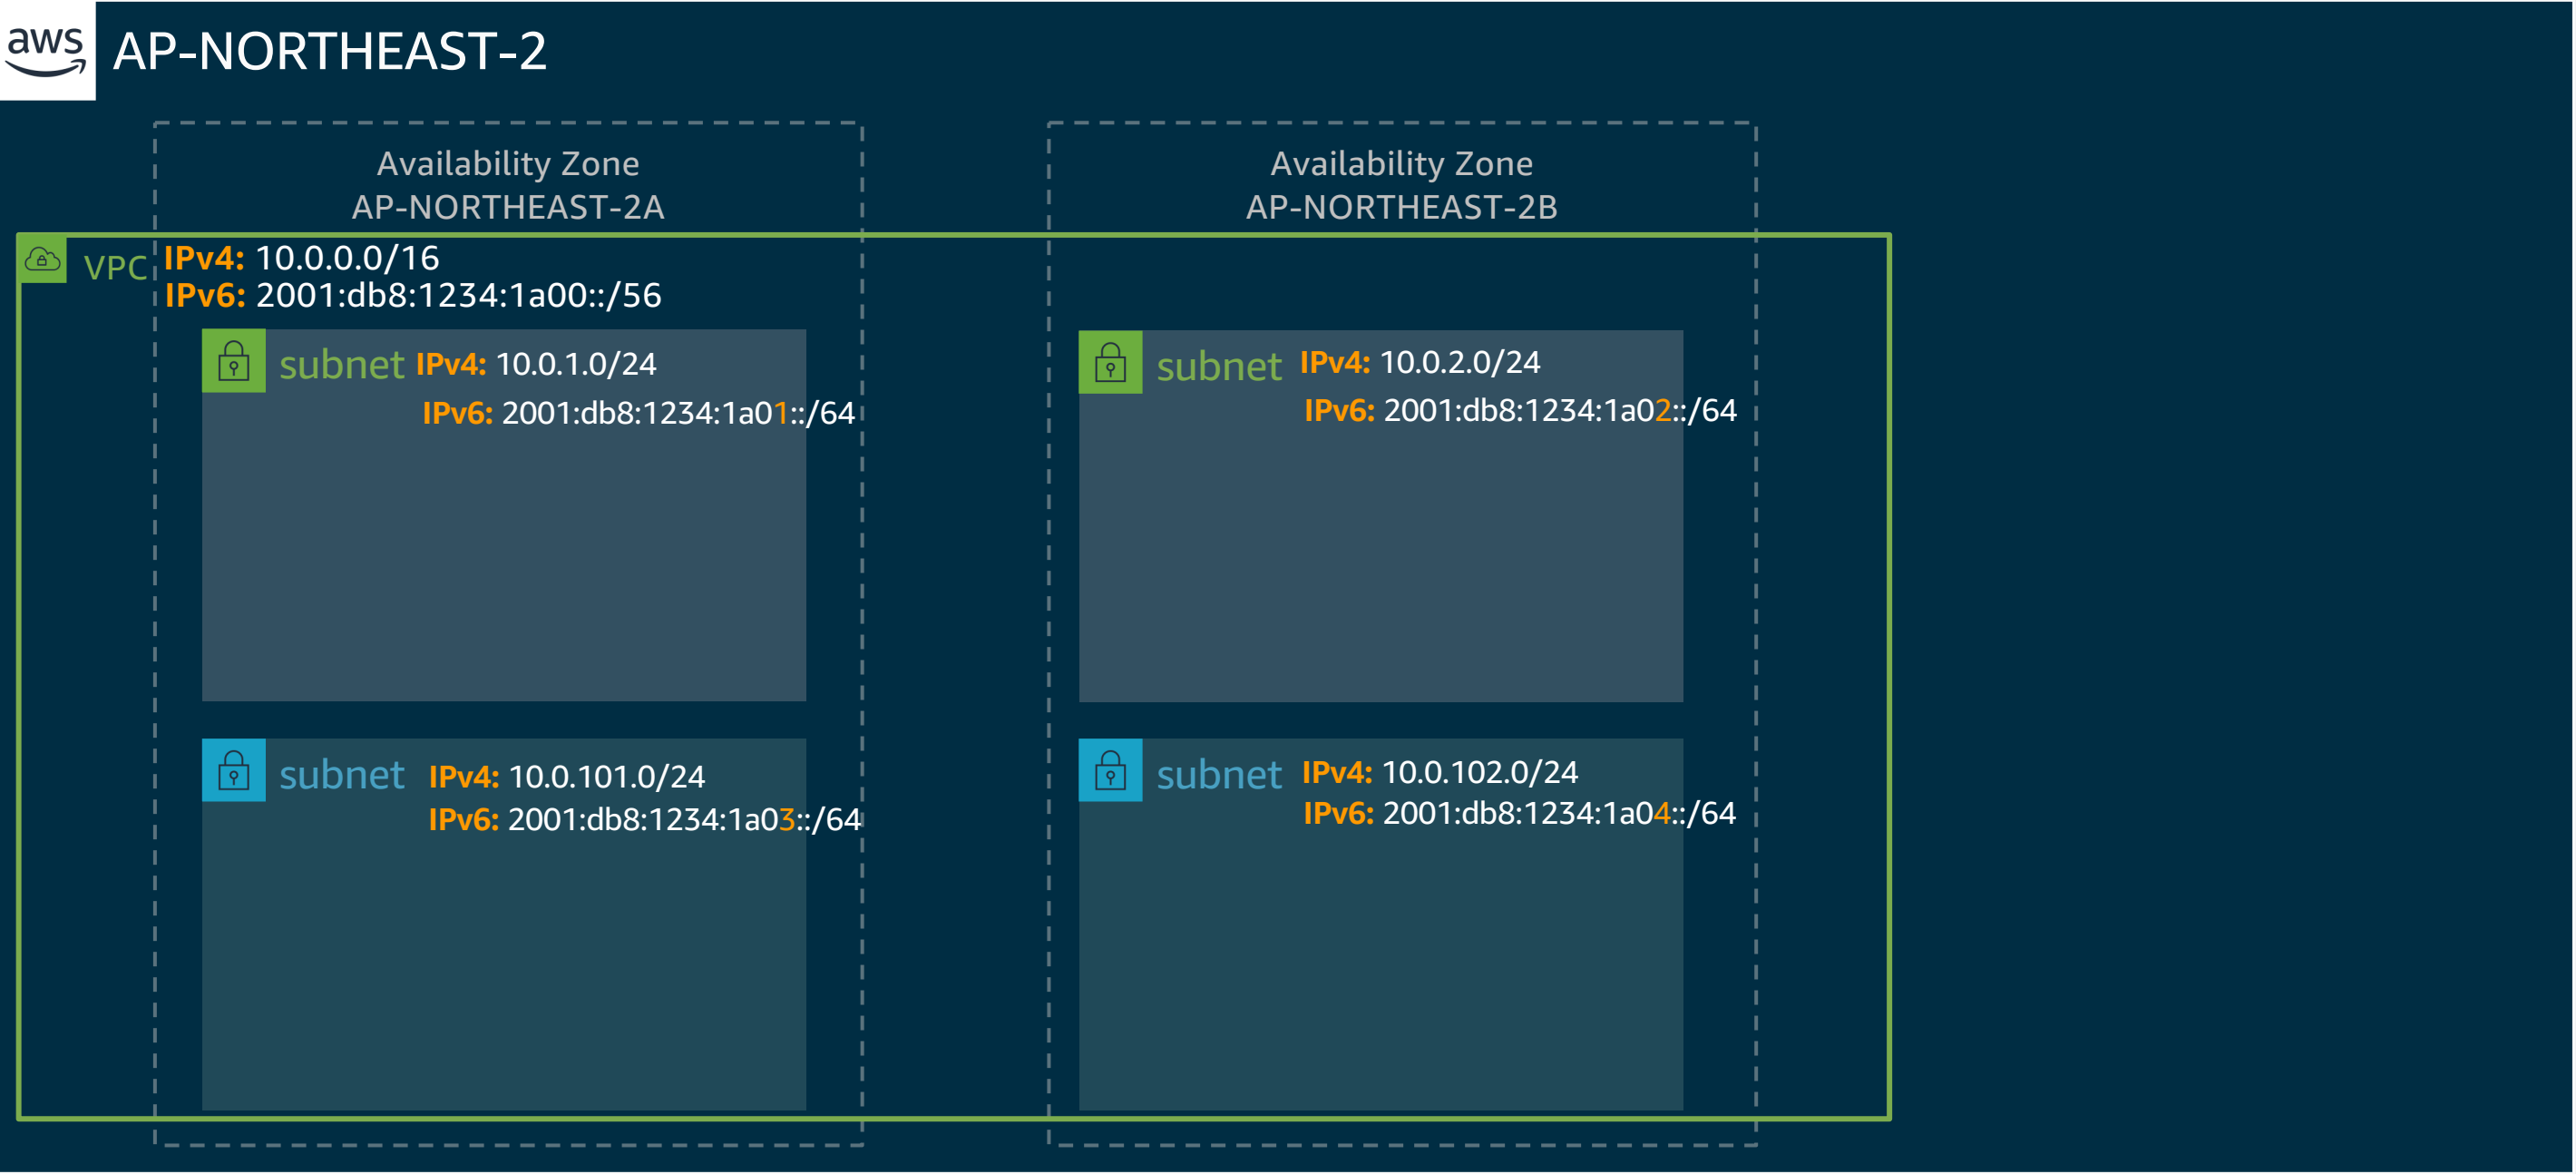
Linux Networking Primitives for Kubernetes



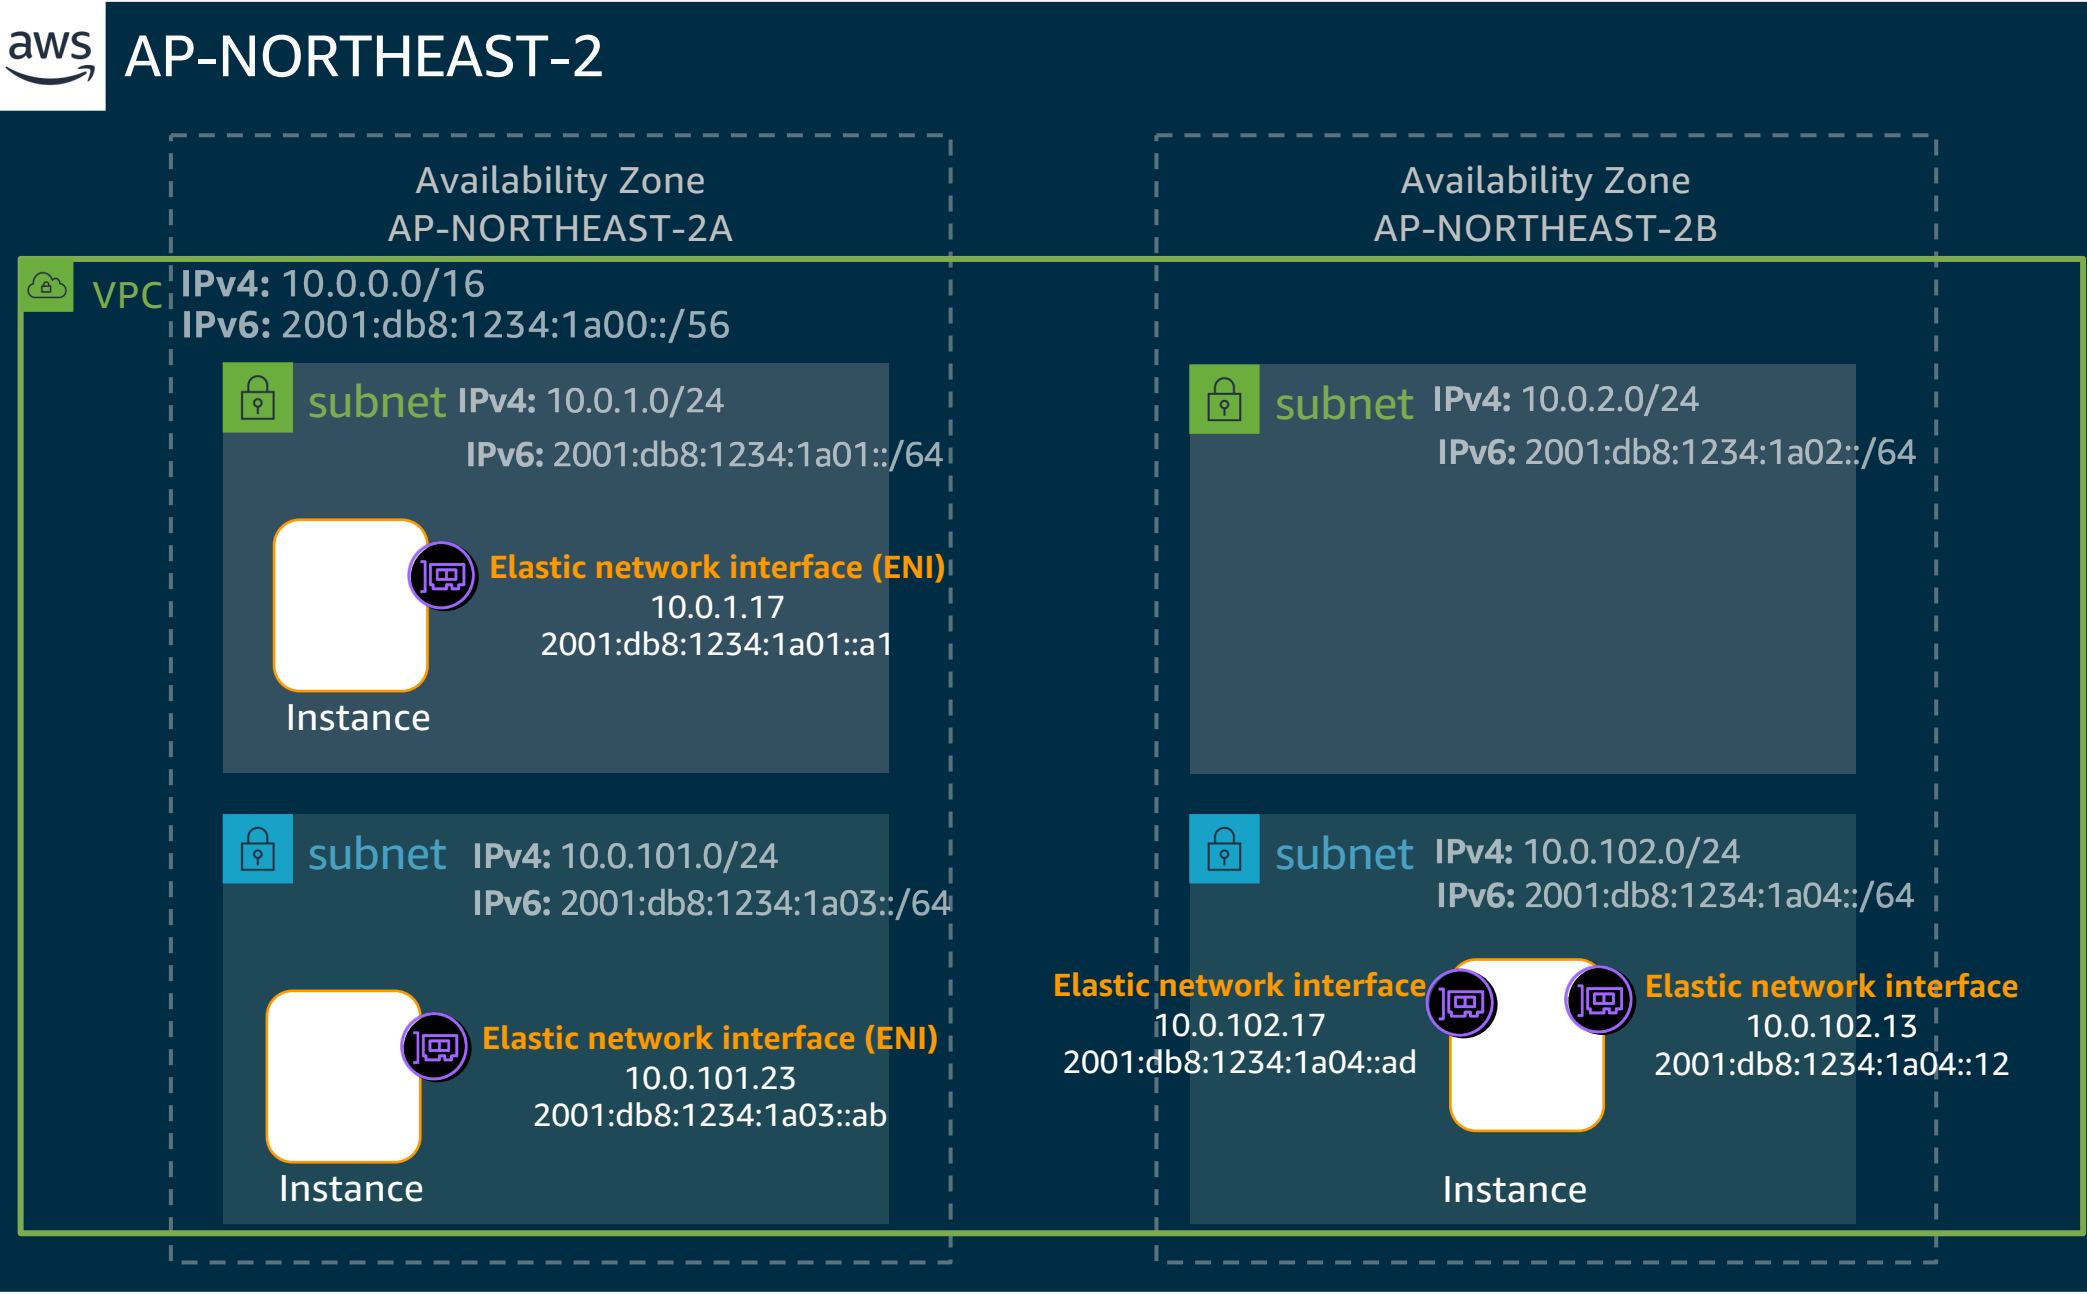
Amazon VPC – IP range (CIDR)



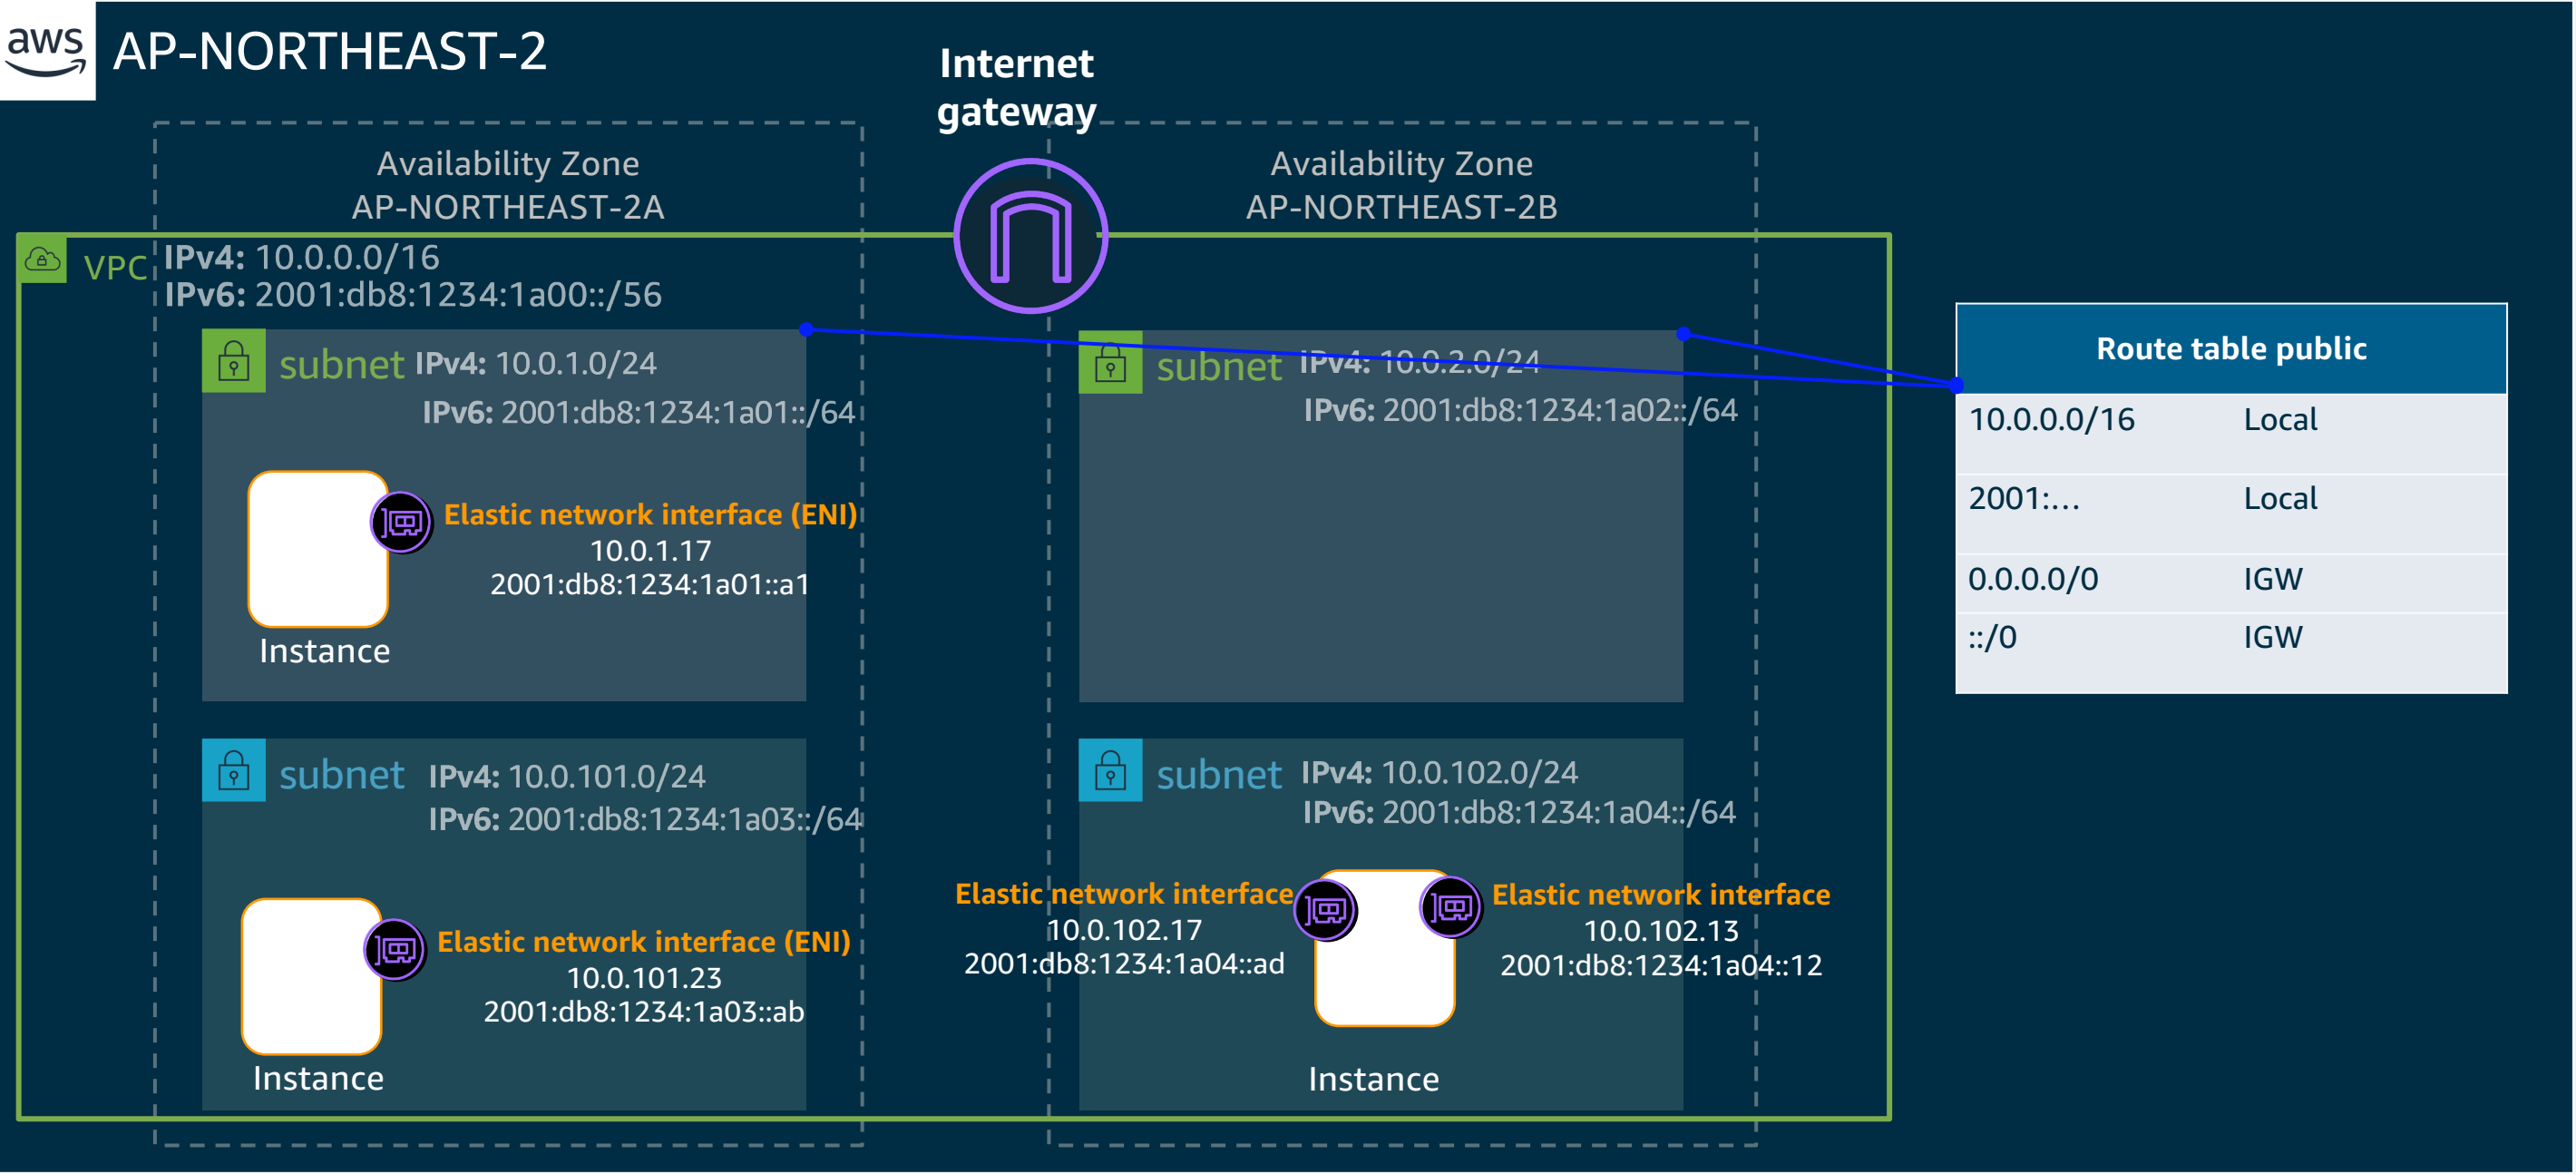
Amazon VPC – Subnets



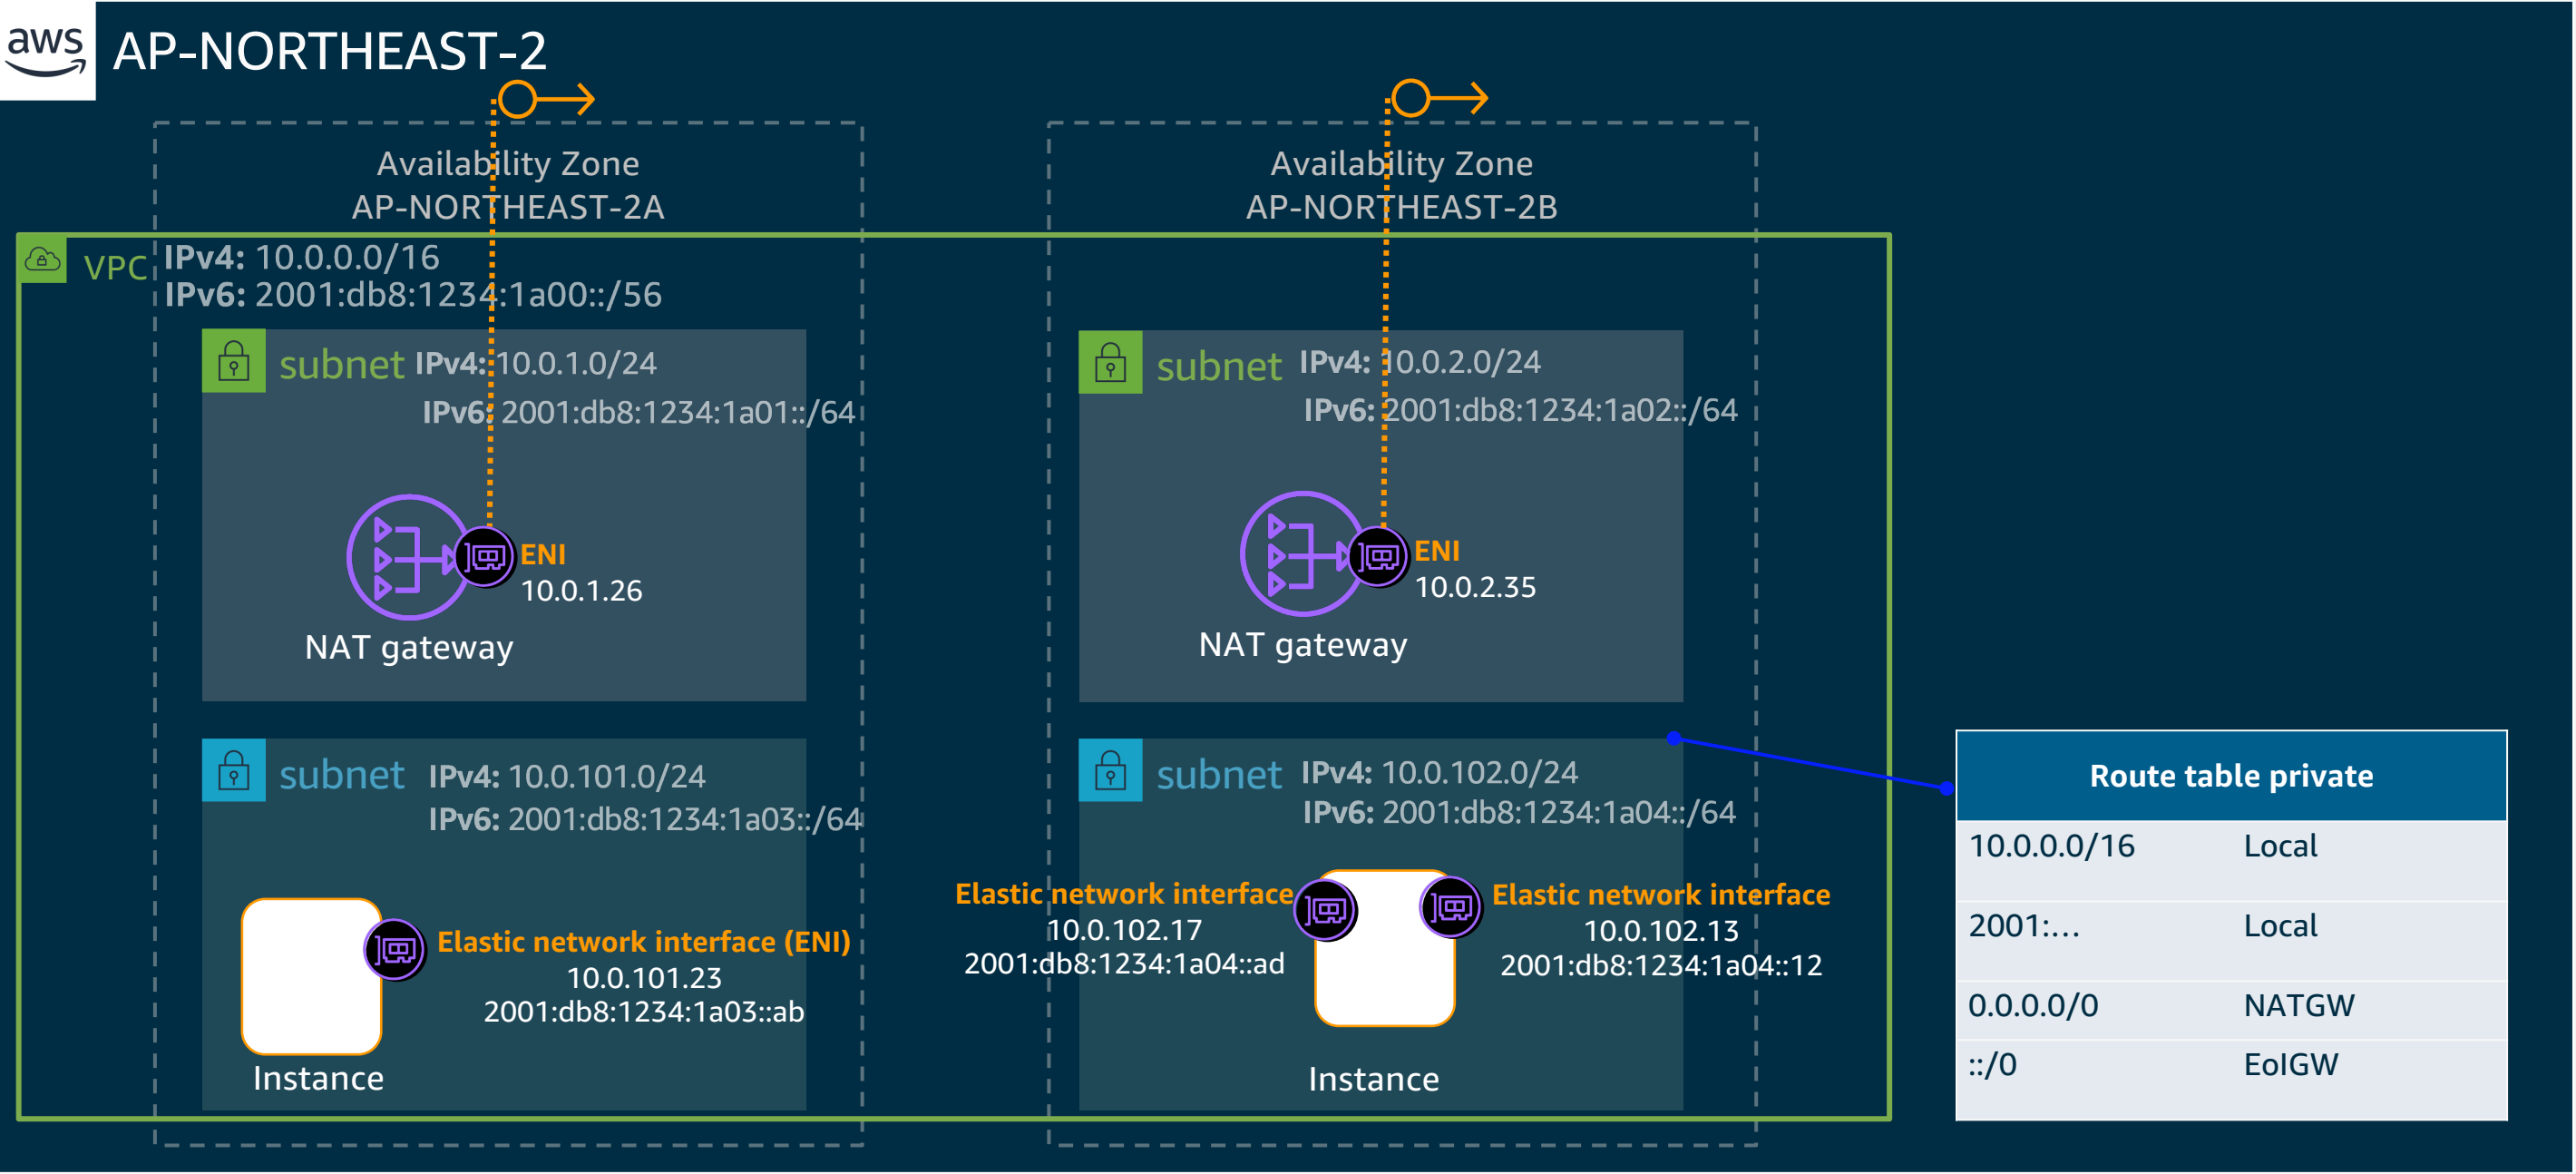
Amazon VPC – Interfaces




Amazon VPC – Connectivity to Internet



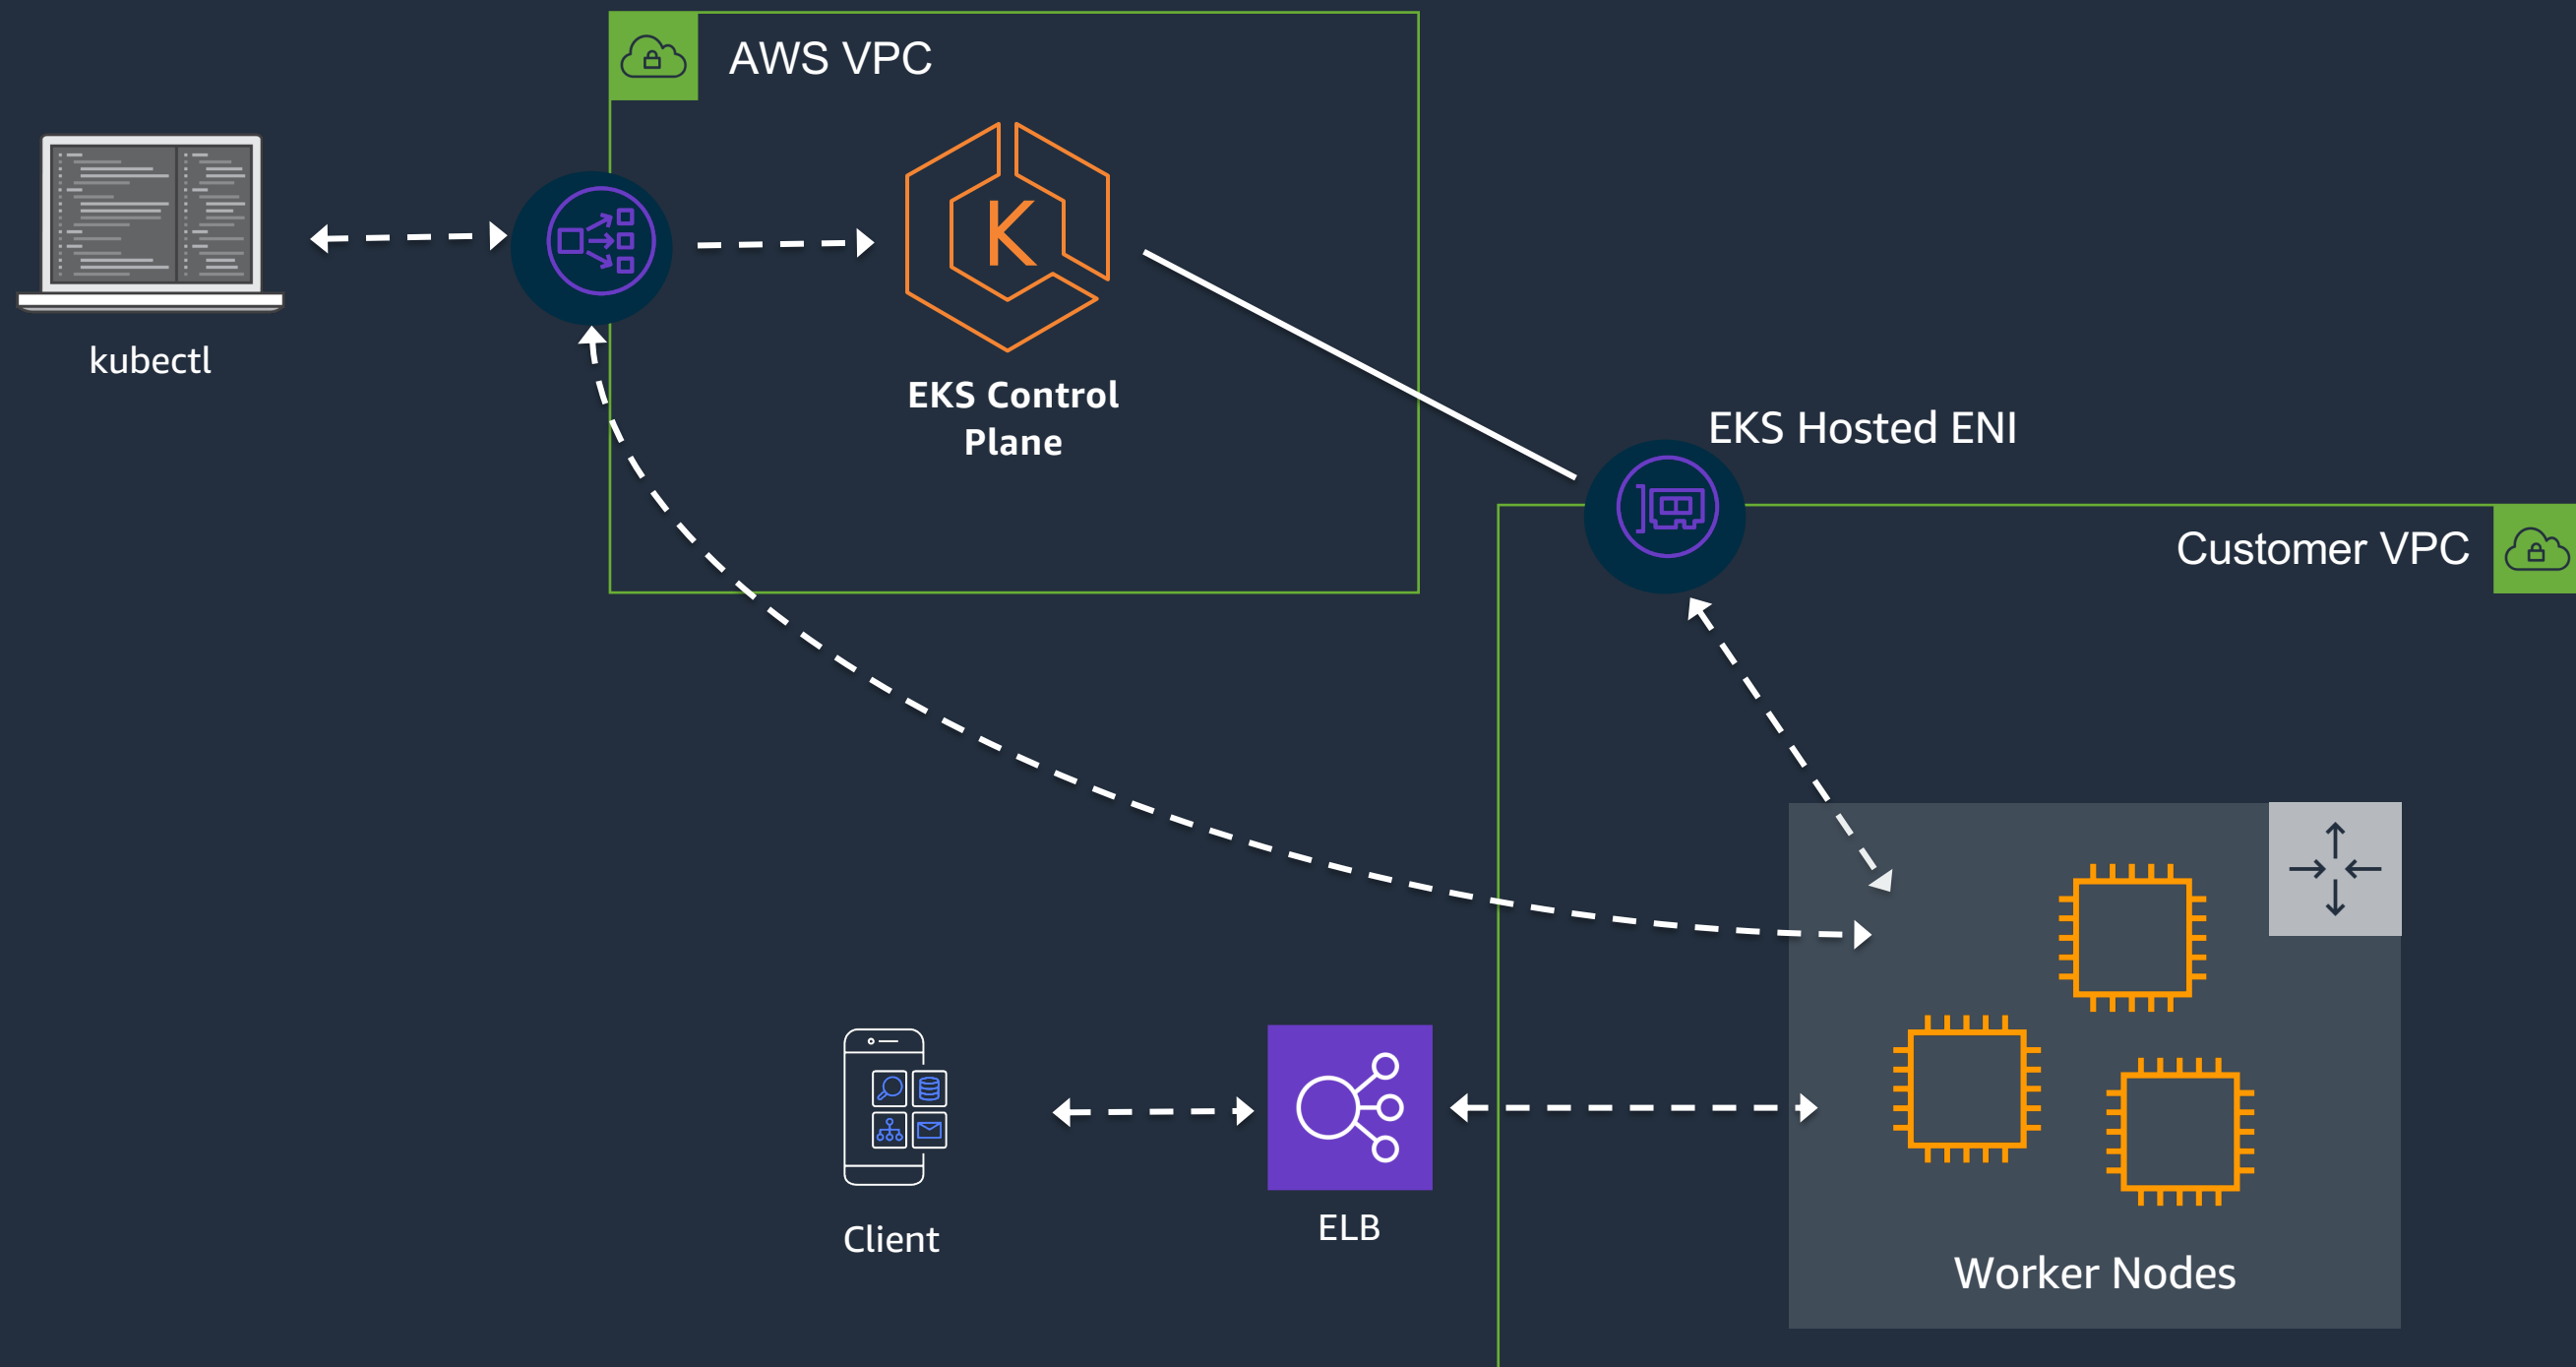
Amazon VPC – Connectivity to Internet



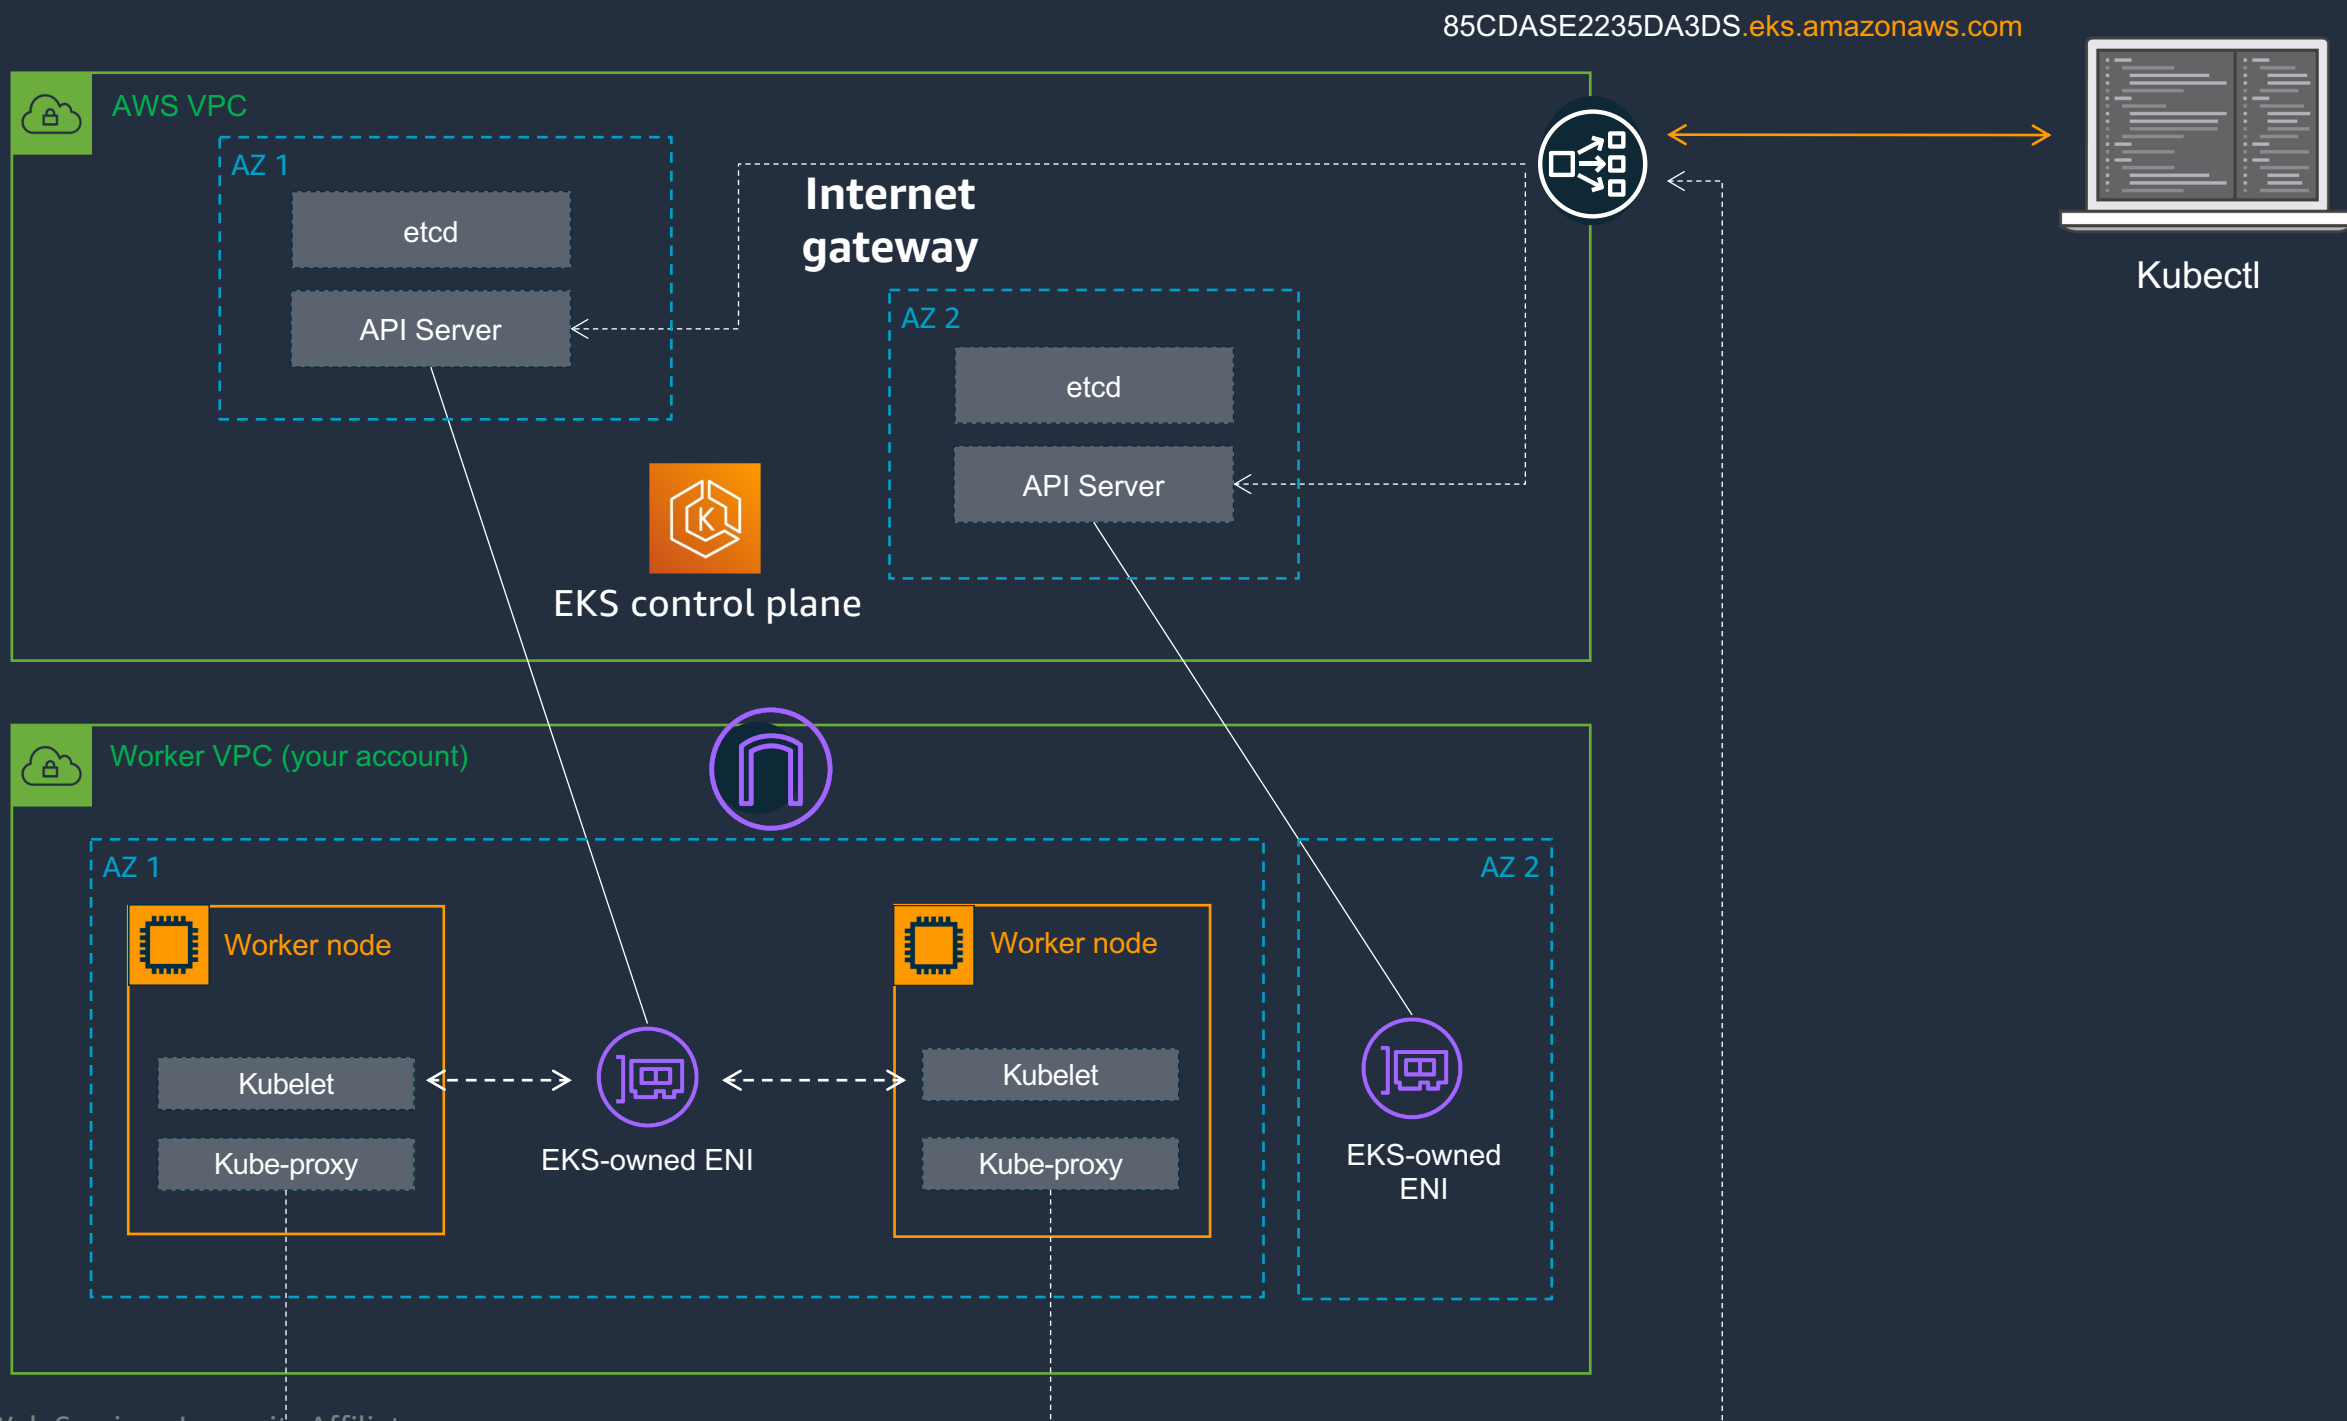
EKS Control Plane Network



EKS Network Architecture

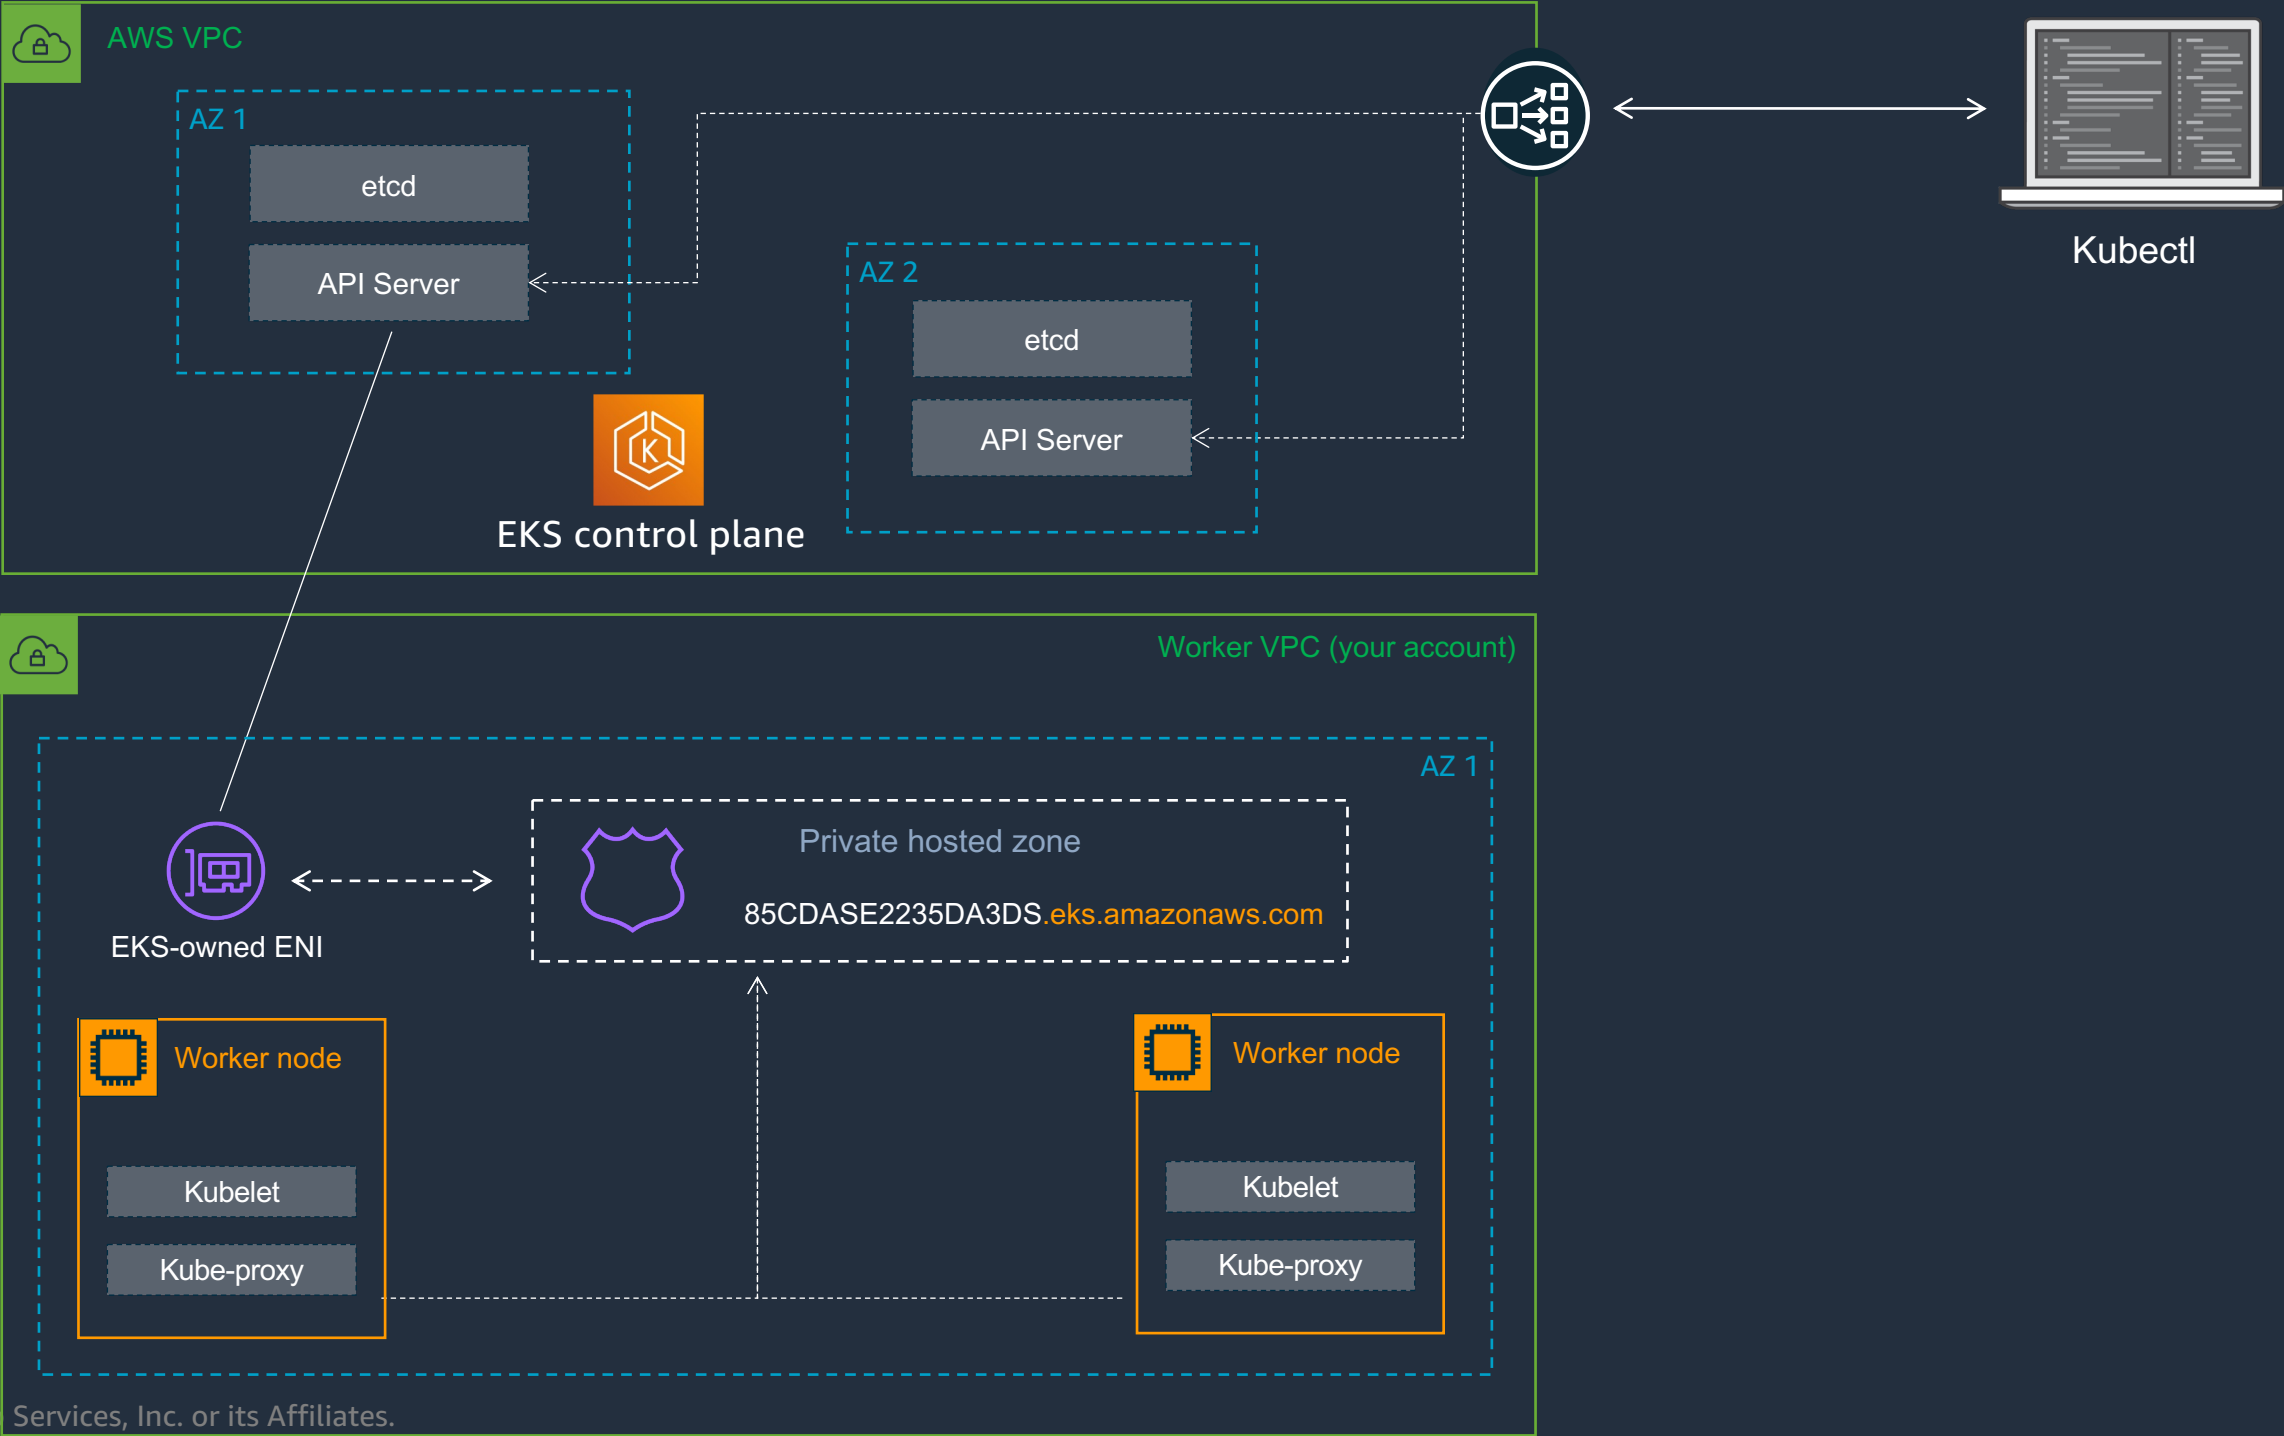


EKS cluster endpoint access – Public

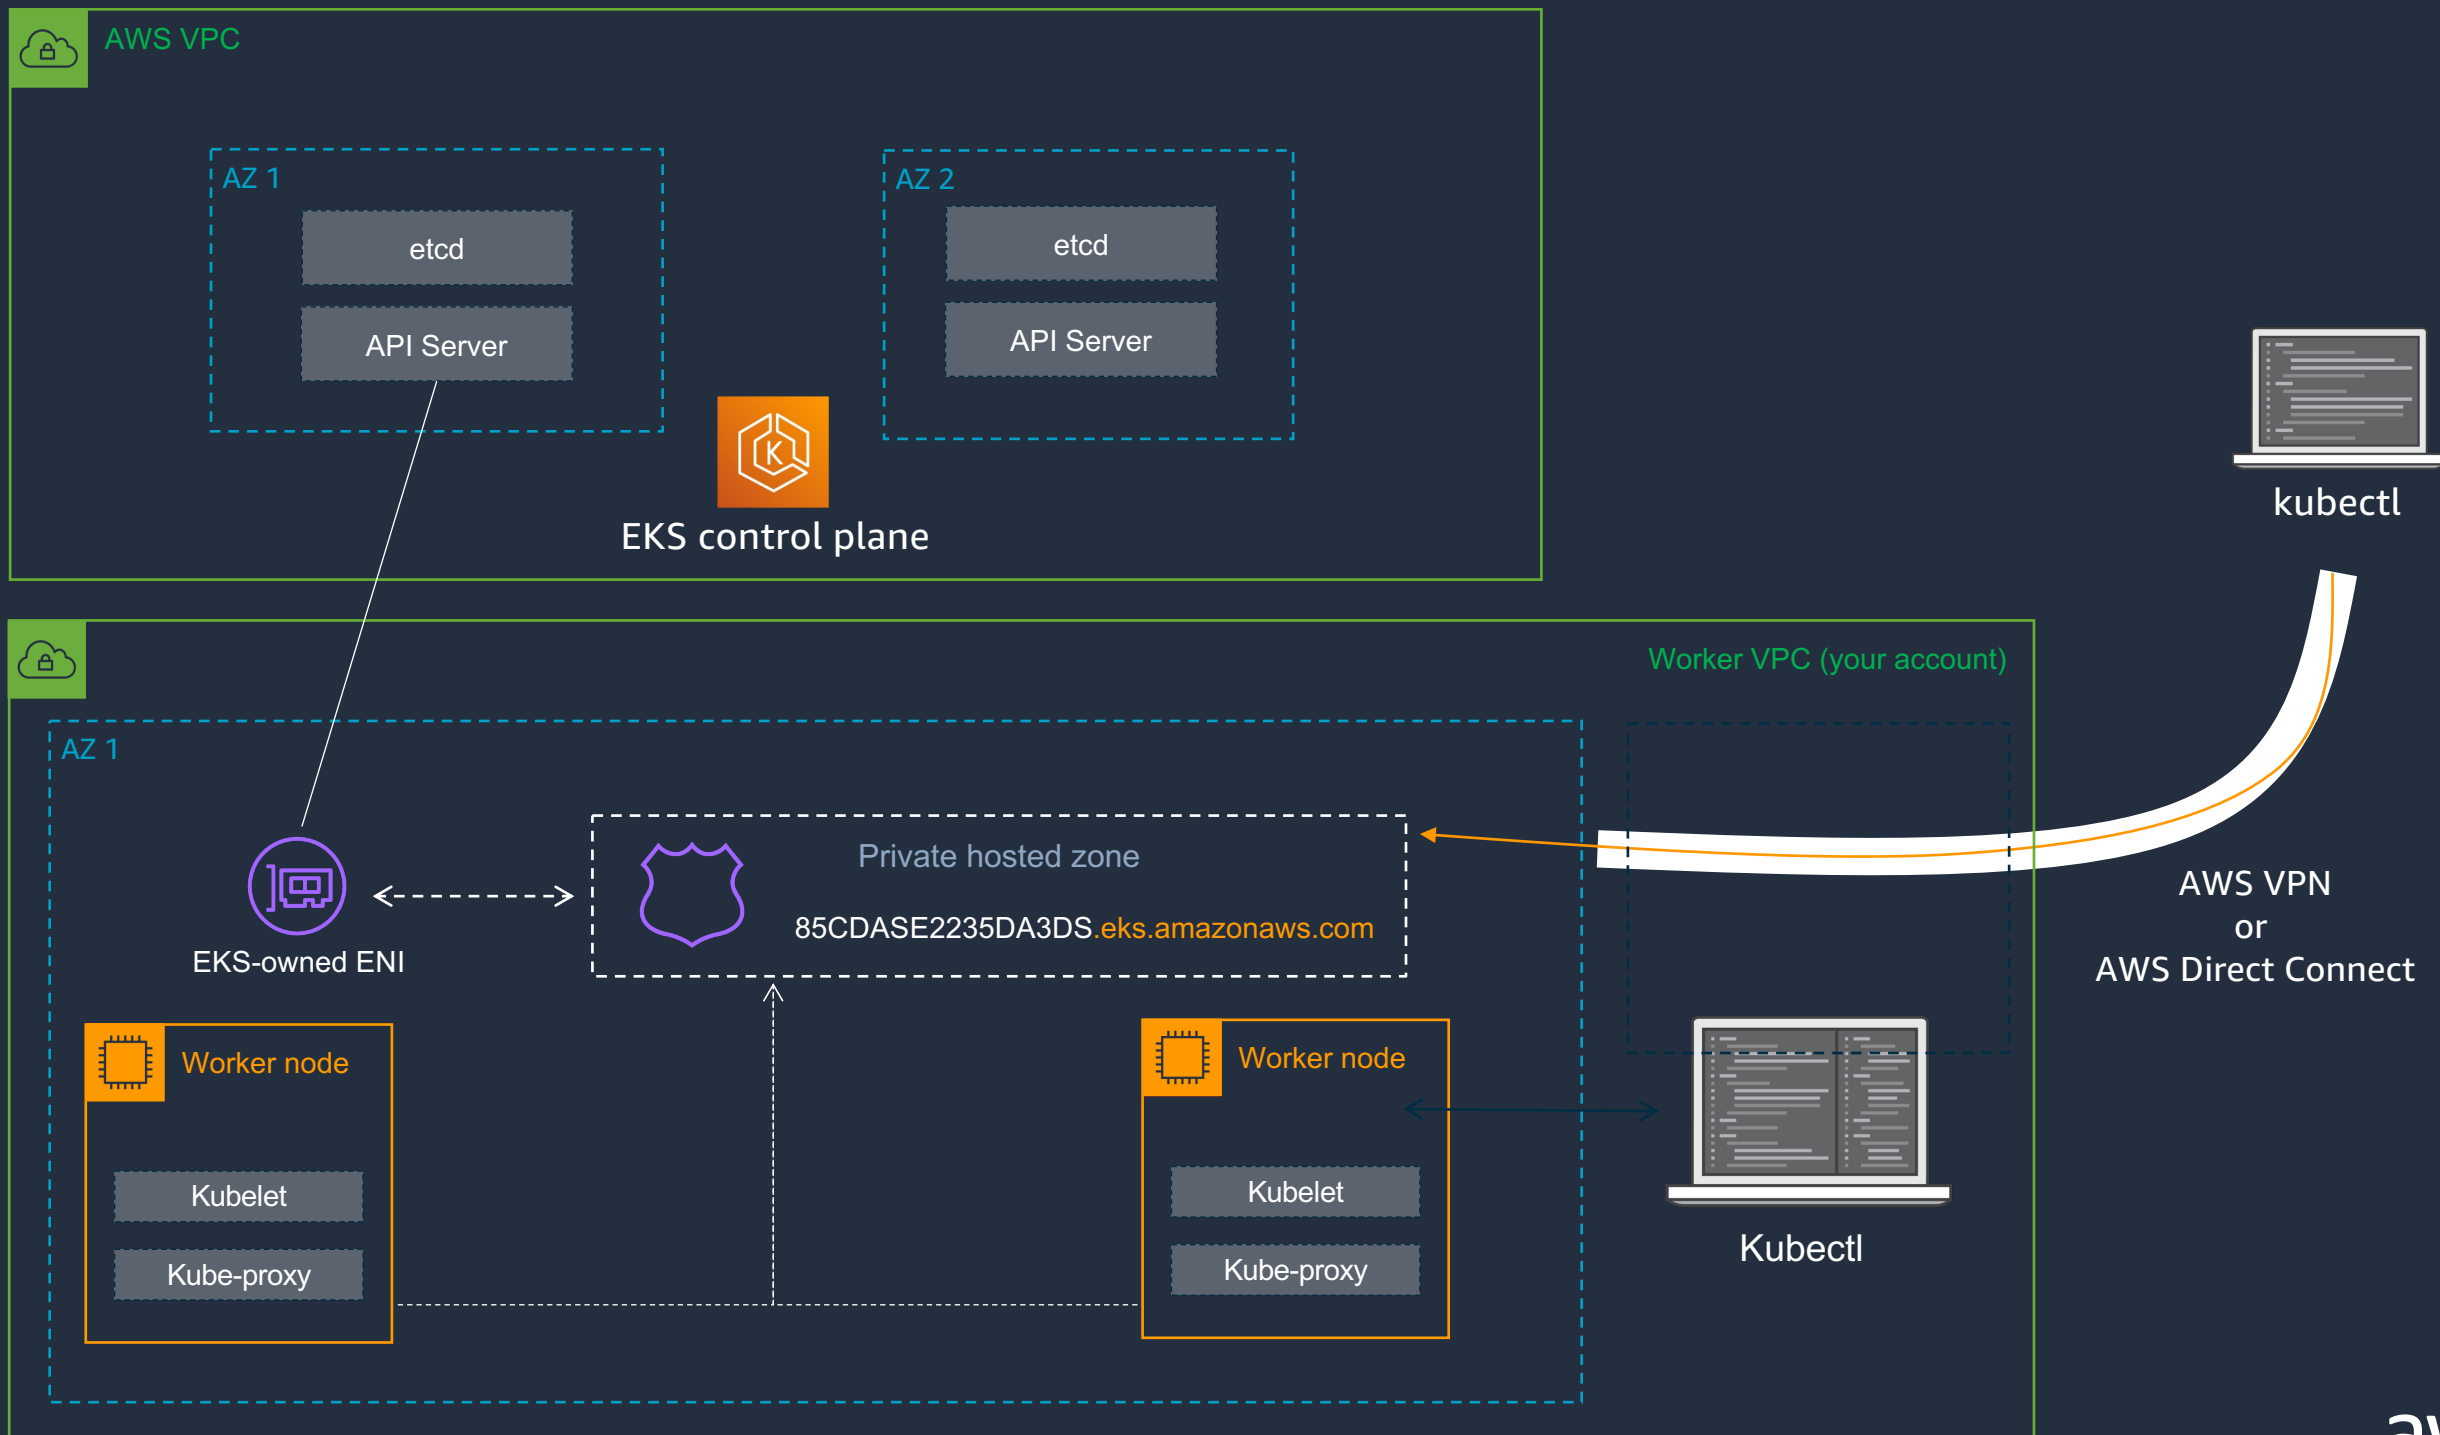


EKS cluster endpoint access – Public & private


85CDASE2235DA3DS.eks.amazonaws.com



EKS cluster endpoint access – Private only



EKS cluster endpoint access



서비스

서비스, 기능, 블로그, 설명서 등을 검색합니다. [Option+S]

EKS > 클러스터 > eksctl-cluster > 네트워킹 관리

네트워킹 관리: eksctl-cluster

클러스터 엔드포인트 액세스 정보

Kubernetes API 서버 엔드포인트에 대한 액세스 권한을 구성합니다.

☒ 퍼블릭

VPC 외부에서 클러스터 엔드포인트에 액세스할 수 있습니다. 작업자 노드 트래픽은 엔드포인트에 연결하기 위해 VPC를 벗어납니다.

☐ 퍼블릭 및 프라이빗

VPC 외부에서 클러스터 엔드포인트에 액세스할 수 있습니다. 엔드포인트에 대한 작업자 노드 트래픽은 VPC 내에 유지됩니다.

☐ 프라이빗

클러스터 엔드포인트는 VPC를 통해서만 액세스할 수 있습니다. 엔드포인트에 대한 작업자 노드 트래픽은 VPC 내에 유지됩니다.

▼ 고급 설정

퍼블릭 액세스 엔드포인트에 소스를 추가/편집합니다. [정보](#)

모든 트래픽에 개방된 API 엔드포인트입니다(0.0.0.0/0).


소스 추가

추가할 수 있는 나머지 CIDR 블록 수: 40

취소

변경 사항 저장

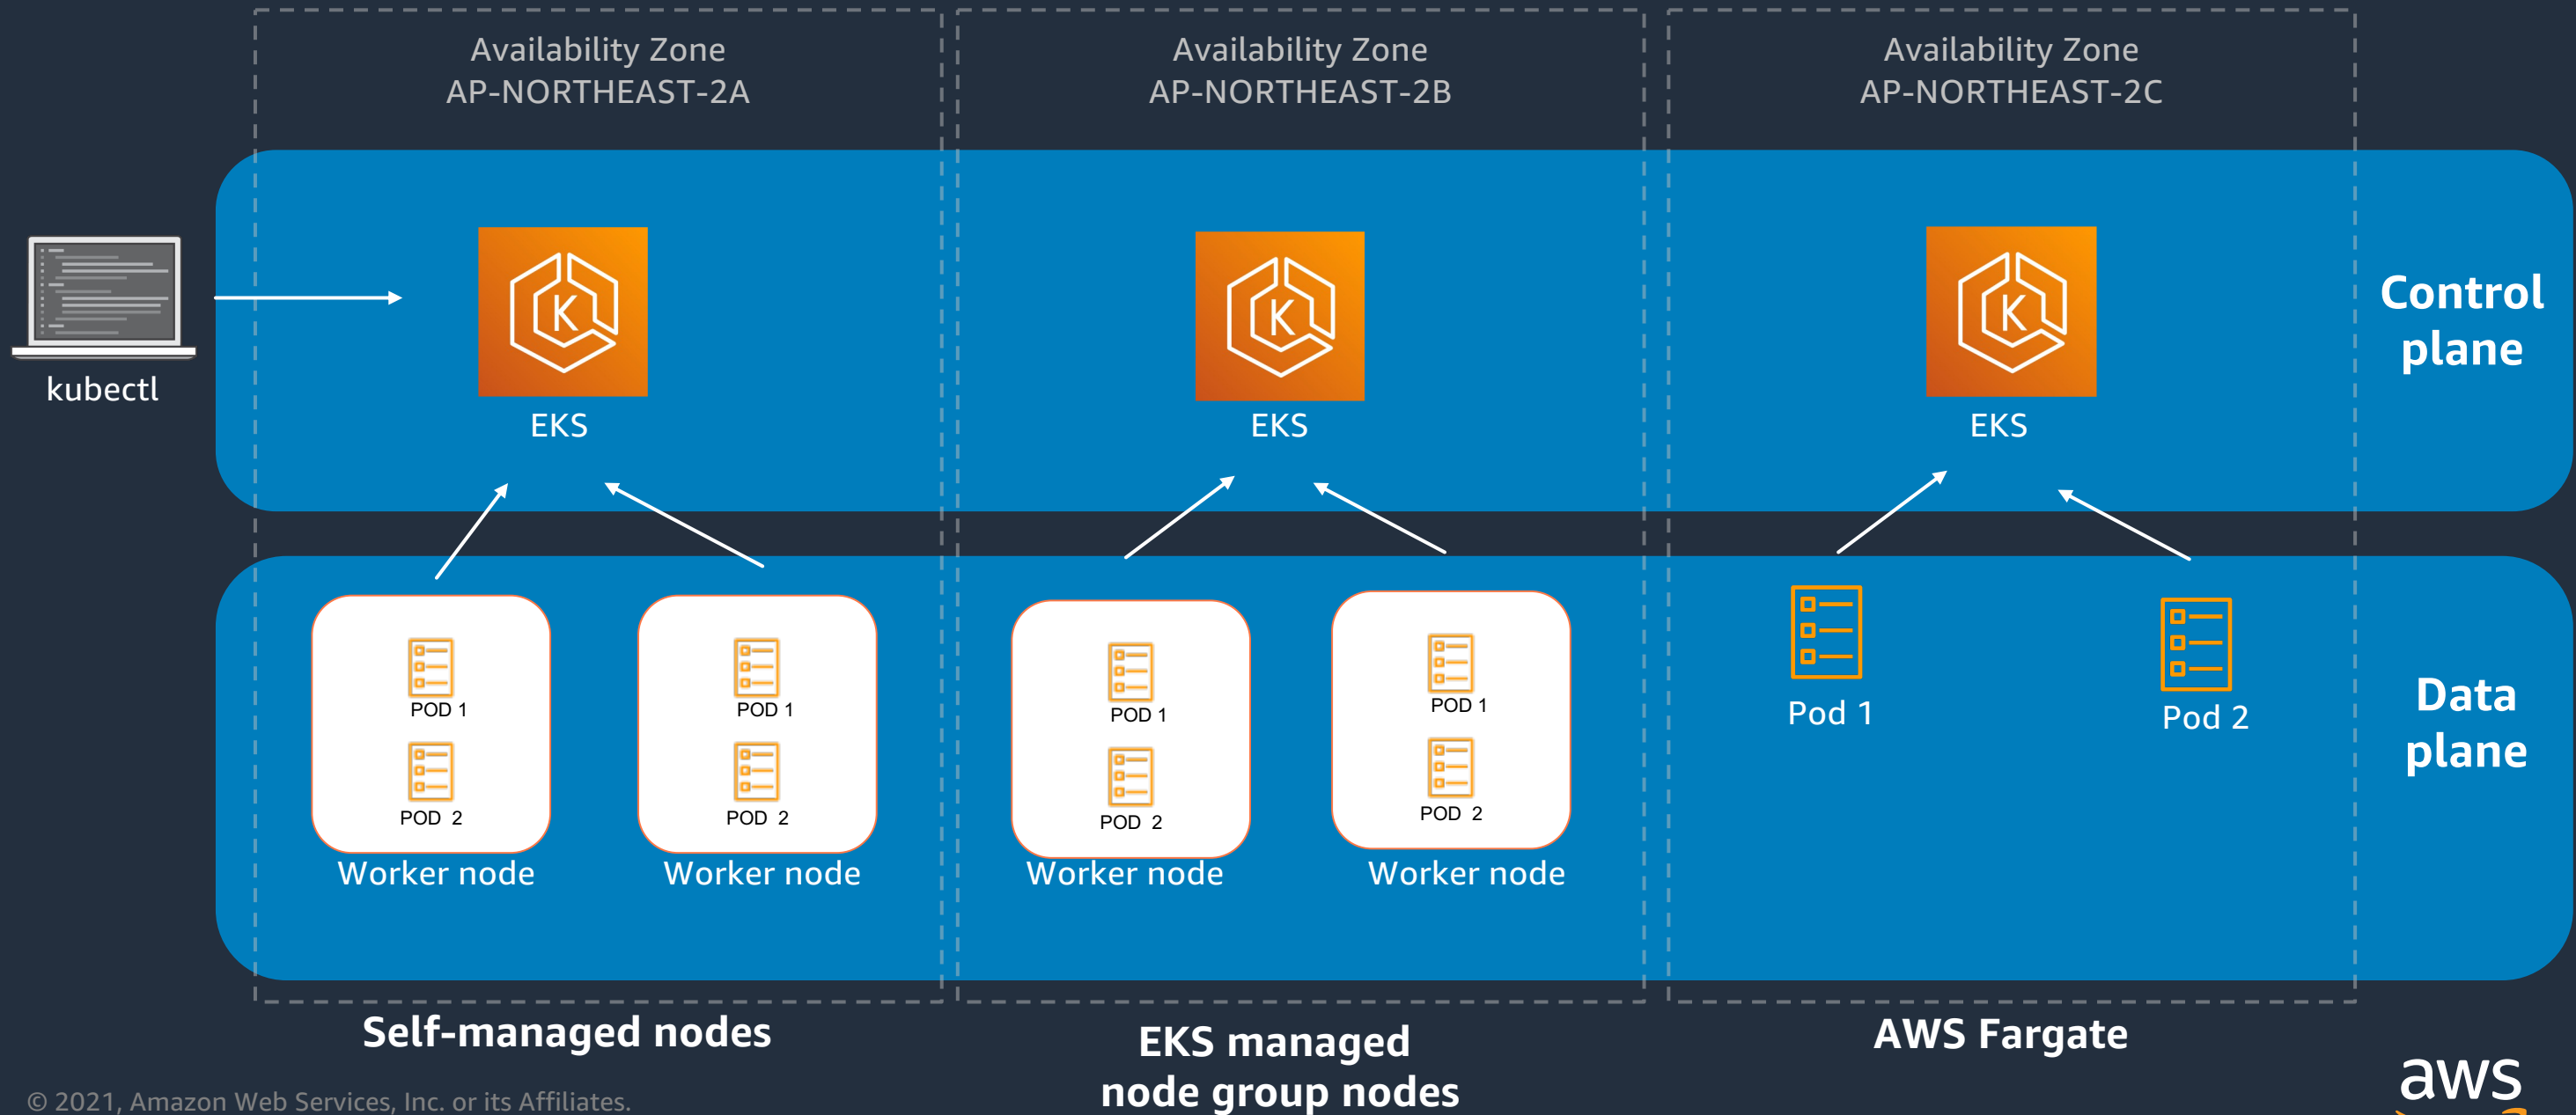
© 2021, Amazon Web Services, Inc. or its Affiliates.



EKS Data Plane Network (CNI)




Amazon EKS



Kubernetes networking model

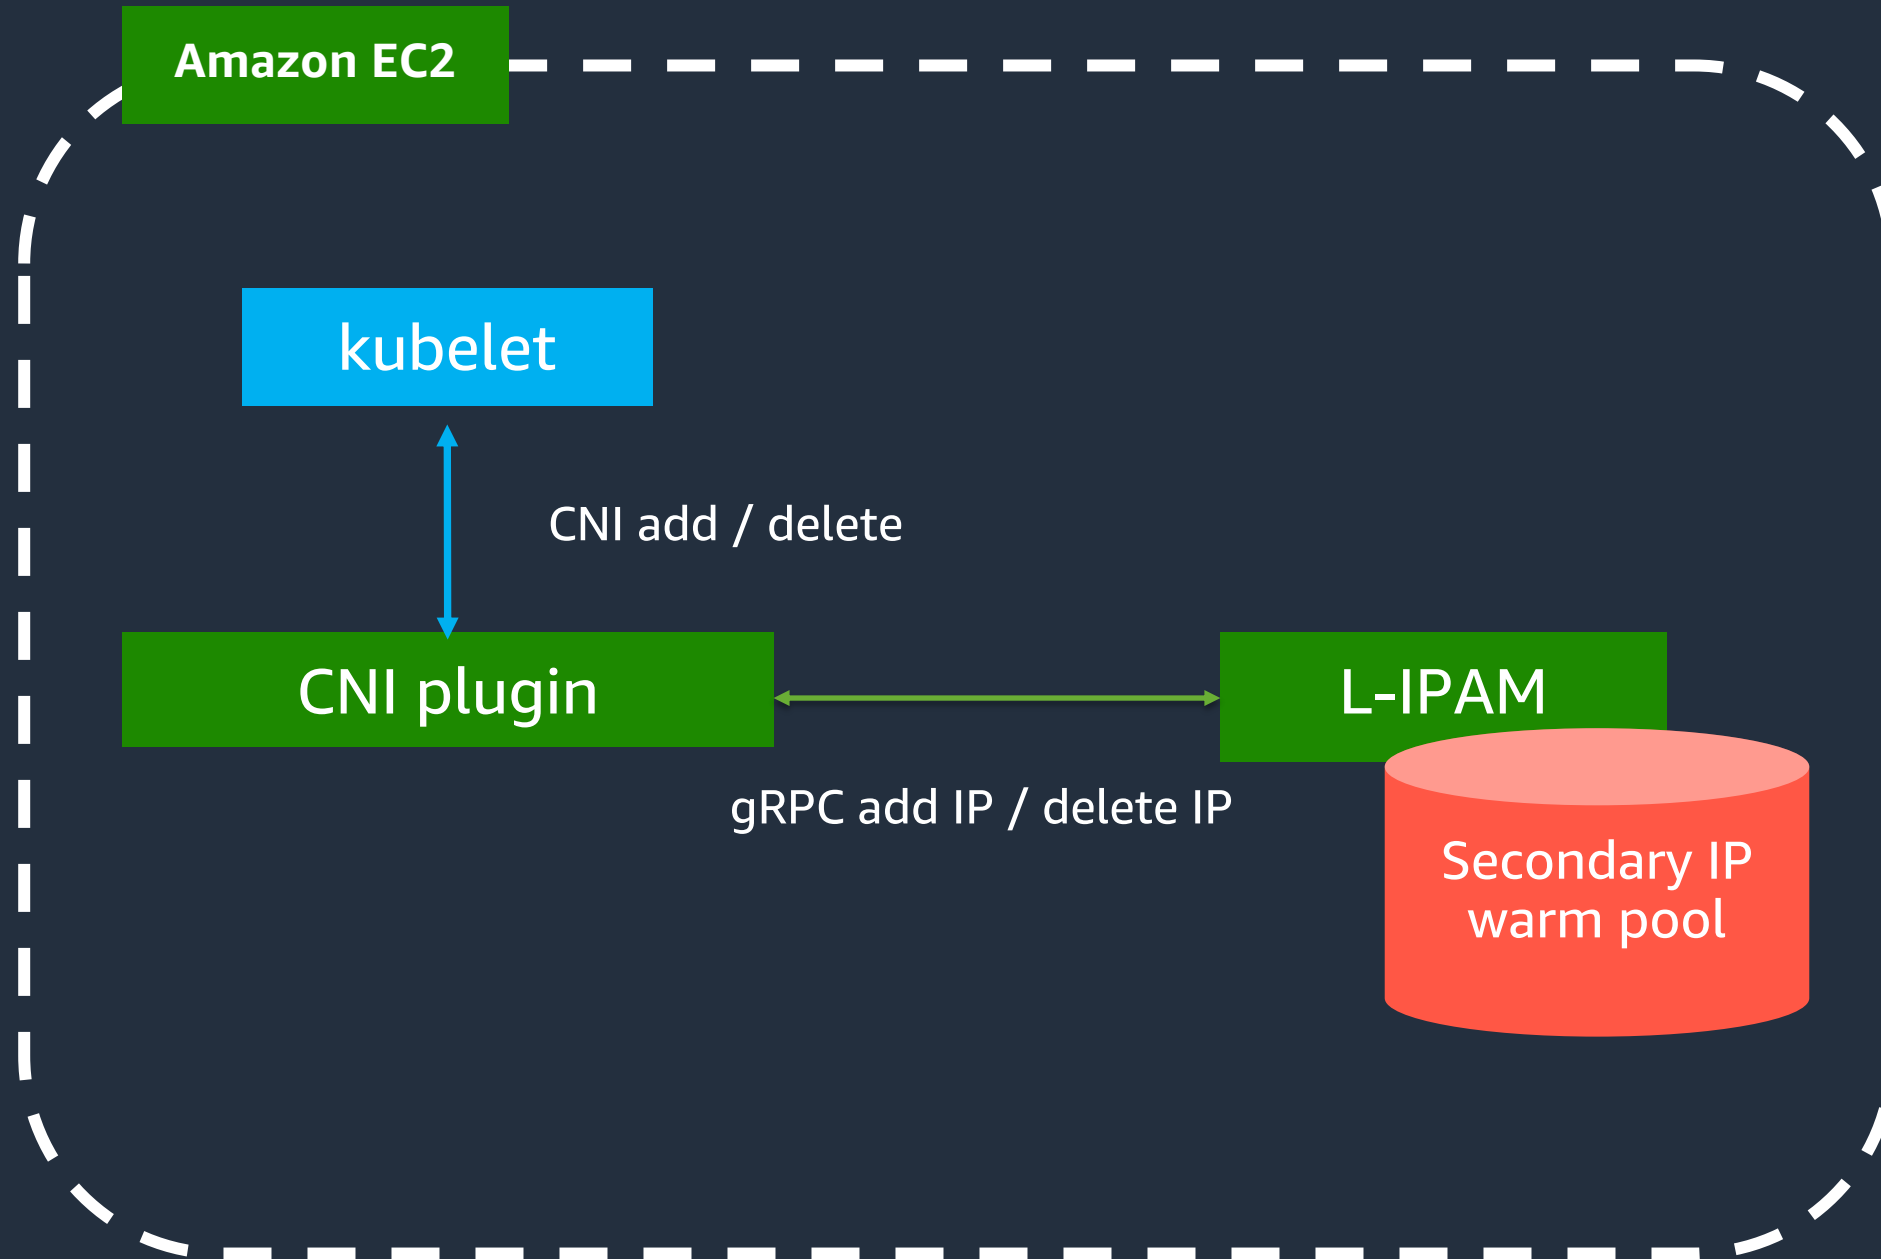
- Network Interface Specification for Kubernetes
- Assign a network namespace and a network interface to the Pod at startup
- Every Pod gets its own IP address
- Containers in the same Pod share the network (IP address)
- Pods communicate to other Pods without NAT
- Clean up the network namespace when the Pod terminates, frees up the IP address



CNI

Container Network
Interface plugin


Amazon VPC CNI




Customizable

WARM_ENI_TARGET
WARM_IP_TARGET
MINIMUM_IP_TARGET
WARM_PREFIX_TARGET

Amazon VPC CNI plugin



Native VPC networking
with CNI plugin



Pods have the same VPC
address inside the pod
as on the VPC



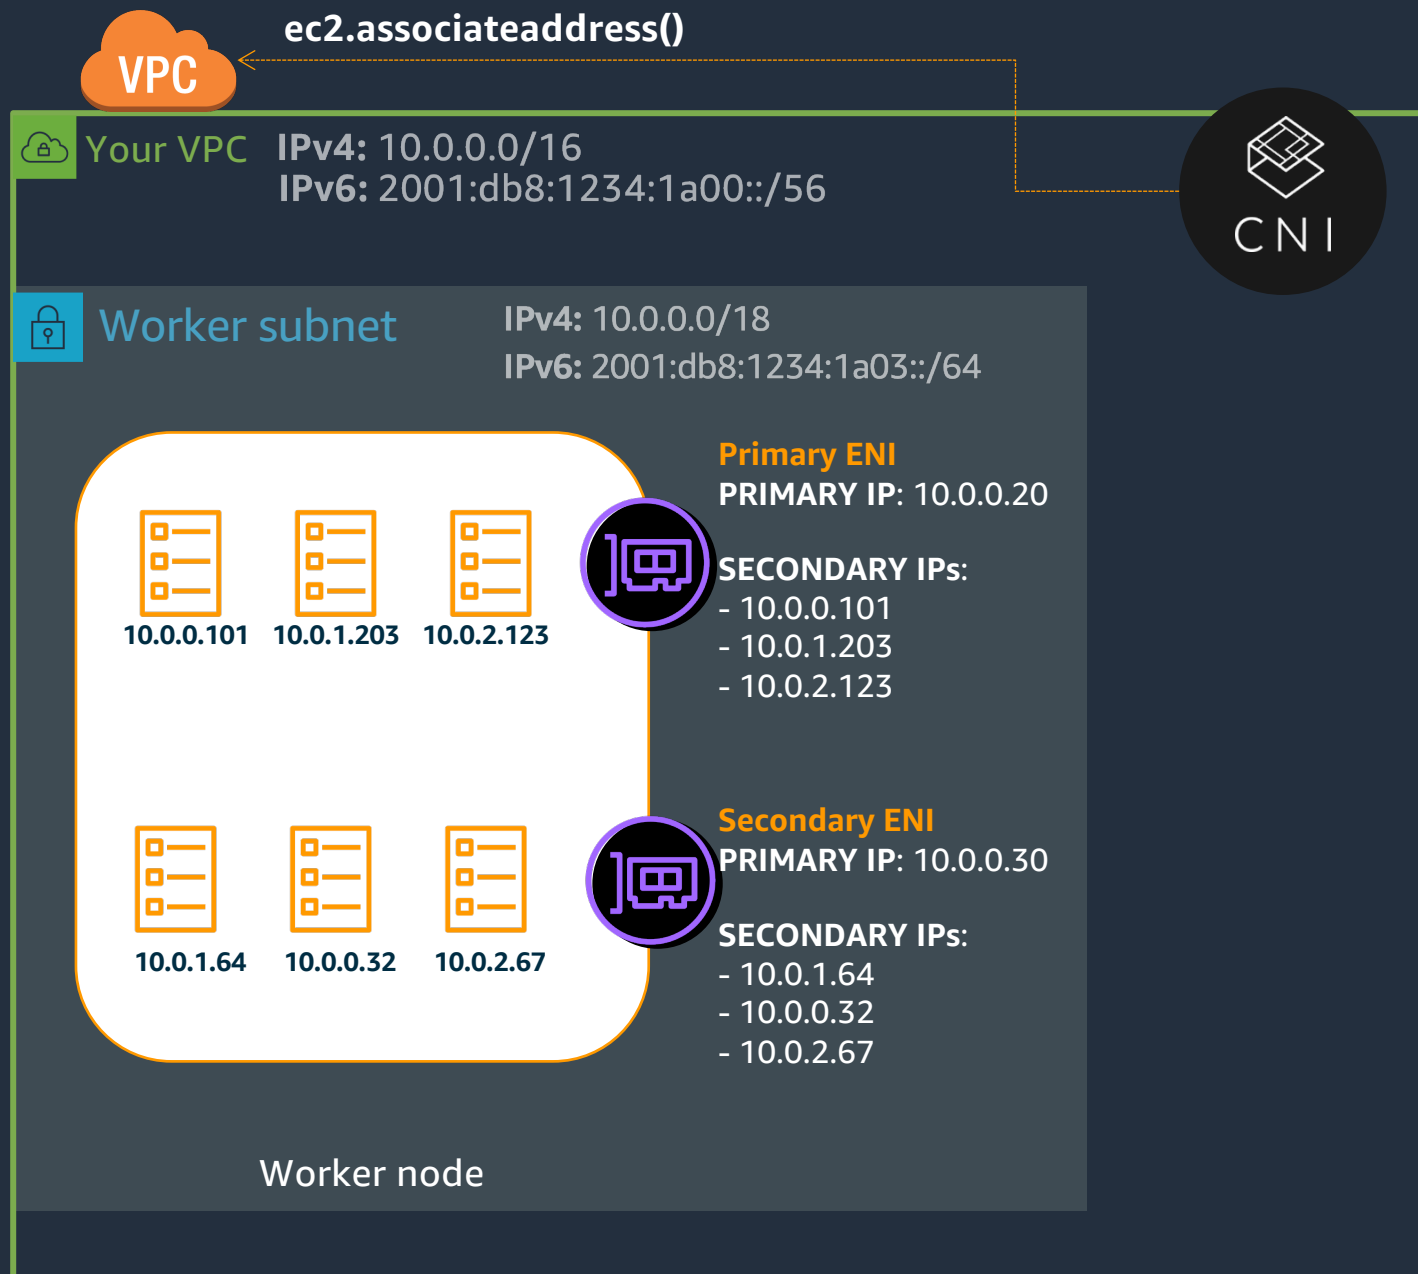
Simple, secure
networking



Open source and
on Github

<https://github.com/aws/amazon-vpc-cni-k8s>

Pod IPv4 networking – CNI plugin



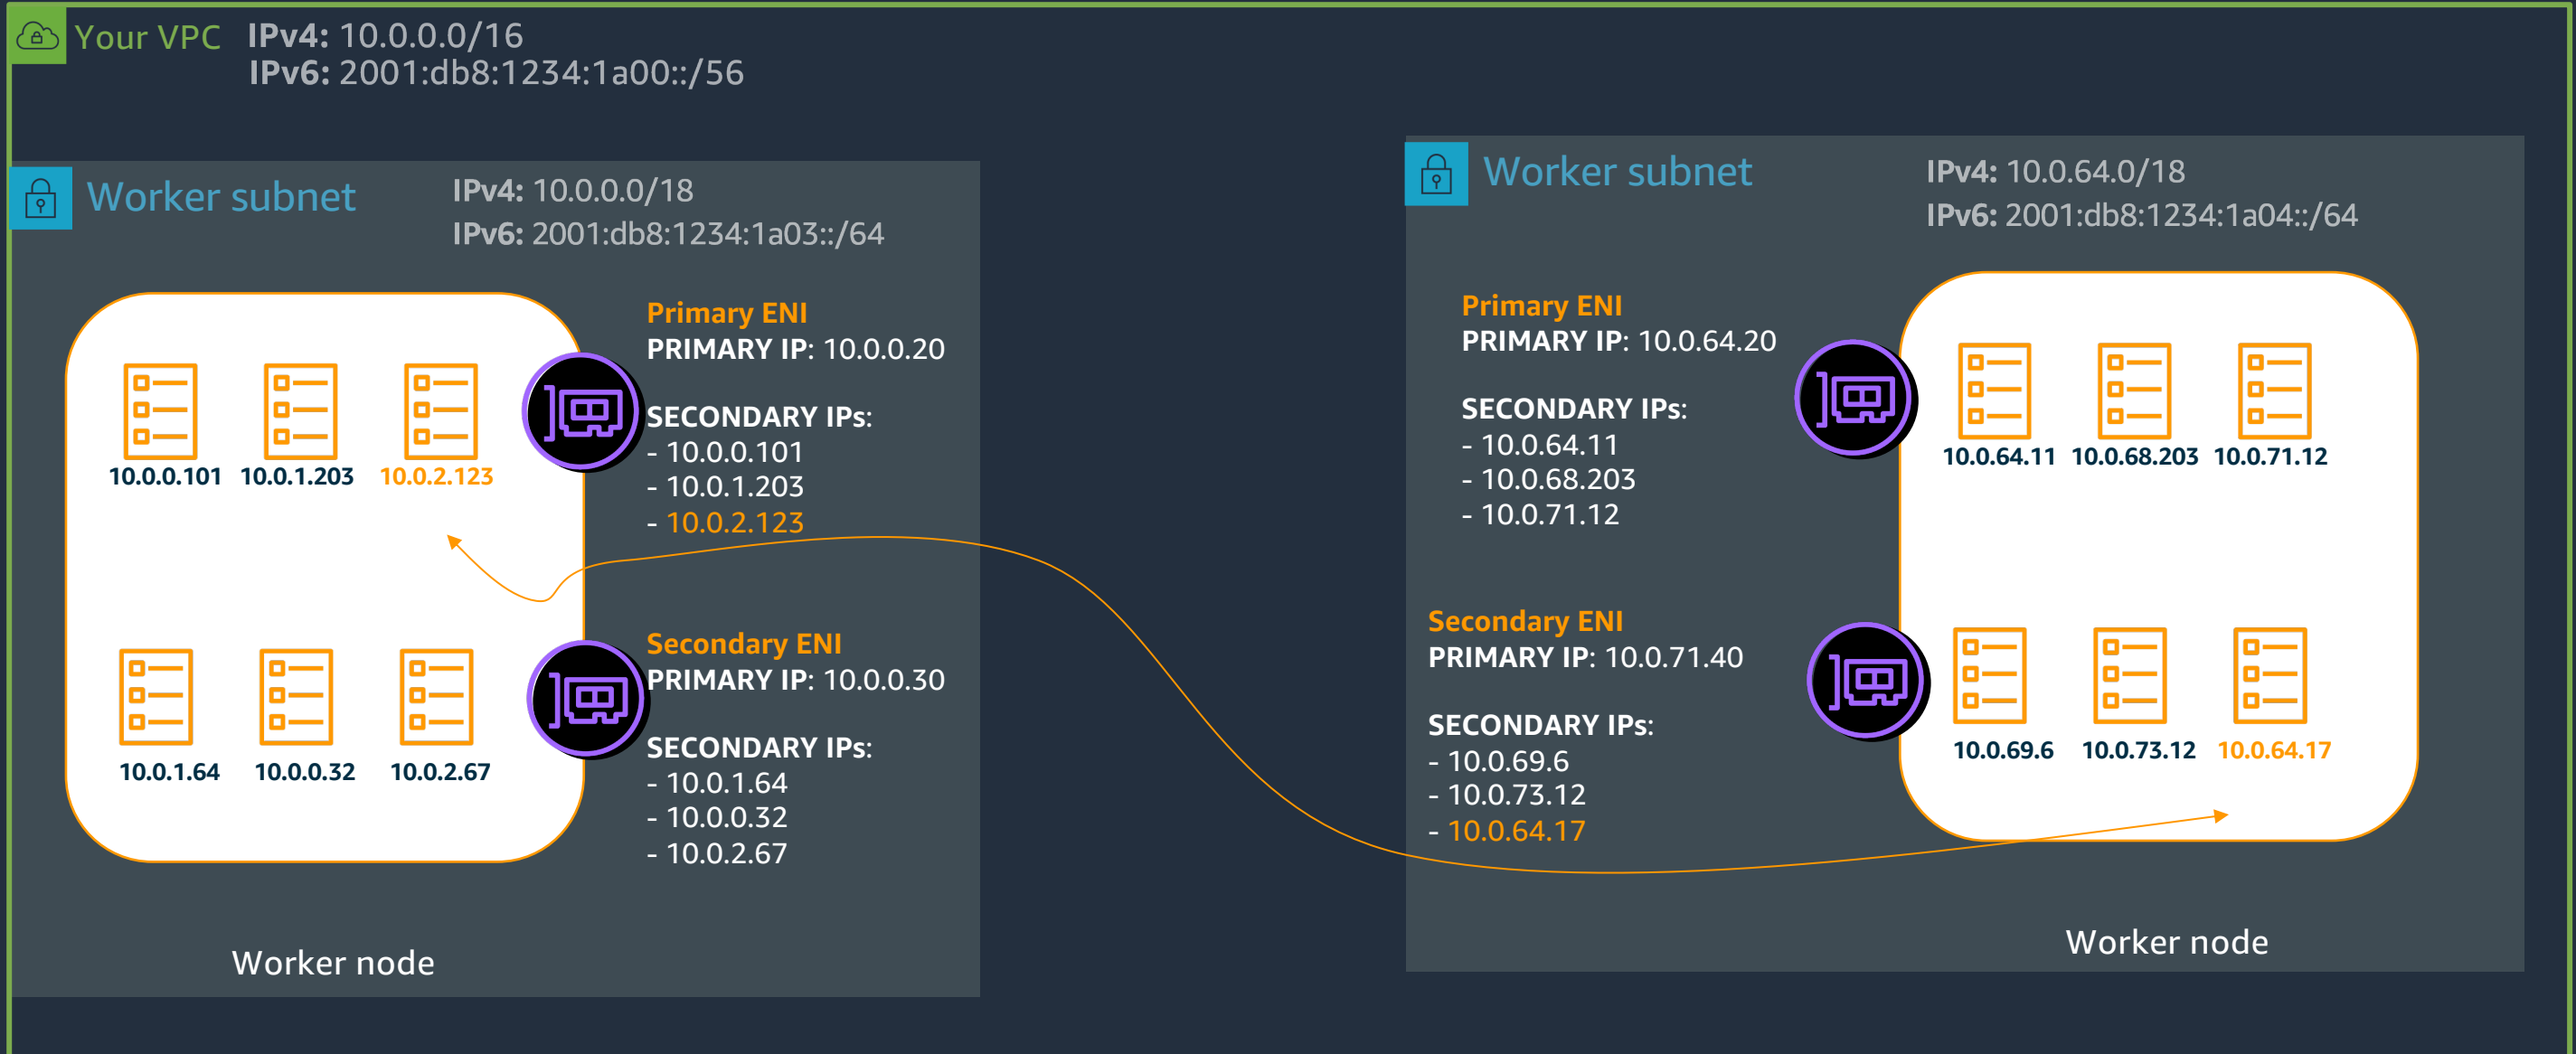
The Amazon VPC Container Network Interface (CNI) plugin is used for

- Creating and attaching ENIs to worker nodes
- Assigning secondary IP addresses for Pods
- Wiring the host network

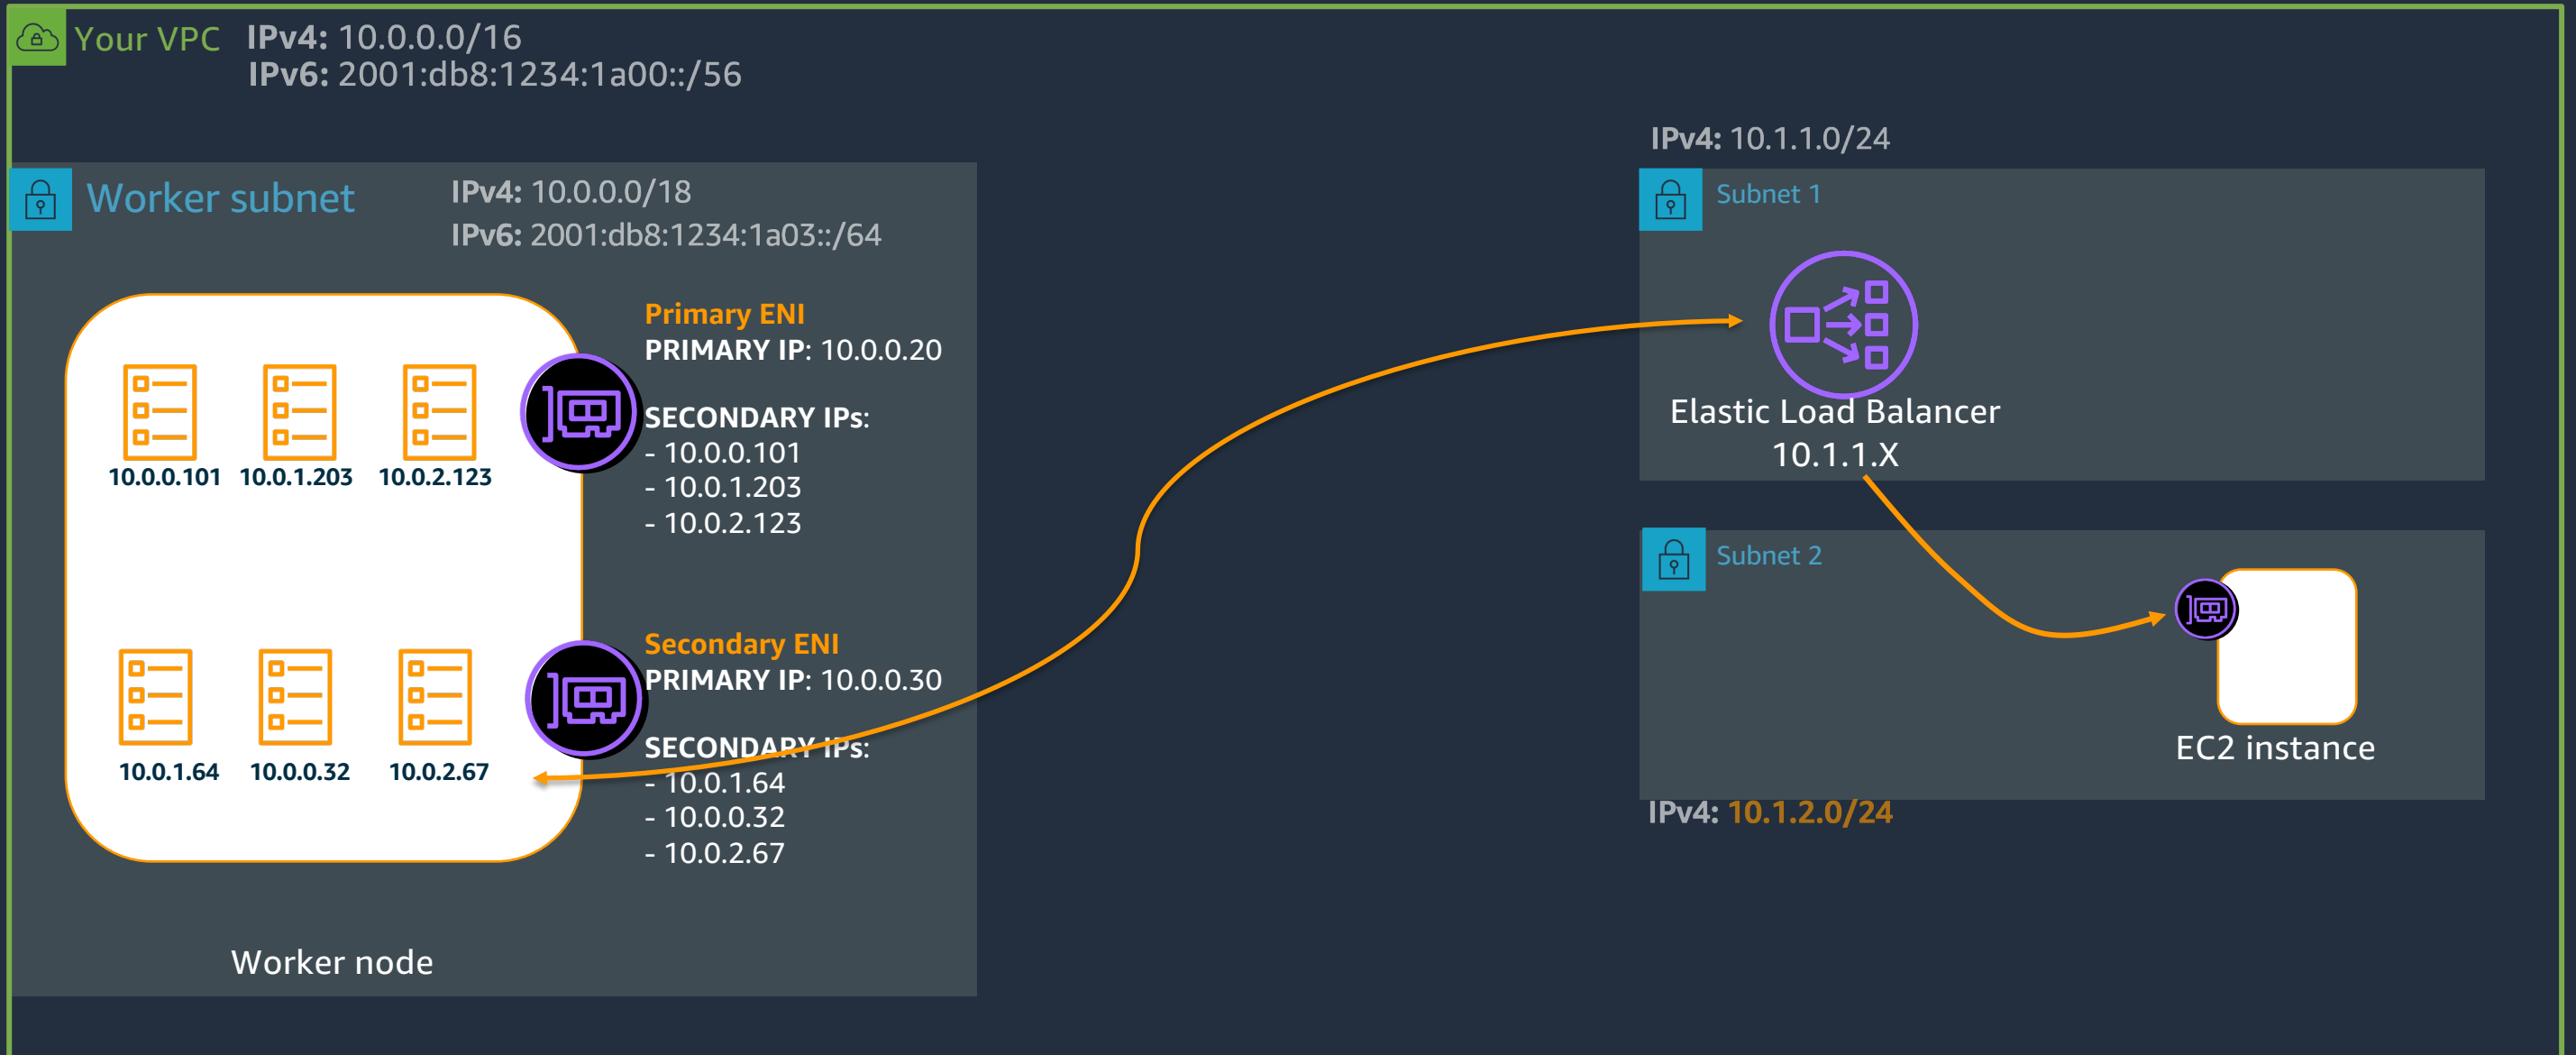
Alternatives to Amazon VPC CNI

- Calico from Tigera
- Cilium from Isovalent
- Weave Net from Weaveworks
- Antrea from VMware

Pod IPv4 networking – Pod to Pod



Pod IPv4 networking – Pod to Pod



Amazon CNI – Configuration options

- Source NAT

```
kubectl set env daemonset aws-node -n kube-system AWS_VPC_K8S_CNI_EXTERNALSNAT=true/false
```

- Prefix delegation

```
kubectl set env daemonset aws-node -n kube-system ENABLE_PREFIX_DELEGATION=true/false
```

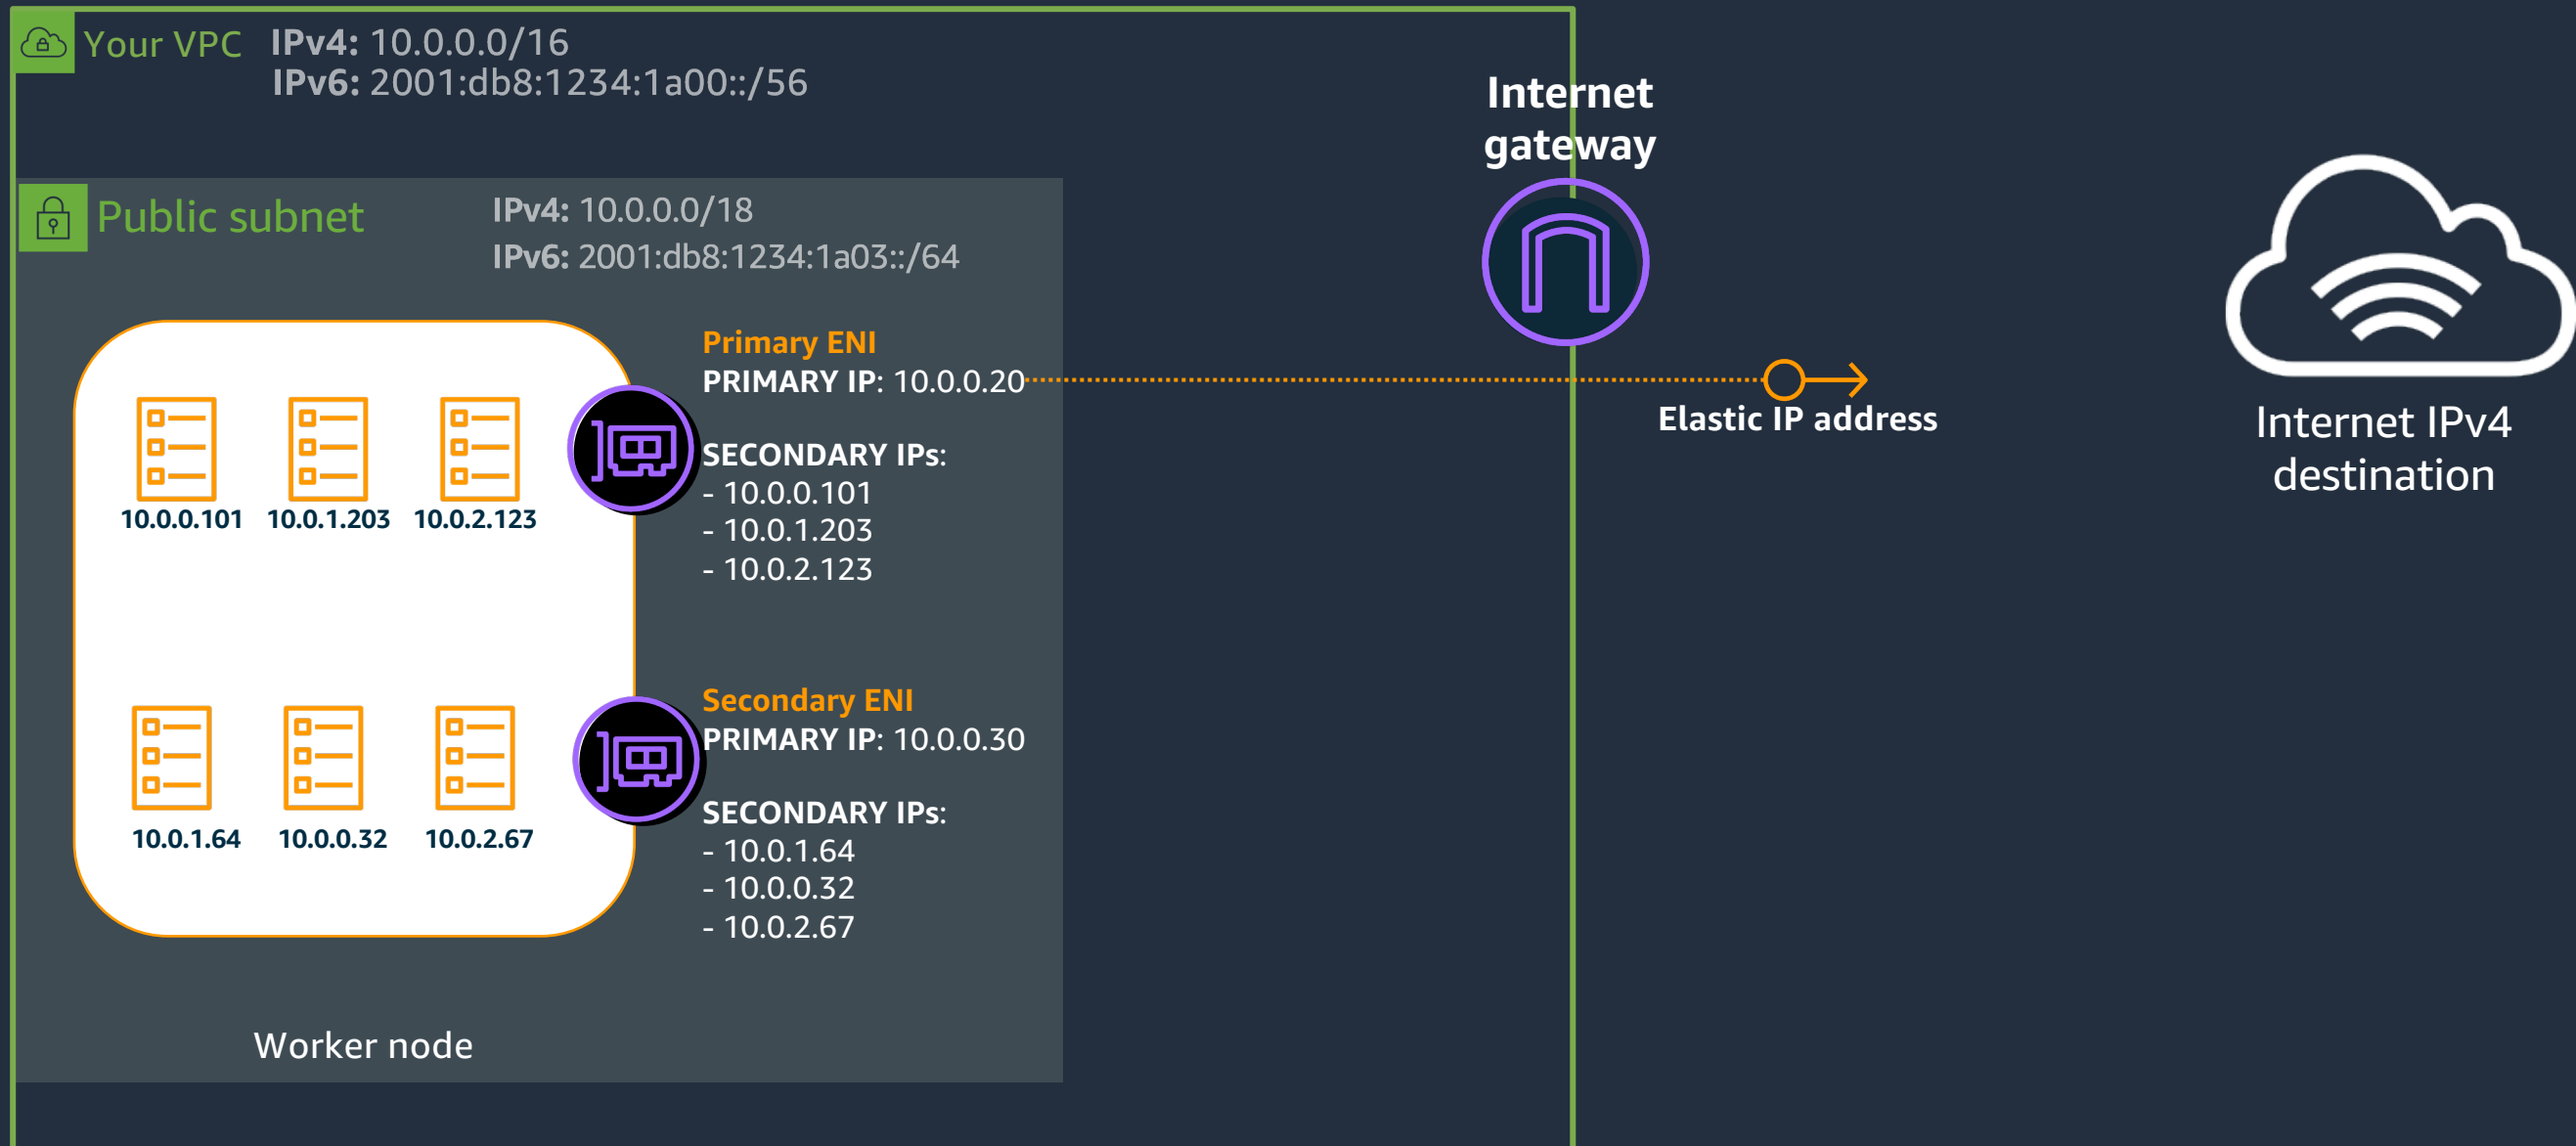
- Security groups for Pods

```
kubectl set env daemonset aws-node -n kube-system ENABLE_POD_ENI=true/false
```

- Custom networking

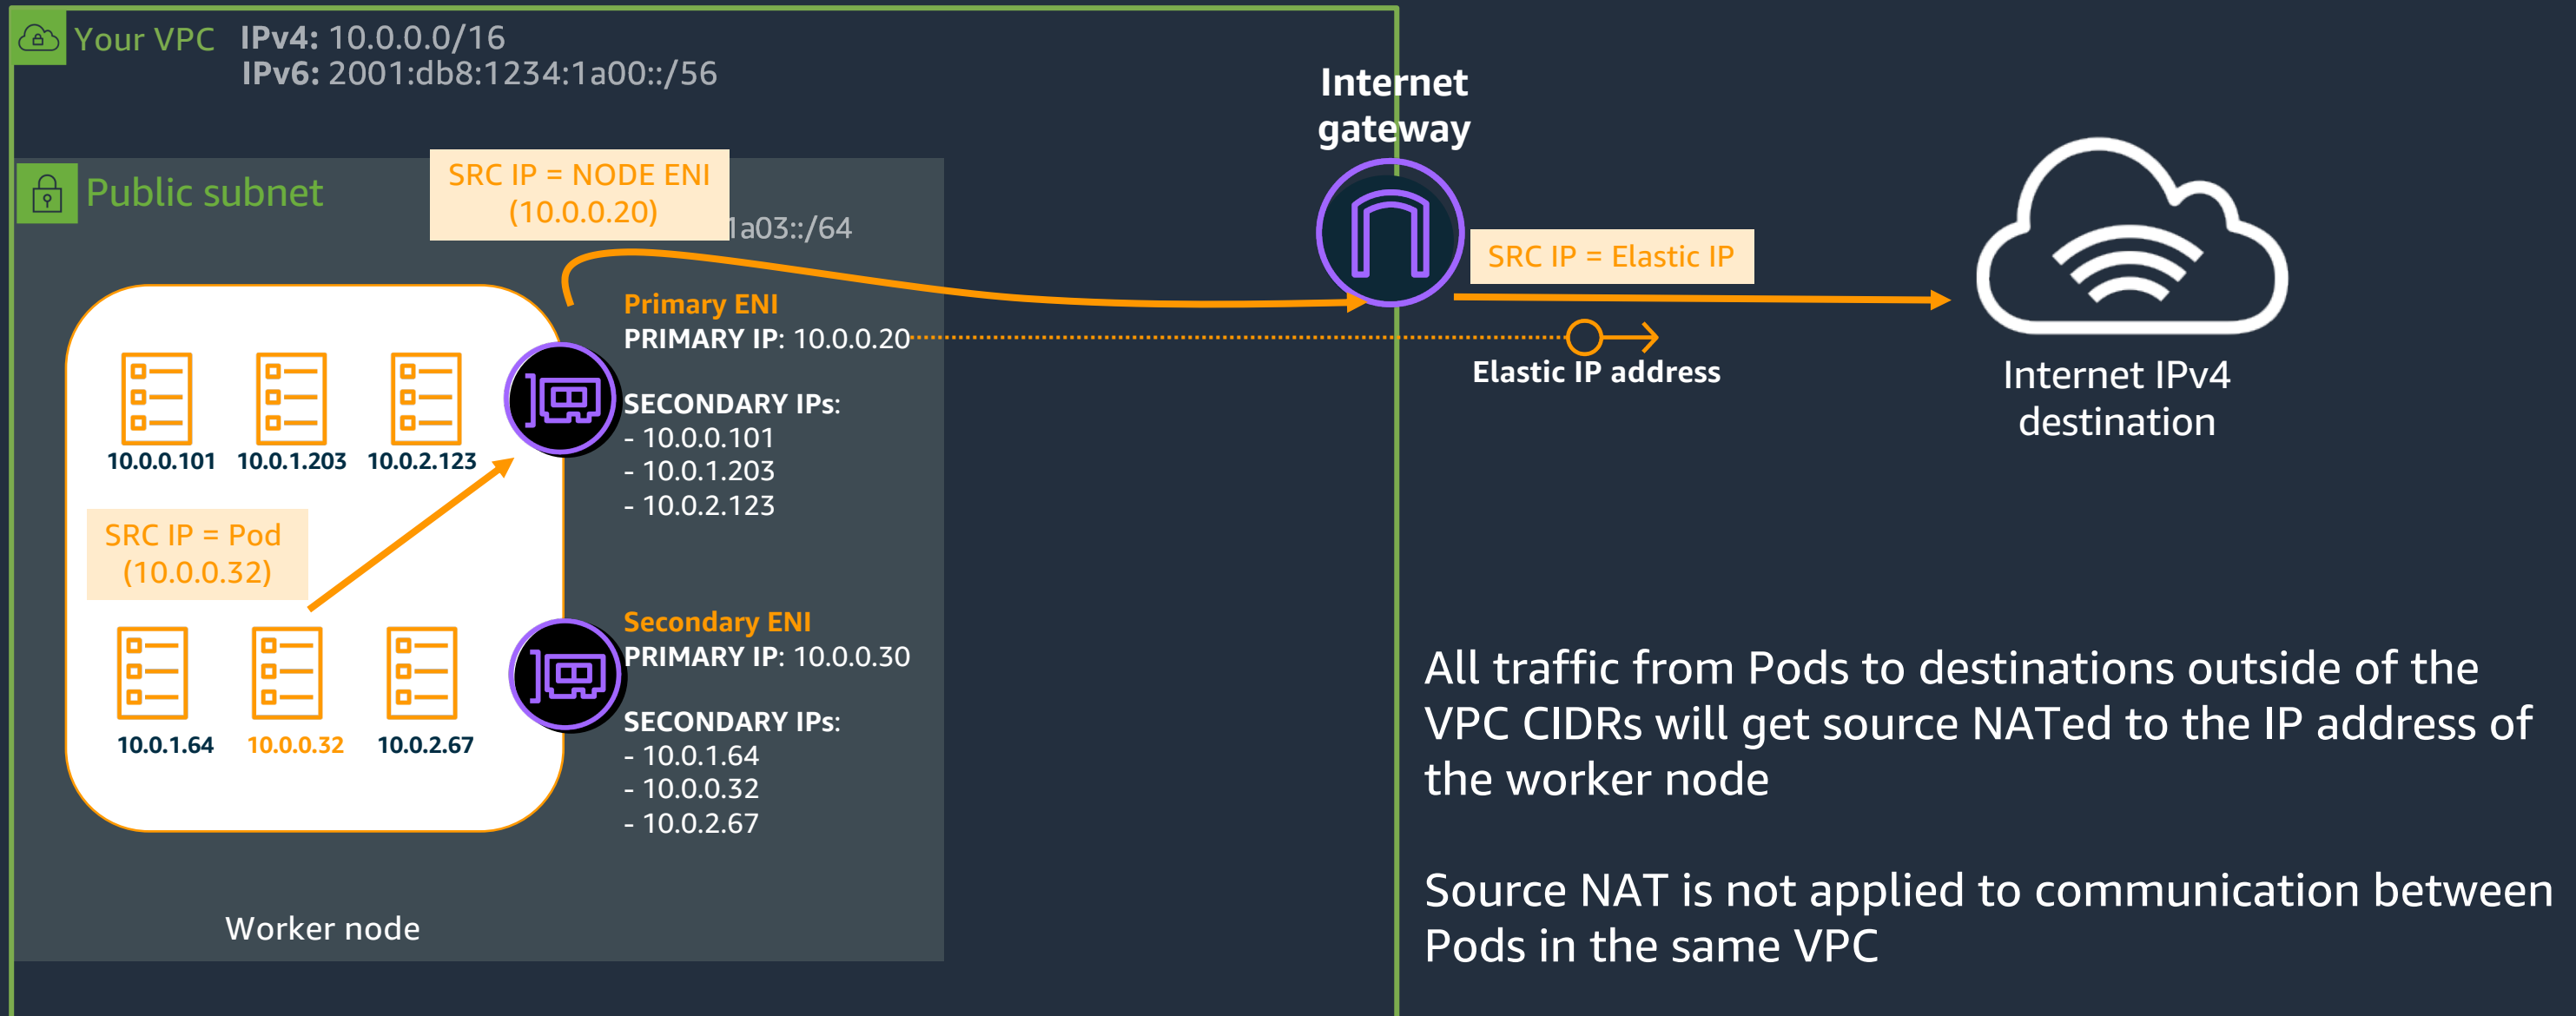
```
kubectl set env daemonset aws-node -n kube-system AWS_VPC_K8S_CNI_CUSTOM_NETWORK_CFG=true/false
```

Pod IPv4 networking – Source NAT enabled



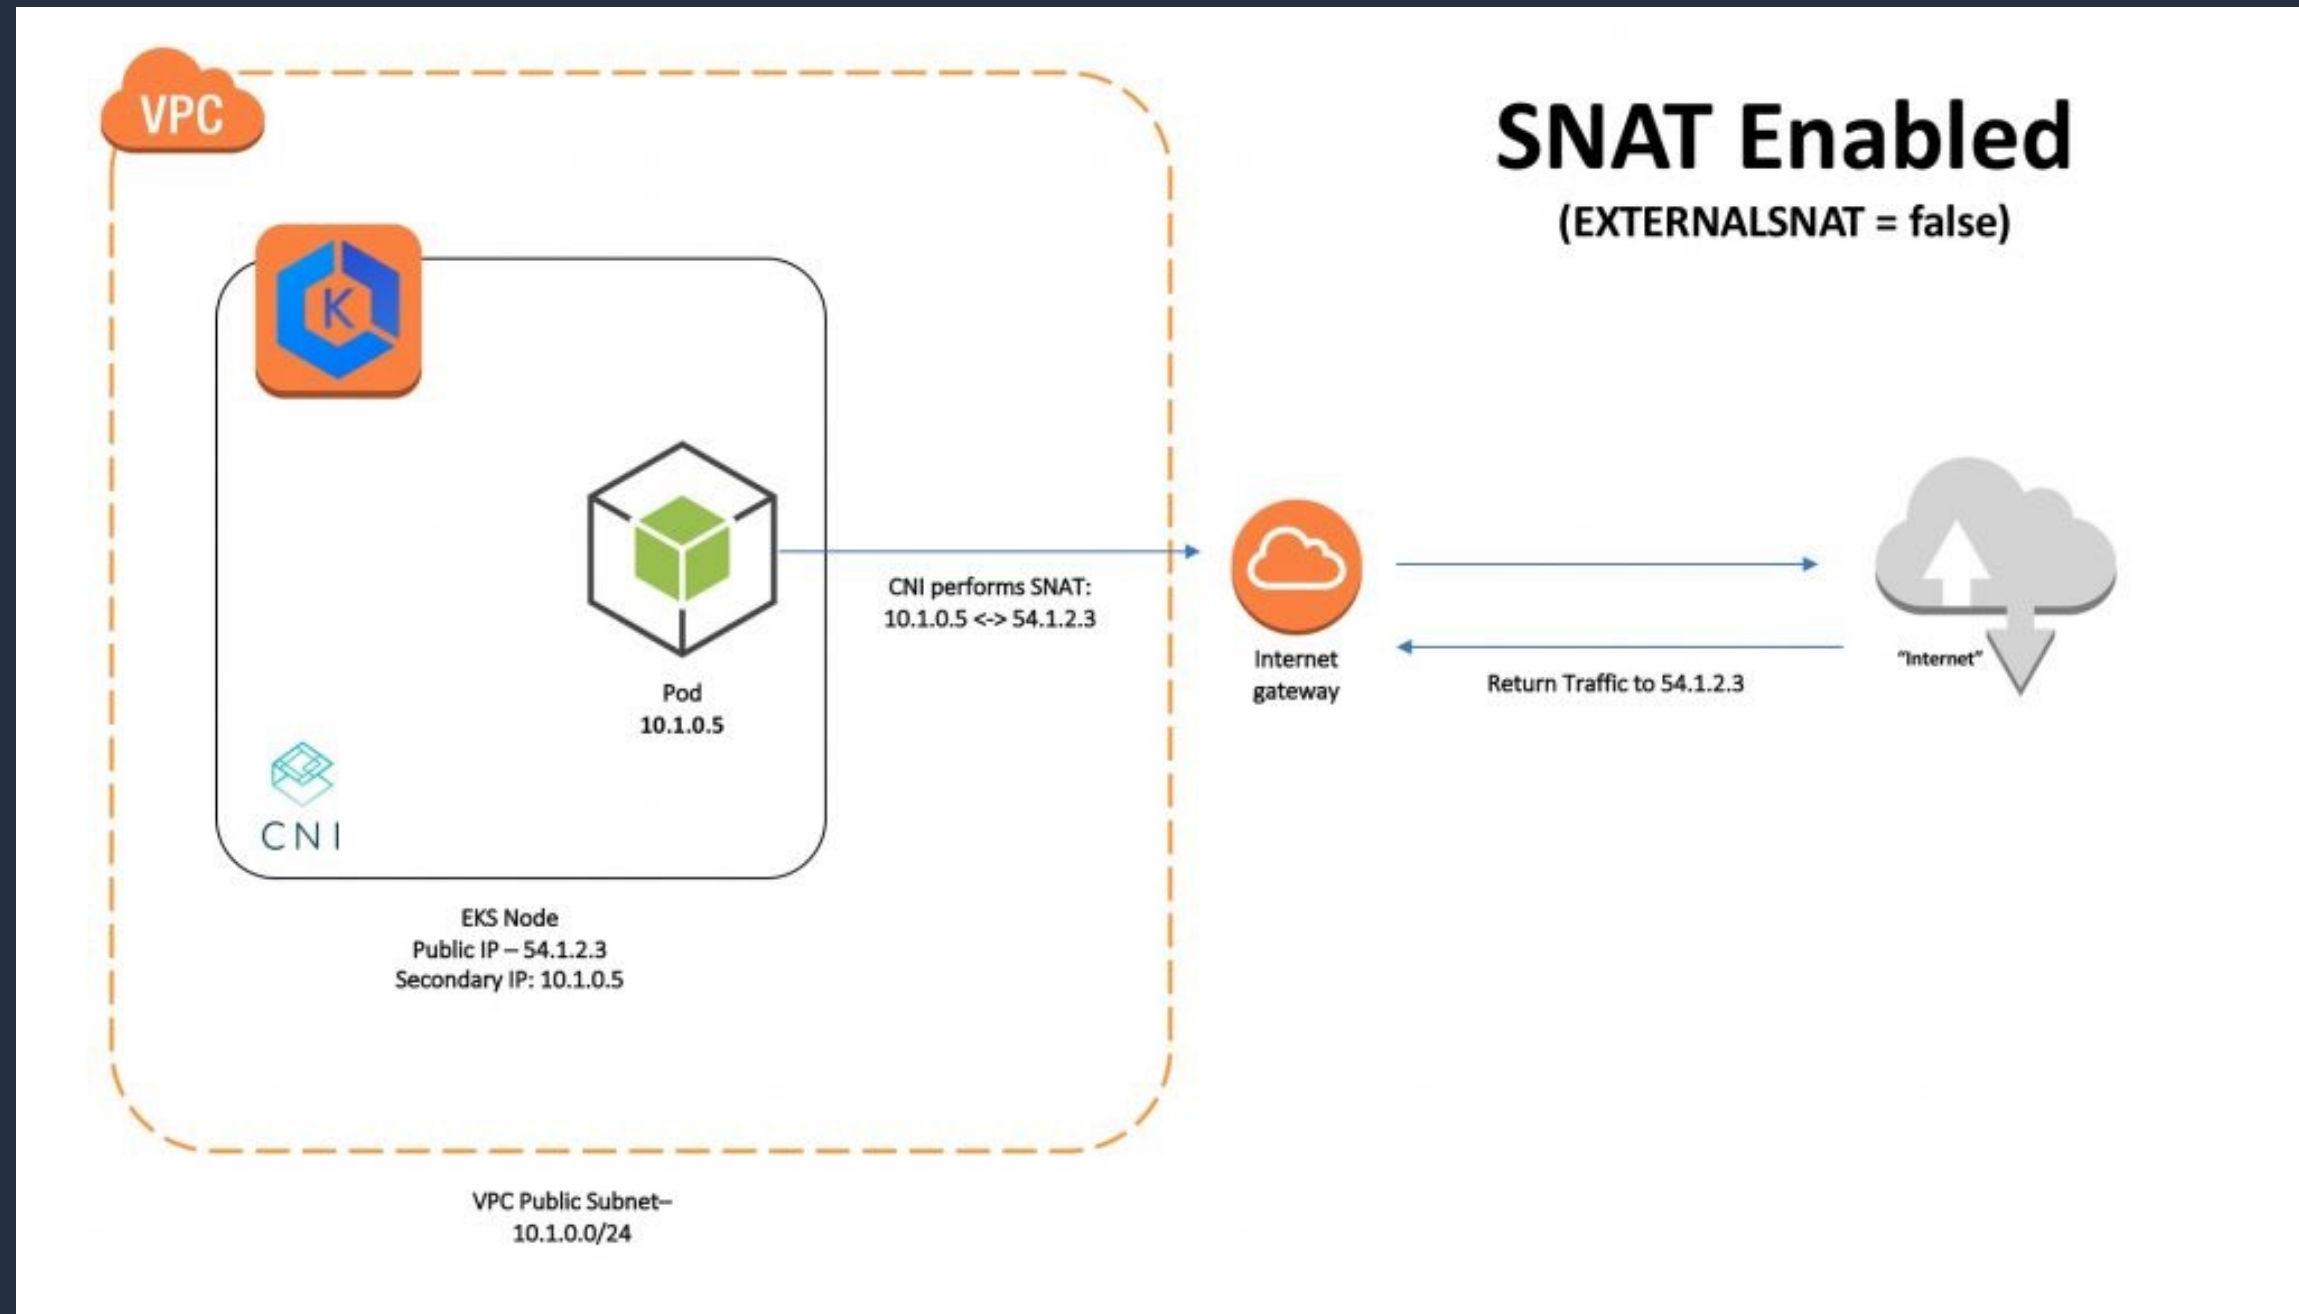
```
kubectl set env daemonset aws-node -n kube-system AWS_VPC_K8S_CNI_EXTERNALSNAT=false
```

Pod IPv4 networking – Source NAT enabled

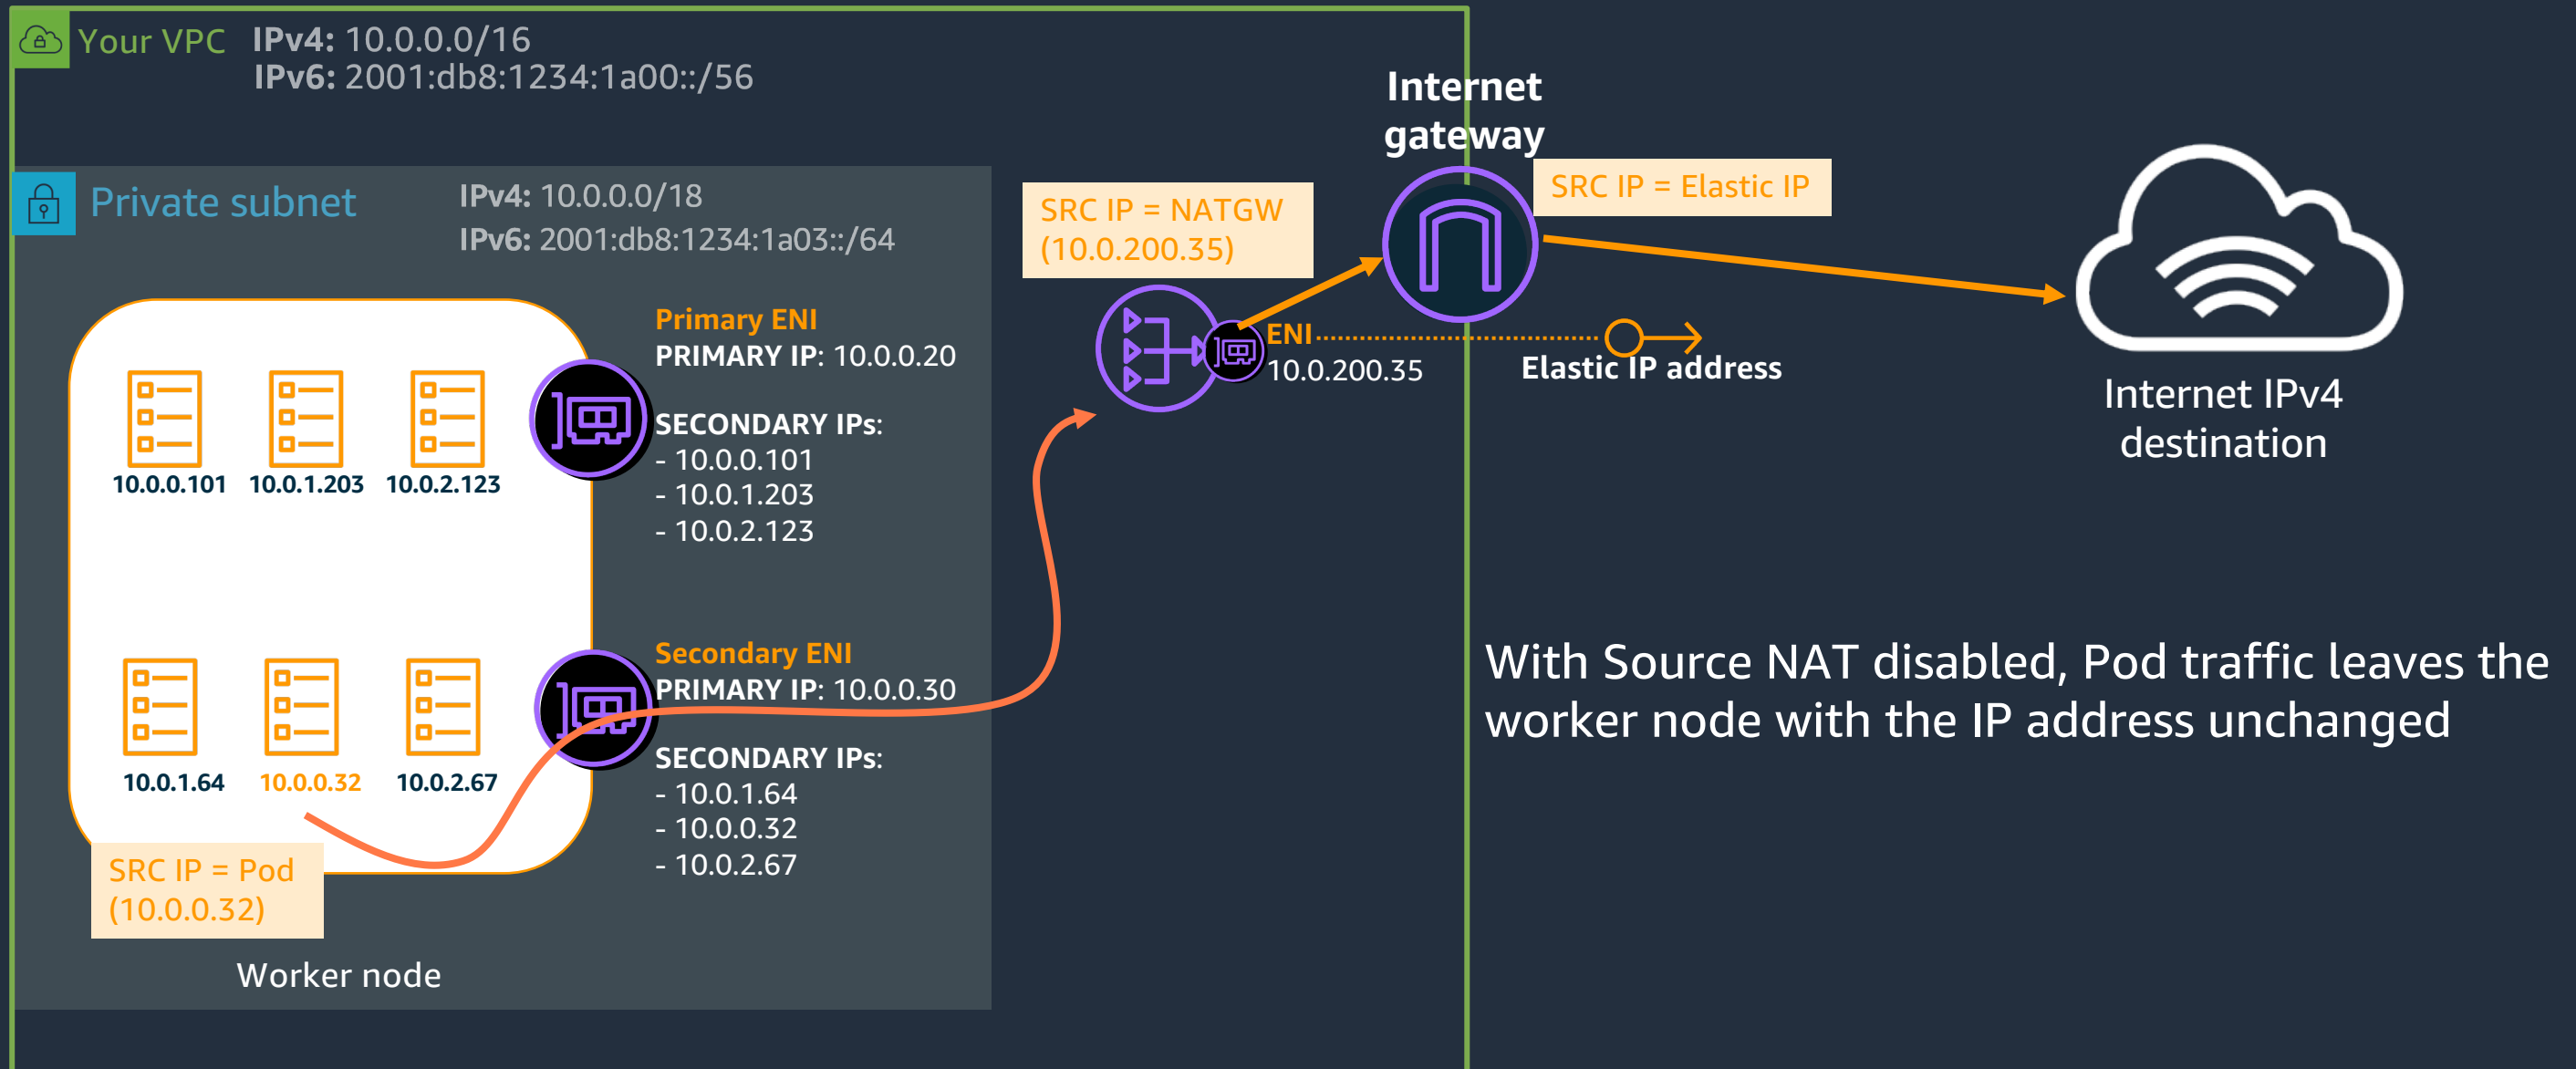


```
kubectl set env daemonset aws-node -n kube-system AWS_VPC_K8S_CNI_EXTERNALSNAT=false
```


Pod IPv4 networking – Source NAT enabled

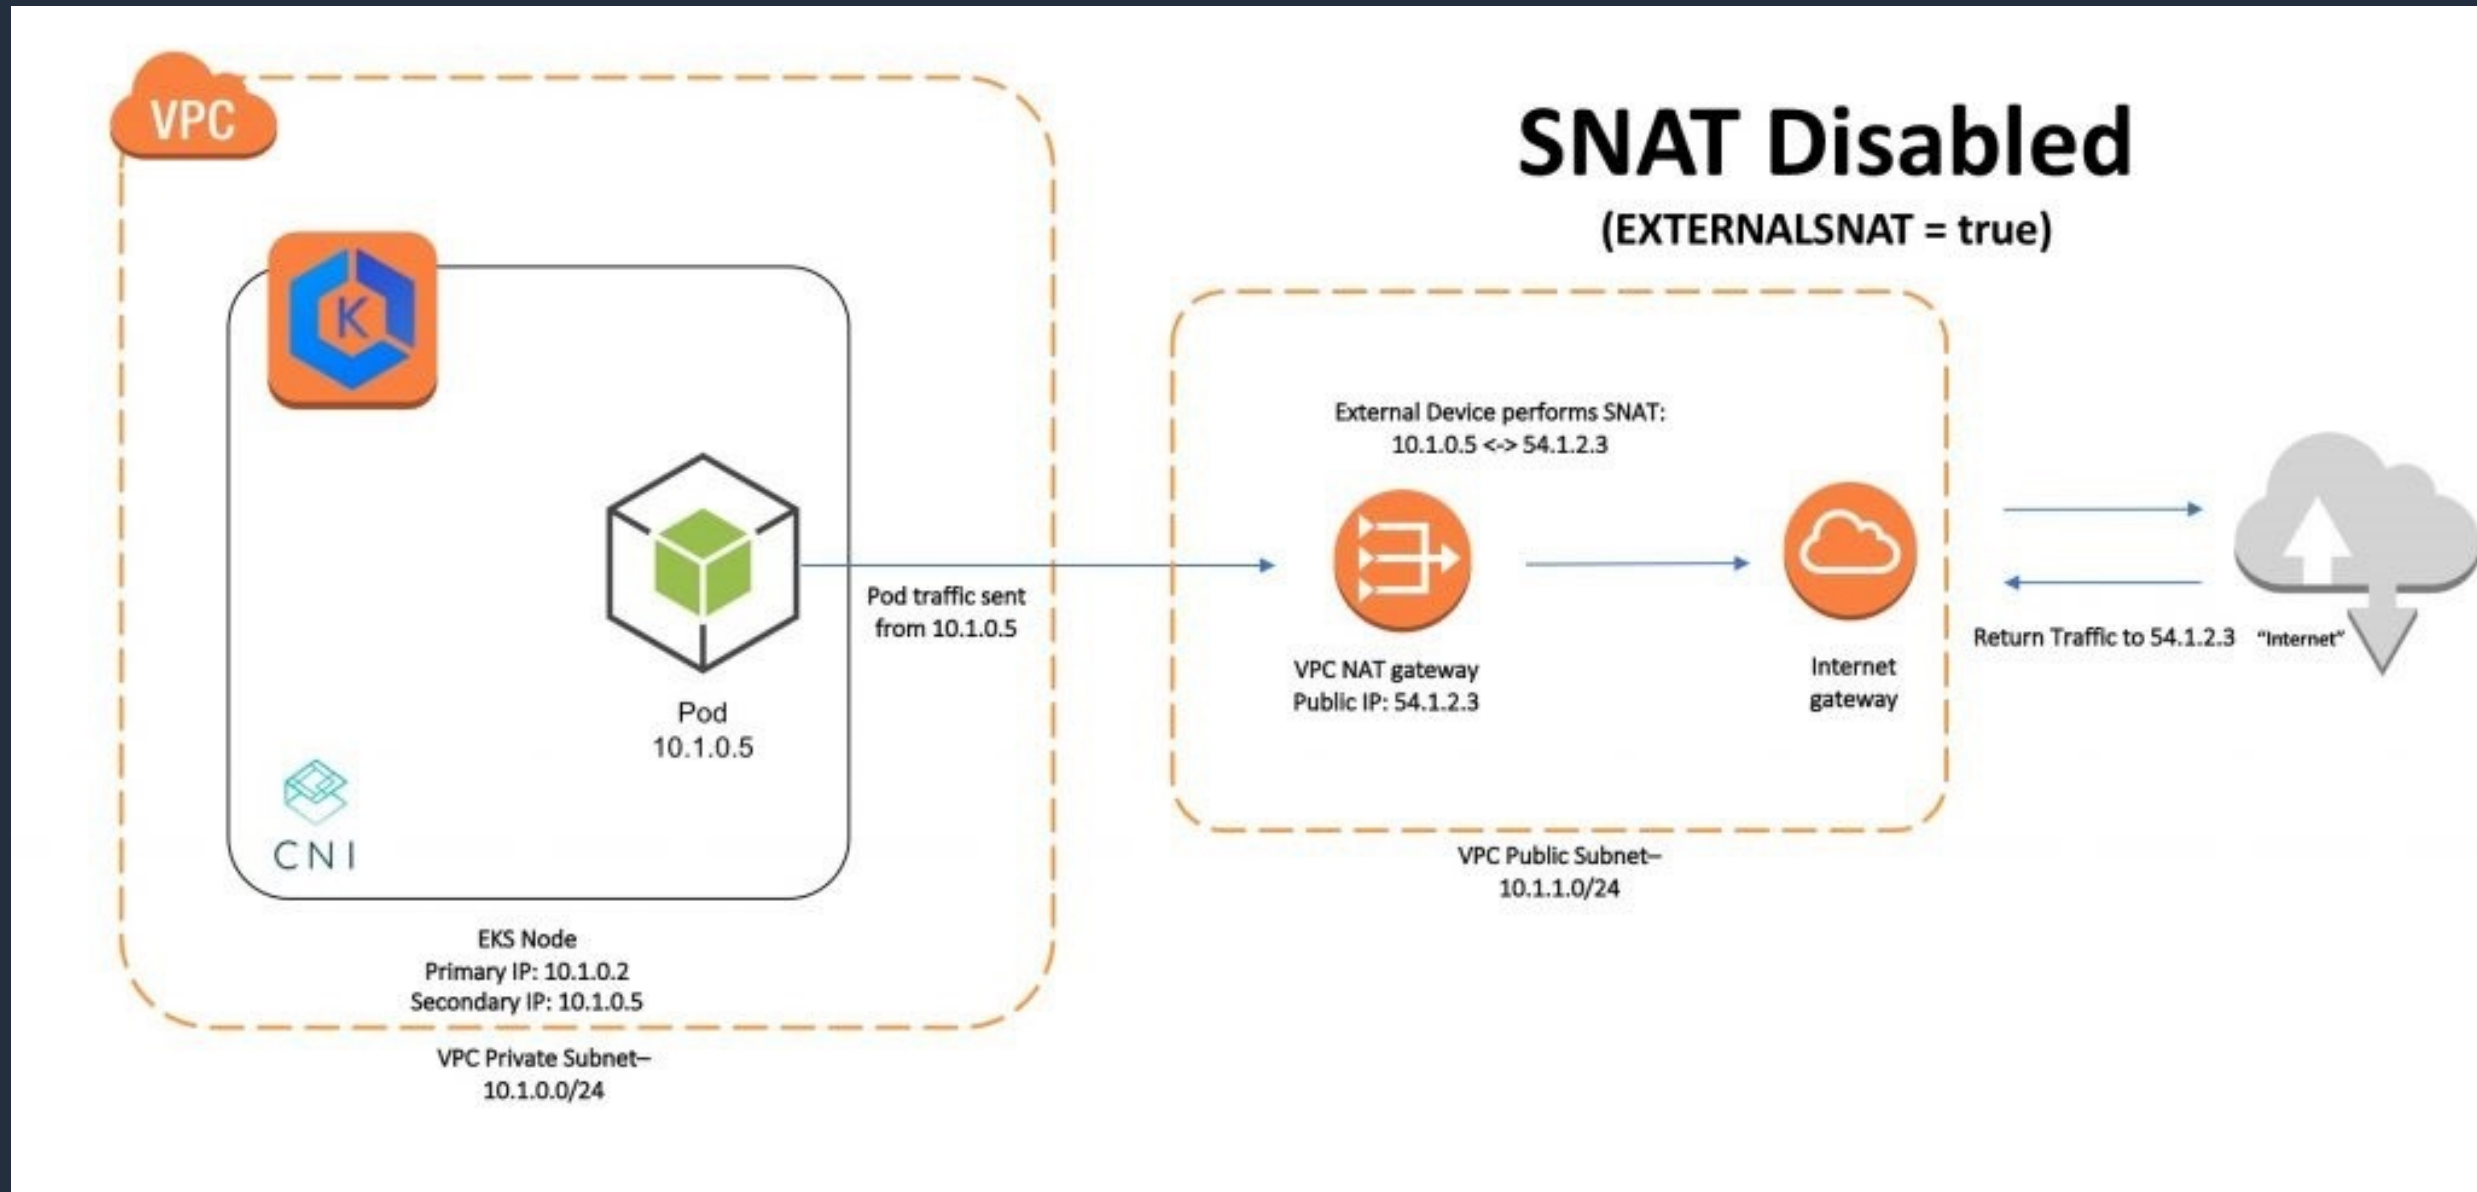


Pod networking – Source NAT disabled



```
kubectl set env daemonset aws-node -n kube-system AWS_VPC_K8S_CNI_EXTERNALSNAT=true
```

Pod networking – Source NAT disabled






Secondary IPs for Pods per EC2 instance

1x ENI IP
9x secondary IPs

1x ENI IP
9x secondary IPs

1x ENI IP
9x secondary IPs



m5.large

Instance type	Maximum network interfaces	Private IPv4 addresses per interface
m5.large	3	10

MAX PODS = (Number of network interfaces × [the number of IP addresses per network interface – 1]) + 2




?

Secondary IPs for Pods per EC2 instance

1x ENI IP
9x secondary IPs

1x ENI IP
9x secondary IPs

1x ENI IP
9x secondary IPs



m5.large

Instance type	Maximum network interfaces	Private IPv4 addresses per interface
m5.large	3	10

<https://github.com/awslabs/amazon-eks-ami/blob/master/files/eni-max-pods.txt>

MAX PODS = (Number of network interfaces × [the number of IP addresses per network interface – 1]) + 2

29

Amazon VPC prefix delegation

Allows for assigning a prefix to an EC2 ENI

- **/28** block for IPv4 (16x IPv4 addresses)
- **/80** block for IPv6 (280 **trillion** IPv6 addresses)

* Prefix delegation is only supported on Nitro instances

Prefix delegation



m5.large

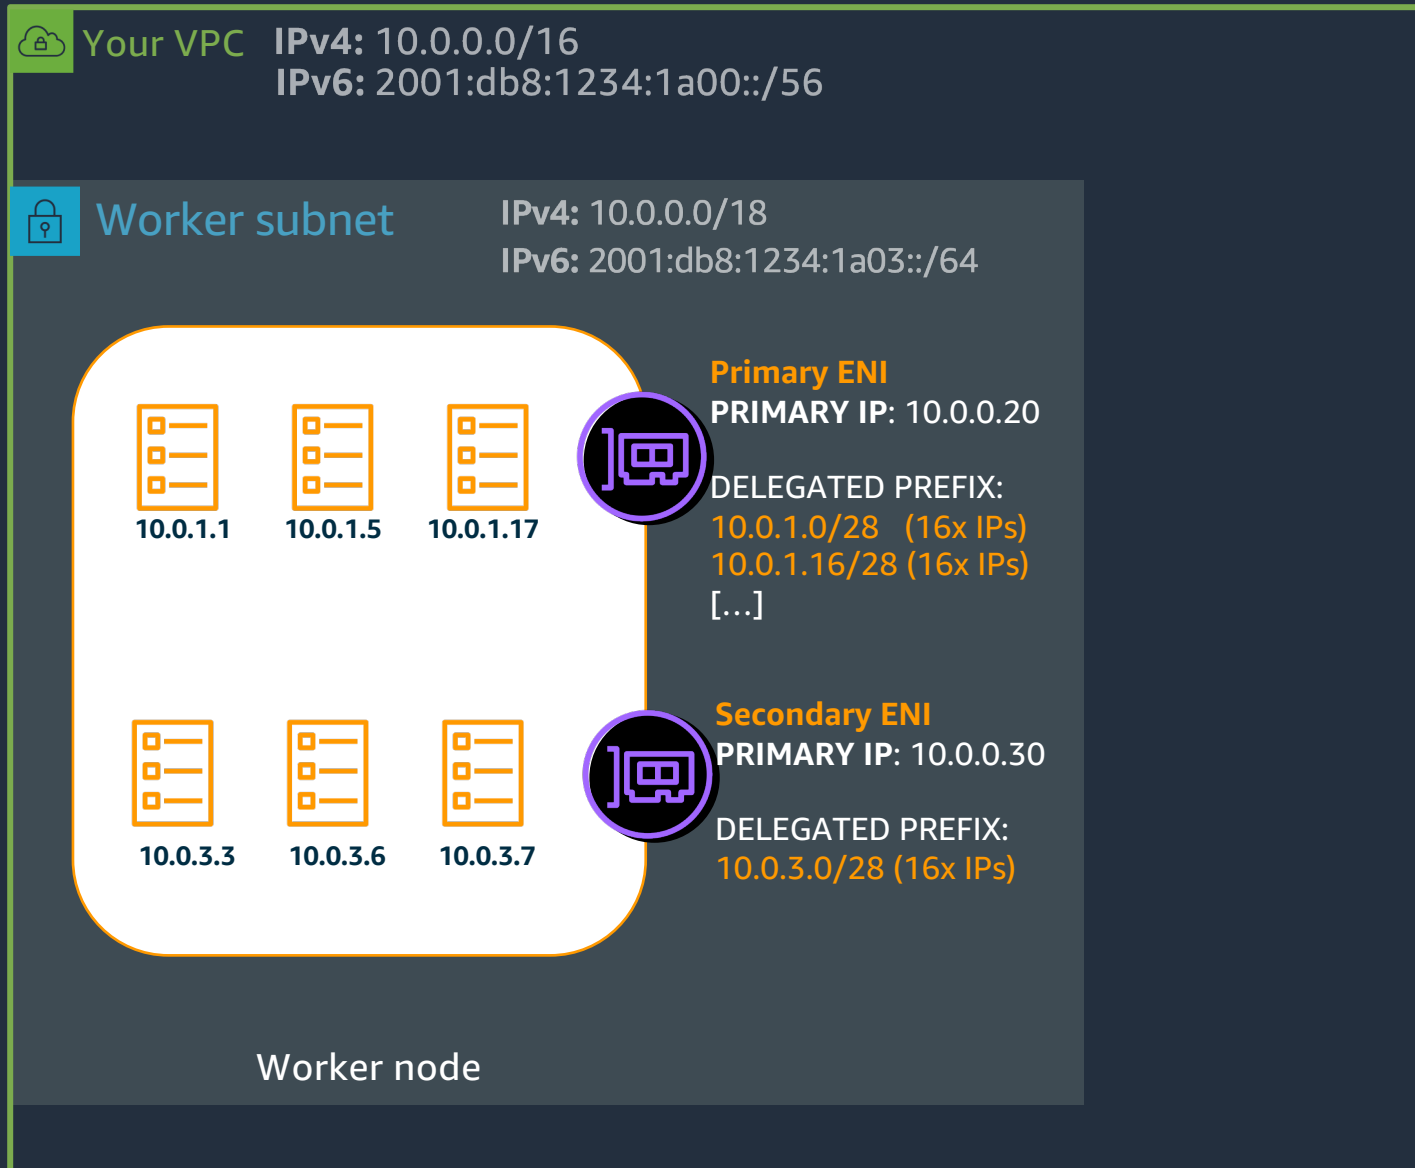
Instance type	Maximum network interfaces	Private IPv4 addresses per interface
m5.large	3	10

MAX IPs ≠ MAX PODS

432 IPv4
Trillions IPv6

110

Pod IPv4 networking – Prefix delegation

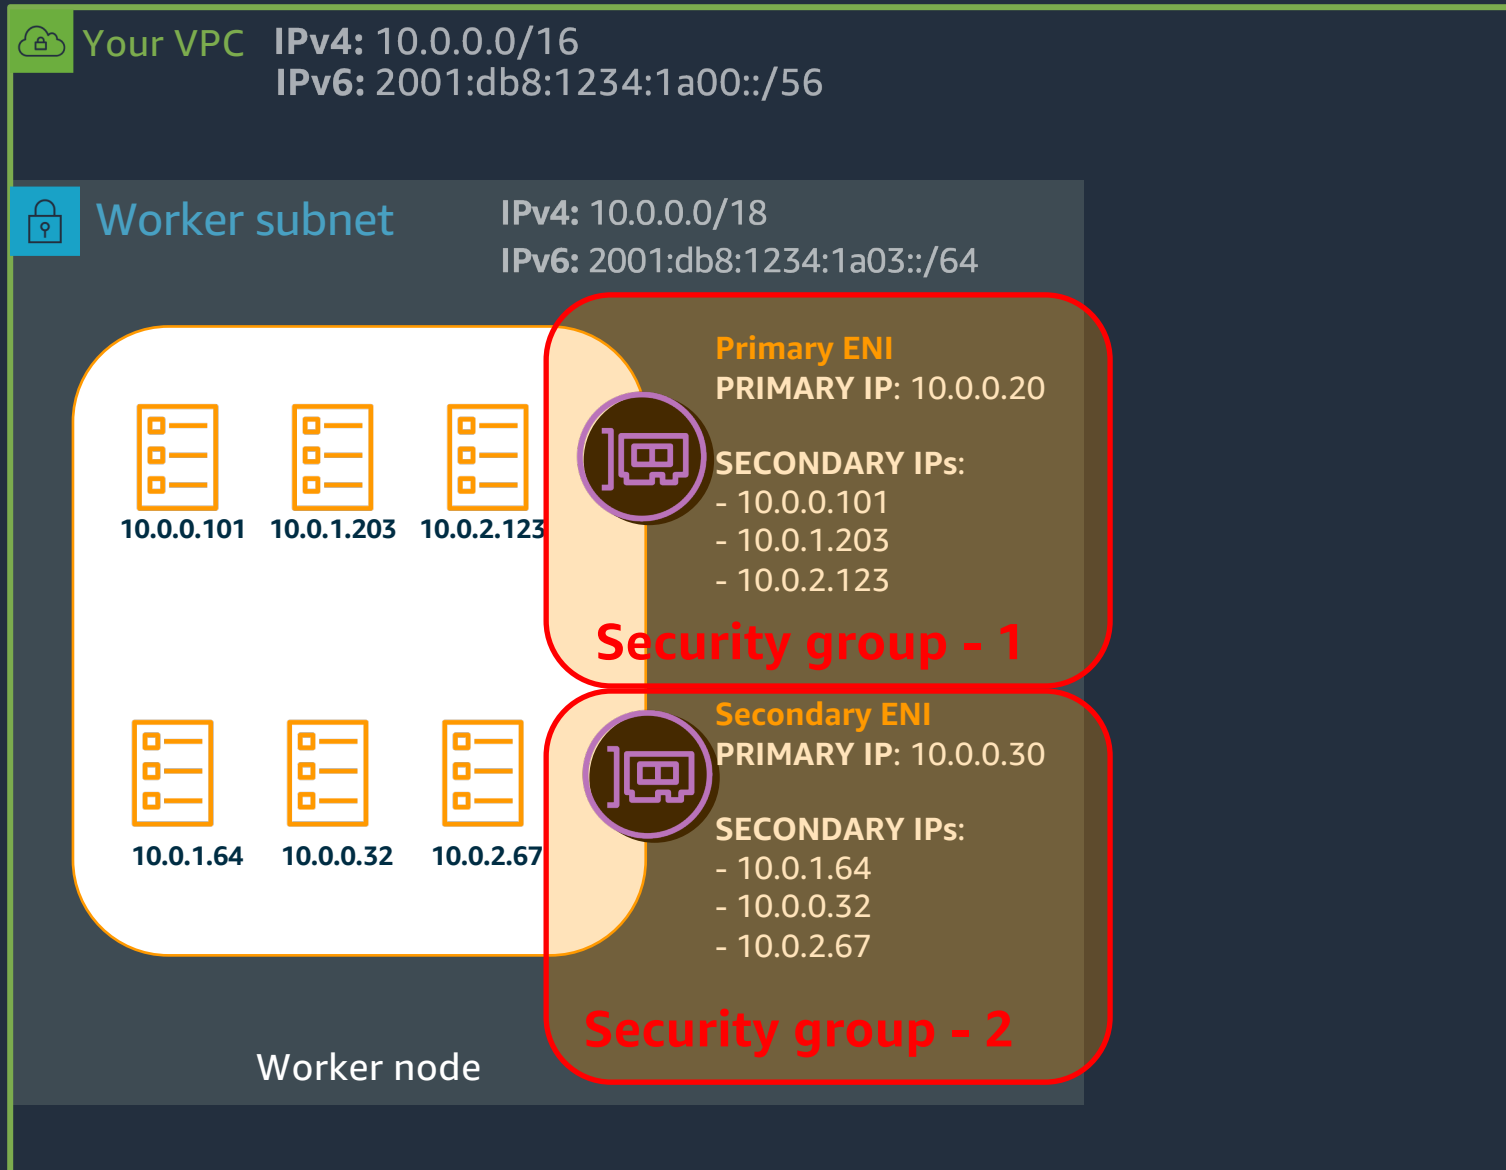


Benefits

- Increased Pod density
- Fewer API calls required to EC2 control plane


```
kubectl set env daemonset aws-node -n kube-system ENABLE_PREFIX_DELEGATION=true
```


Pod networking – Security groups for Pods



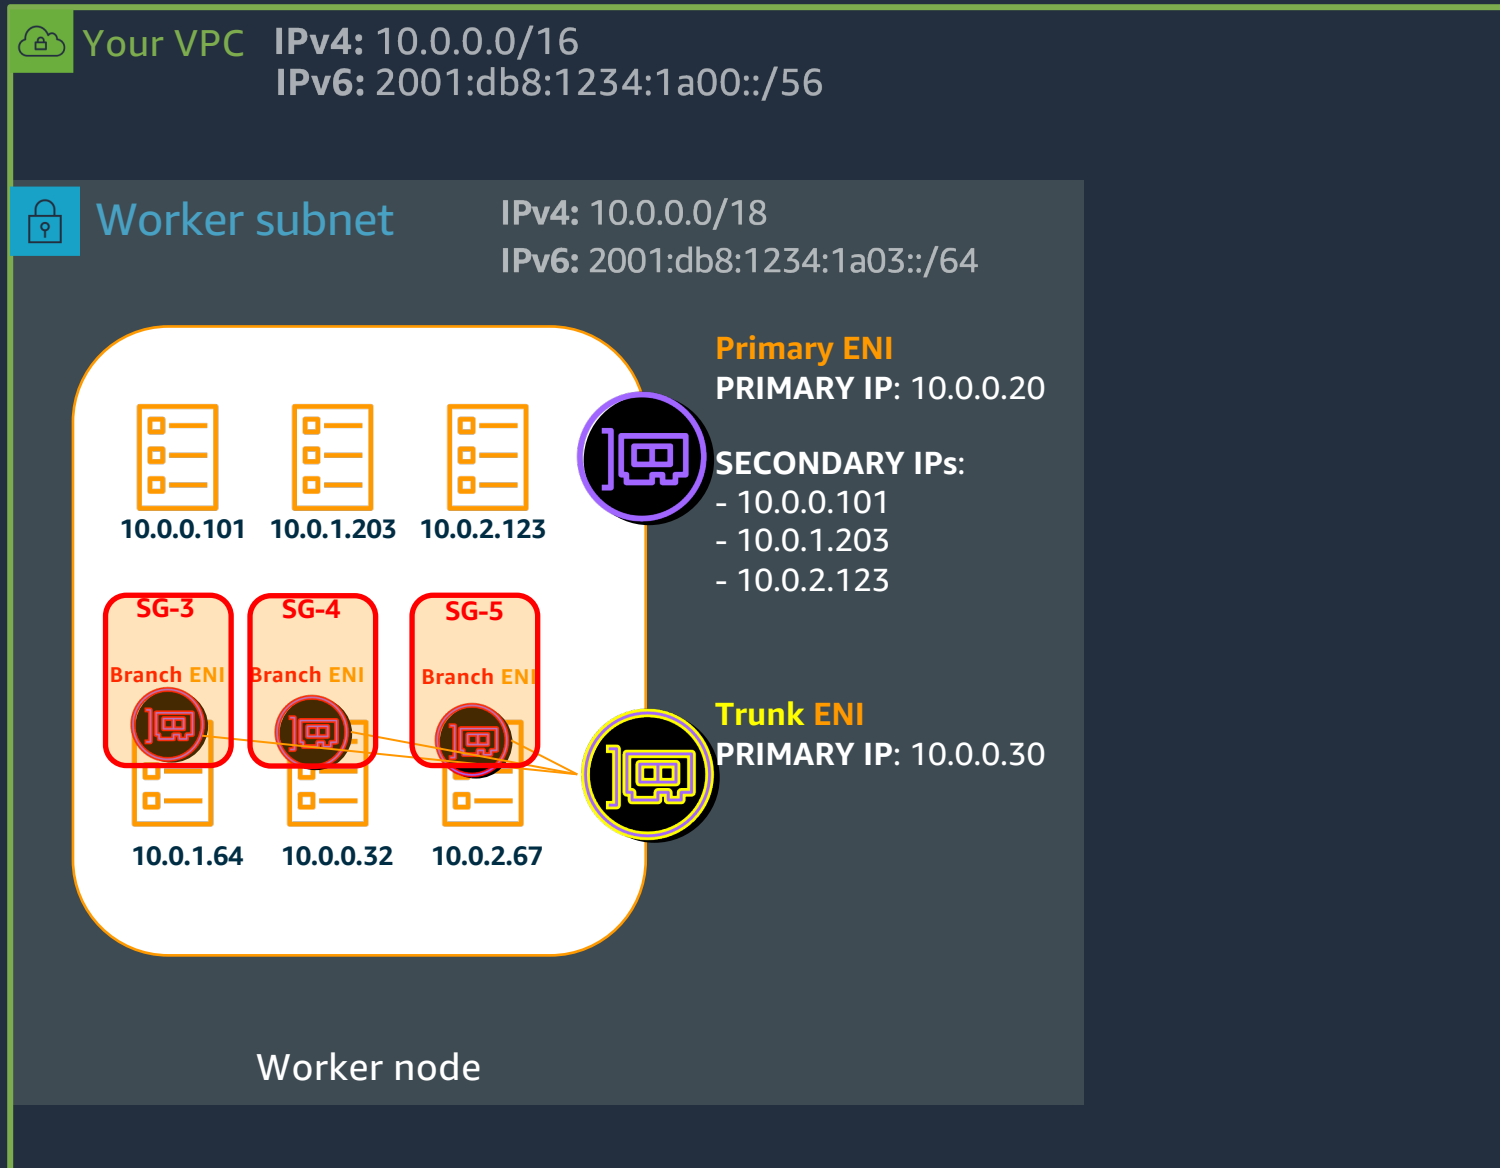
Pods share the security group on their respective ENI

Branch and trunk ENIs



Feature only available through Amazon EKS or Amazon ECS and not directly on Amazon EC2

Pod networking – Security groups for Pods



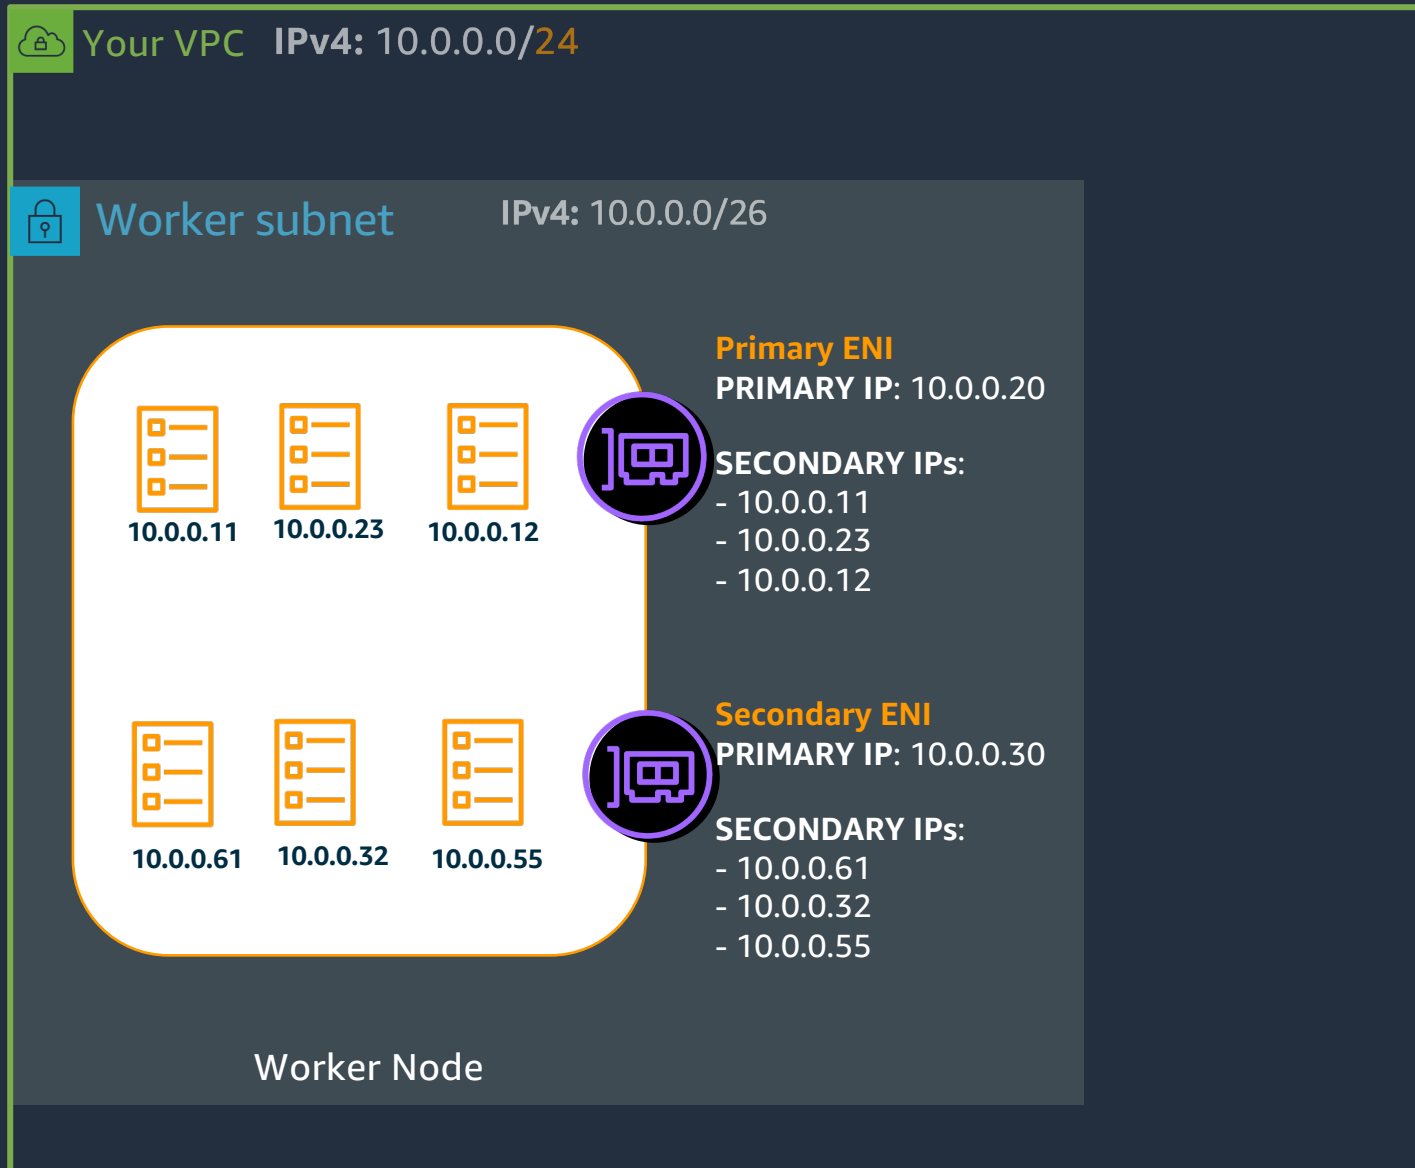
Each Pod gets a dedicated ENI (branch-ENI) mapped to a trunk ENI, allowing for independent security group configurations per Pod (**ENABLE_POD_ENI=true**)

Considerations

- Supported by most Nitro instances
- SNAT is disabled for Pods using branch ENIs
- Cluster version 1.18+
- Review integration with NodePort and LoadBalancer services
- [vpc.amazonaws.com/has-trunk-attached=true](https://docs.aws.amazon.com/ko_kr/AmazonECS/latest/developerguide/container-instance-eni.html)
- DISABLE_TCP_EARLY_DEMUX=true
- Branch ENI Limits
 - https://docs.aws.amazon.com/ko_kr/AmazonECS/latest/developerguide/container-instance-eni.html

```
kubectl set env daemonset aws-node -n kube-system ENABLE_POD_ENI=true
```


Pod networking – Custom networking



If the private IPv4 space is limited, the number of available IP addresses could constrain the number of Pods

/24 VPC CIDR provides 251 unique IPv4 addresses

Pod networking – Custom networking



VPCs can have multiple IPv4 CIDR ranges

100.64.0.0/10 (RFC 6598) can be used in private networks

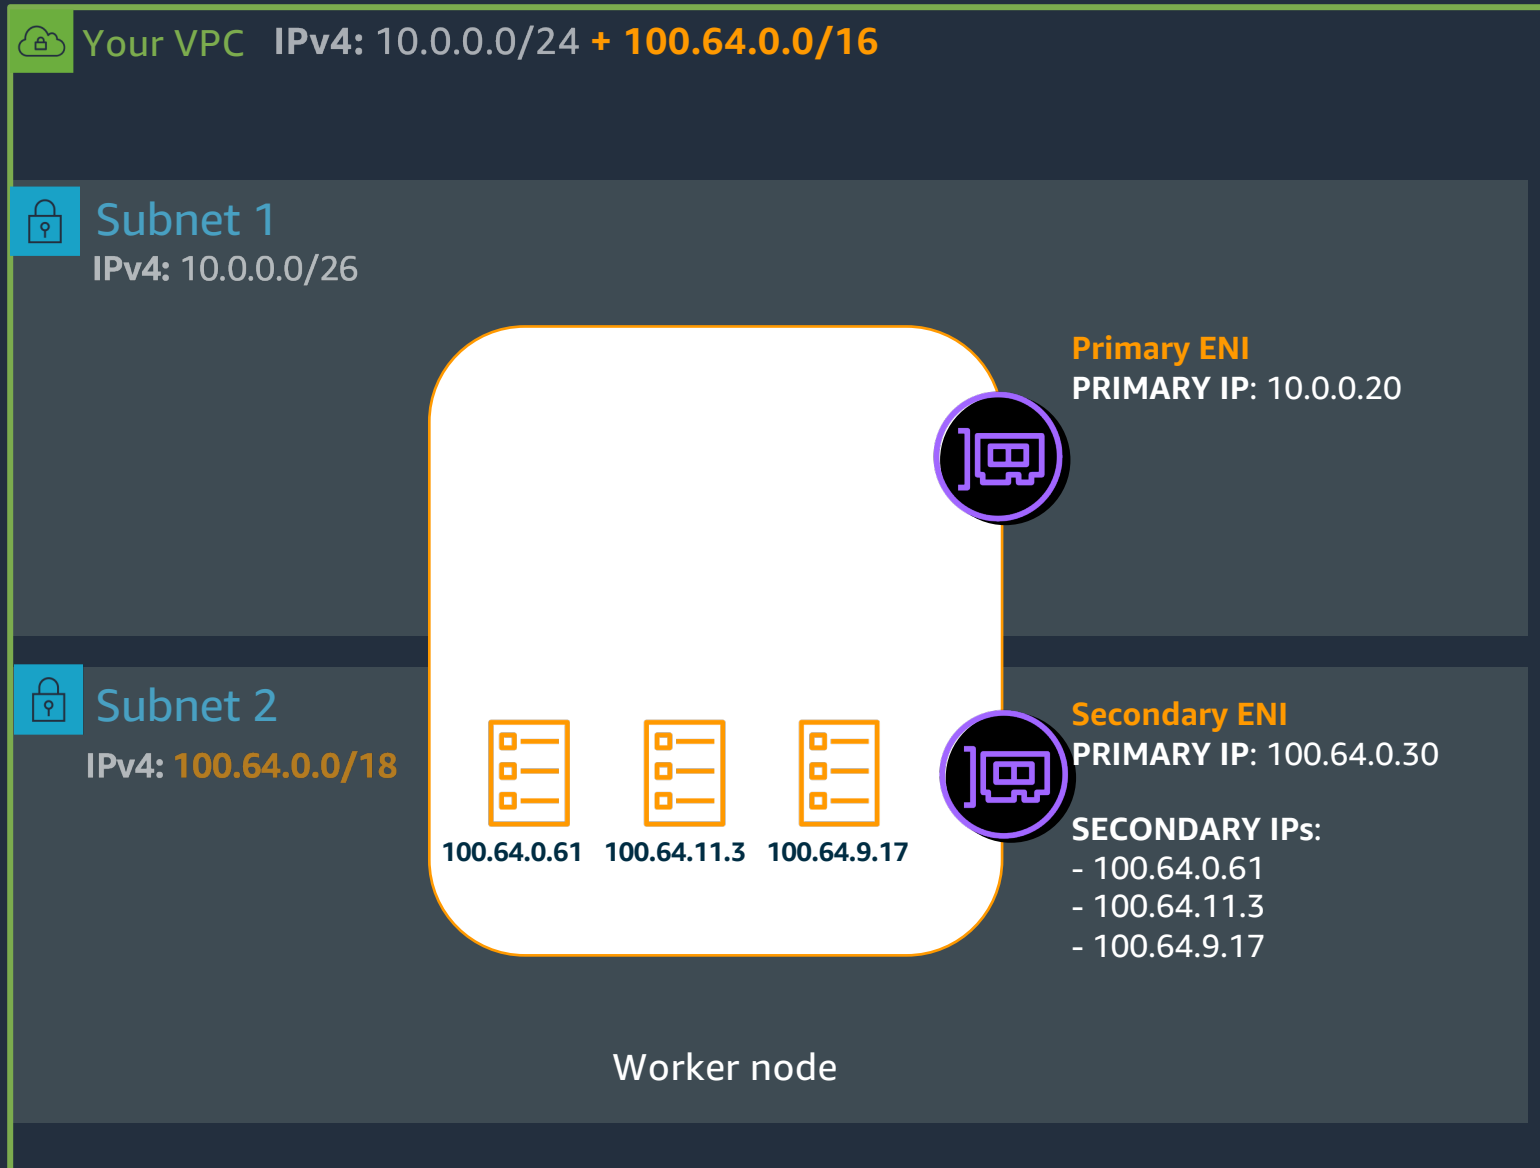
Worker node primary ENI is not used for Pods

Custom networking can be combined with SNAT and prefix delegation

```
apiVersion: crd.k8s.amazonaws.com/v1alpha1
kind: ENIConfig
metadata:
  name: us-west-2a
spec:
  securityGroups:
    - sg-0dff111a1d11c1c11
  subnet: subnet-011b111c1f11fdf11
```

```
kubectl set env daemonset aws-node -n kube-system AWS_VPC_K8S_CNI_CUSTOM_NETWORK_CFG=true
```

Pod networking – Custom networking

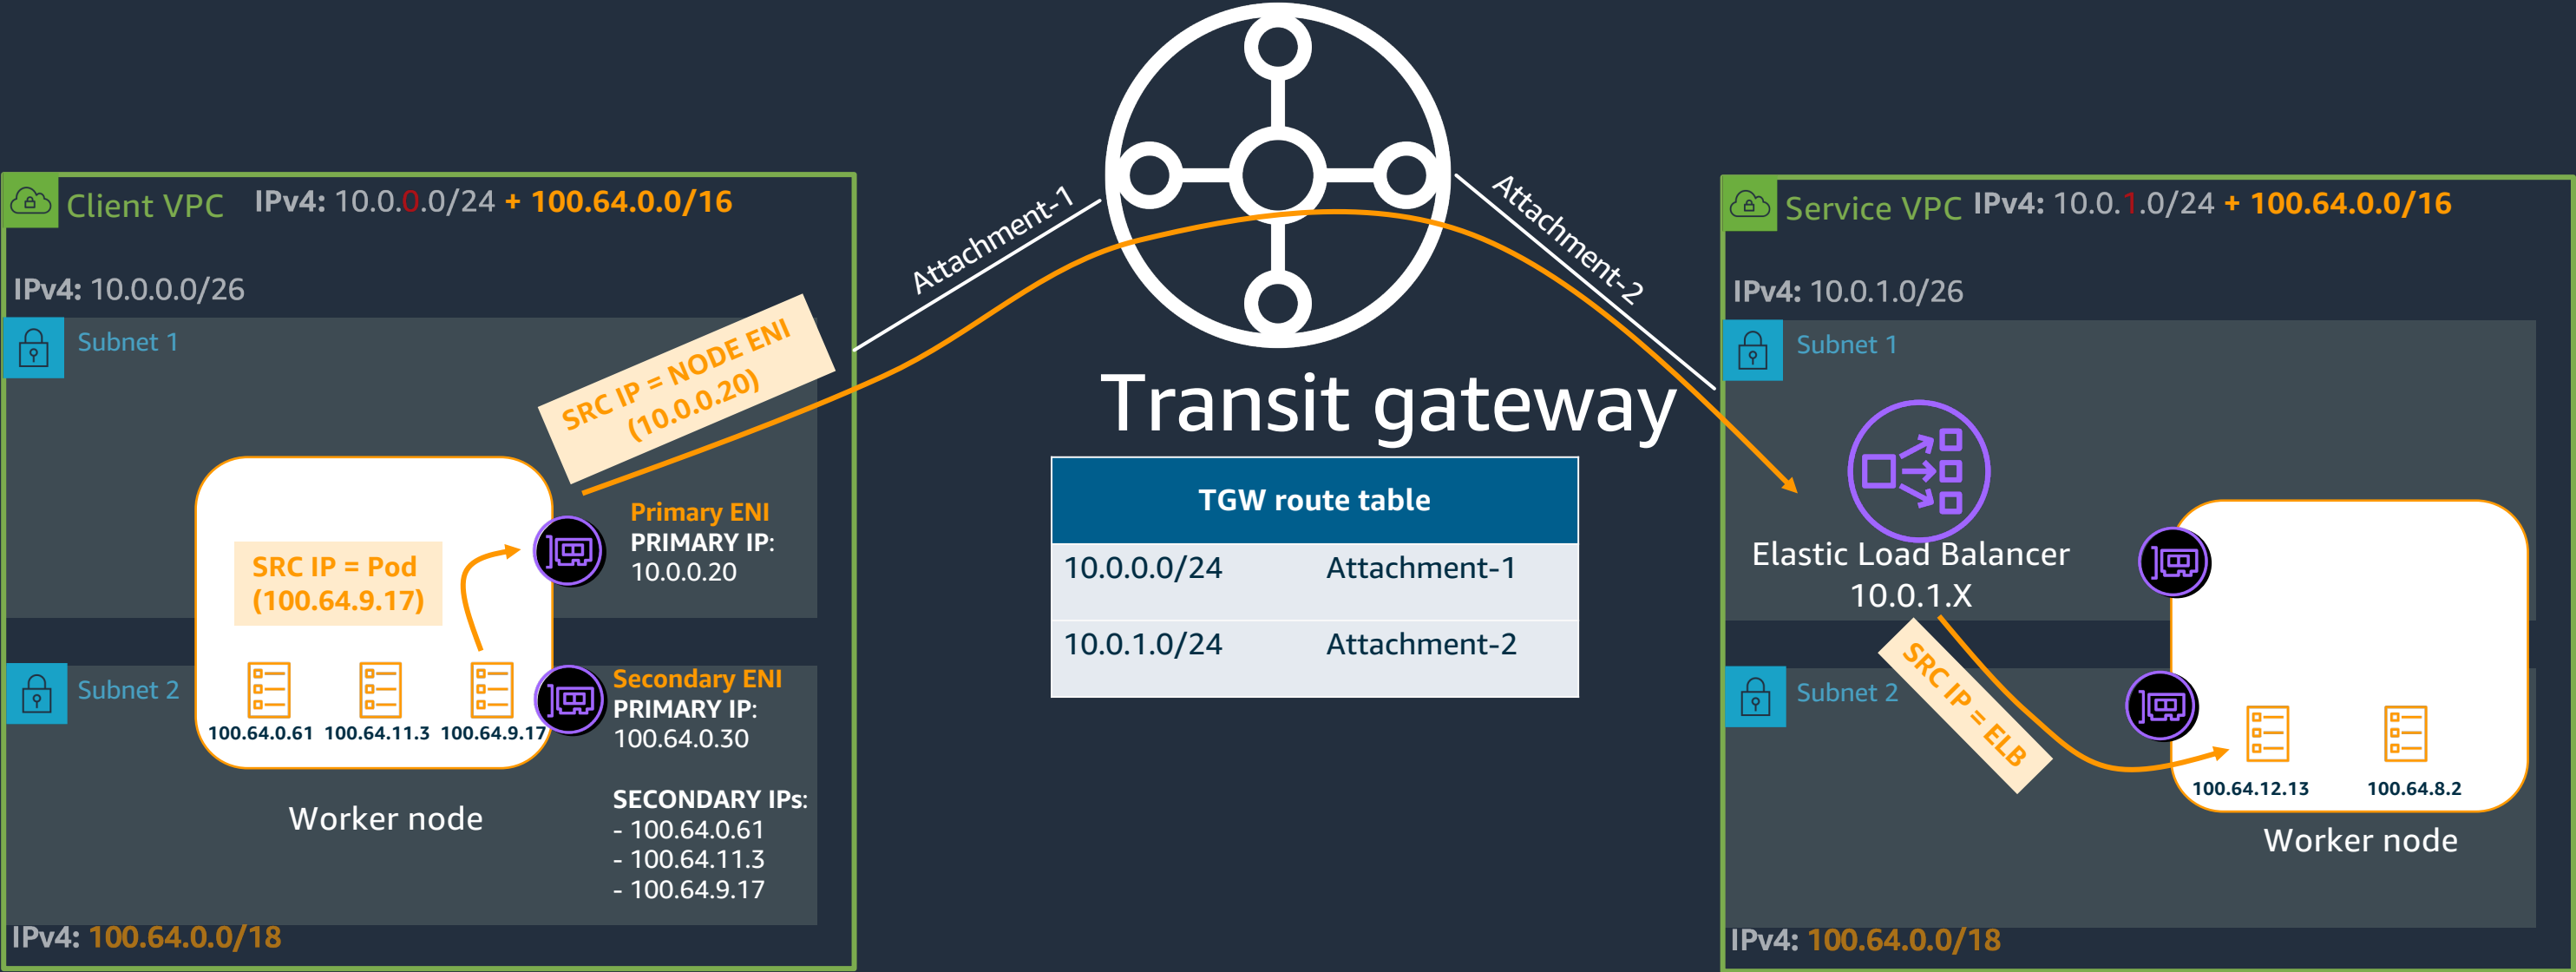


Kubelet BootstrapArguments parameter

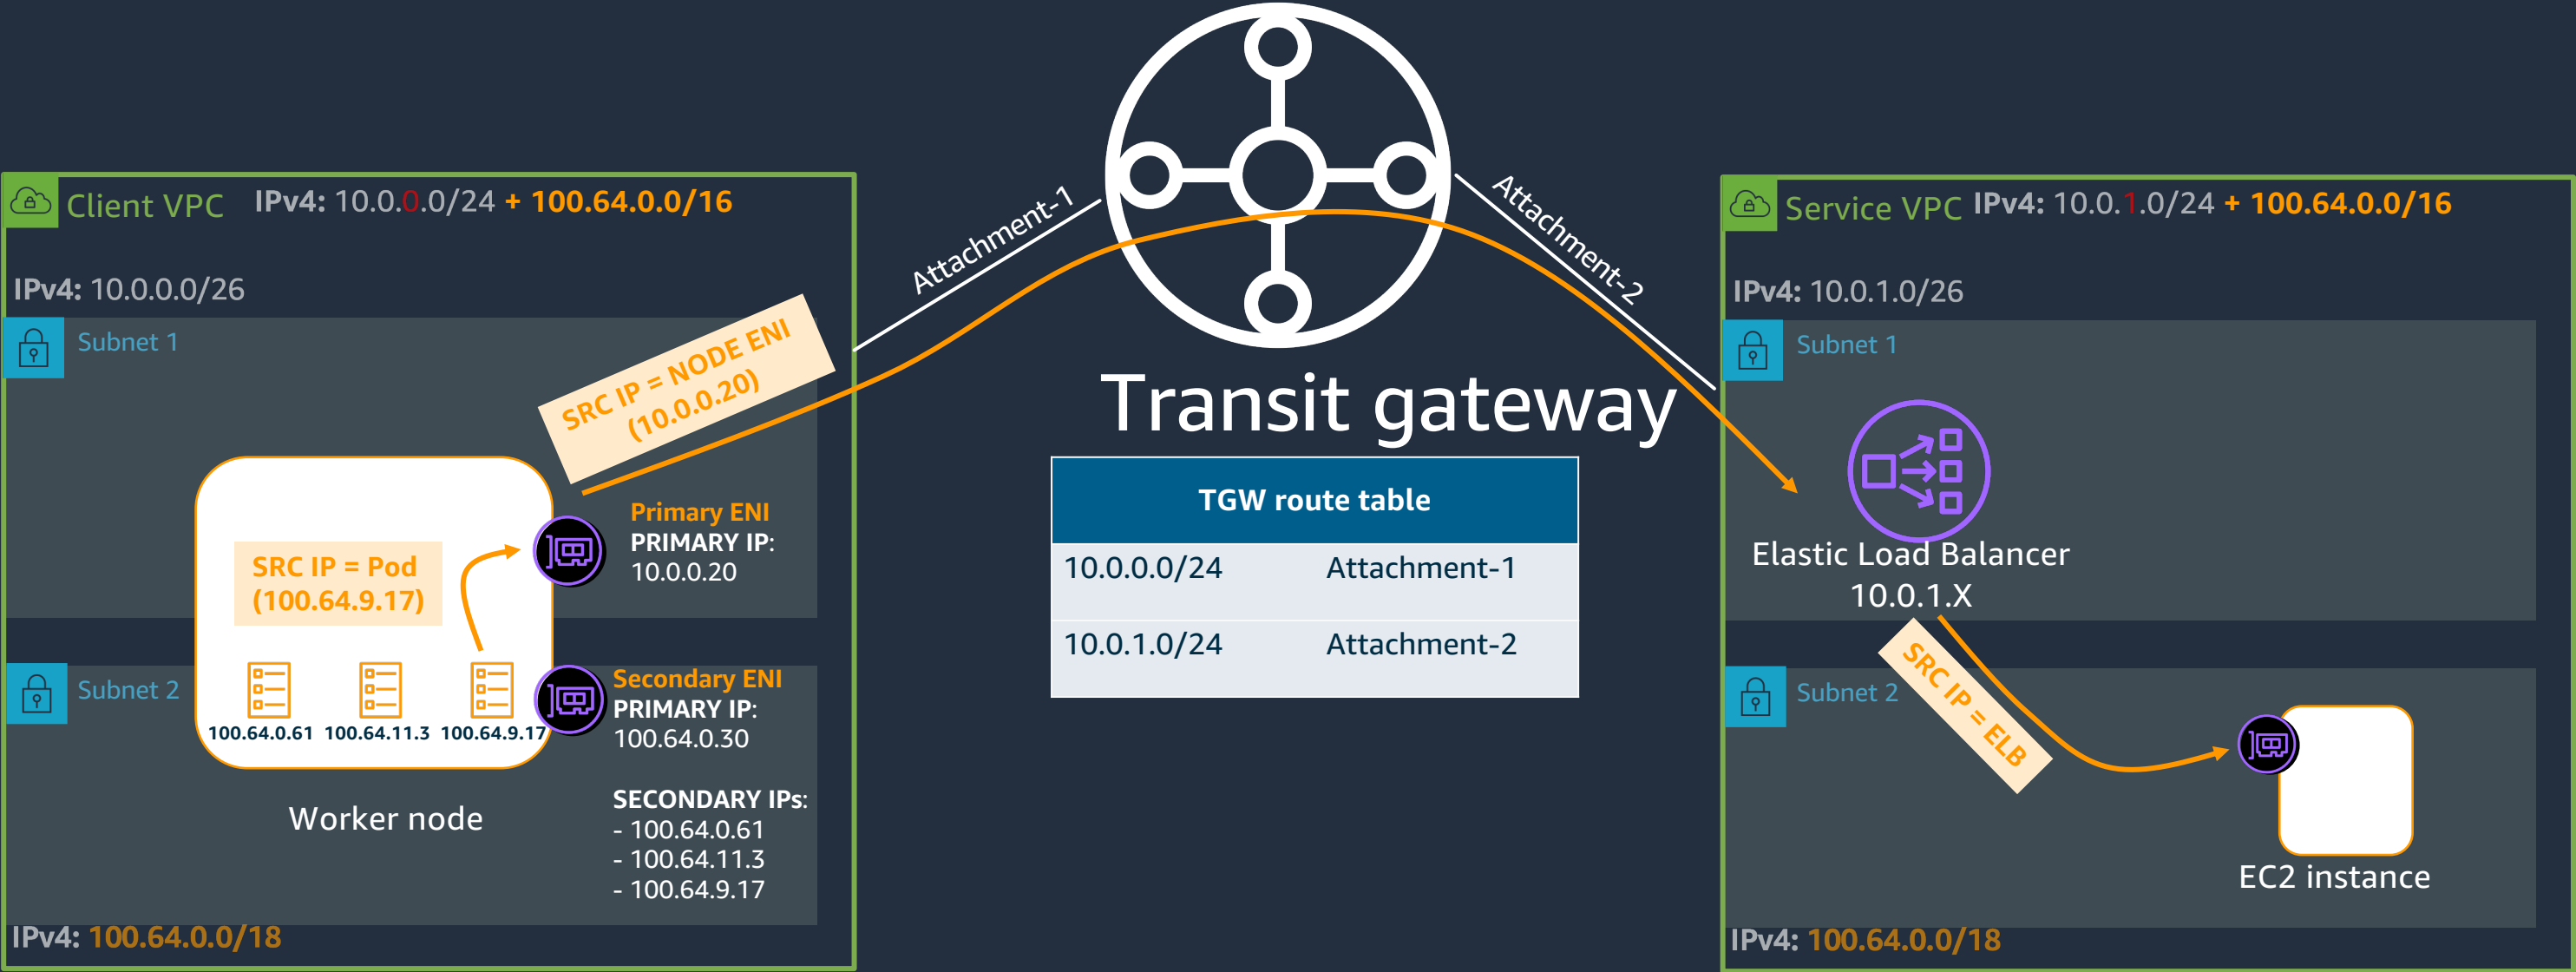
Change "--max-pod=29" to "--max-pod=110"

```
--use-max-pods false --kubelet-extra-args '--max-pods=110'
```

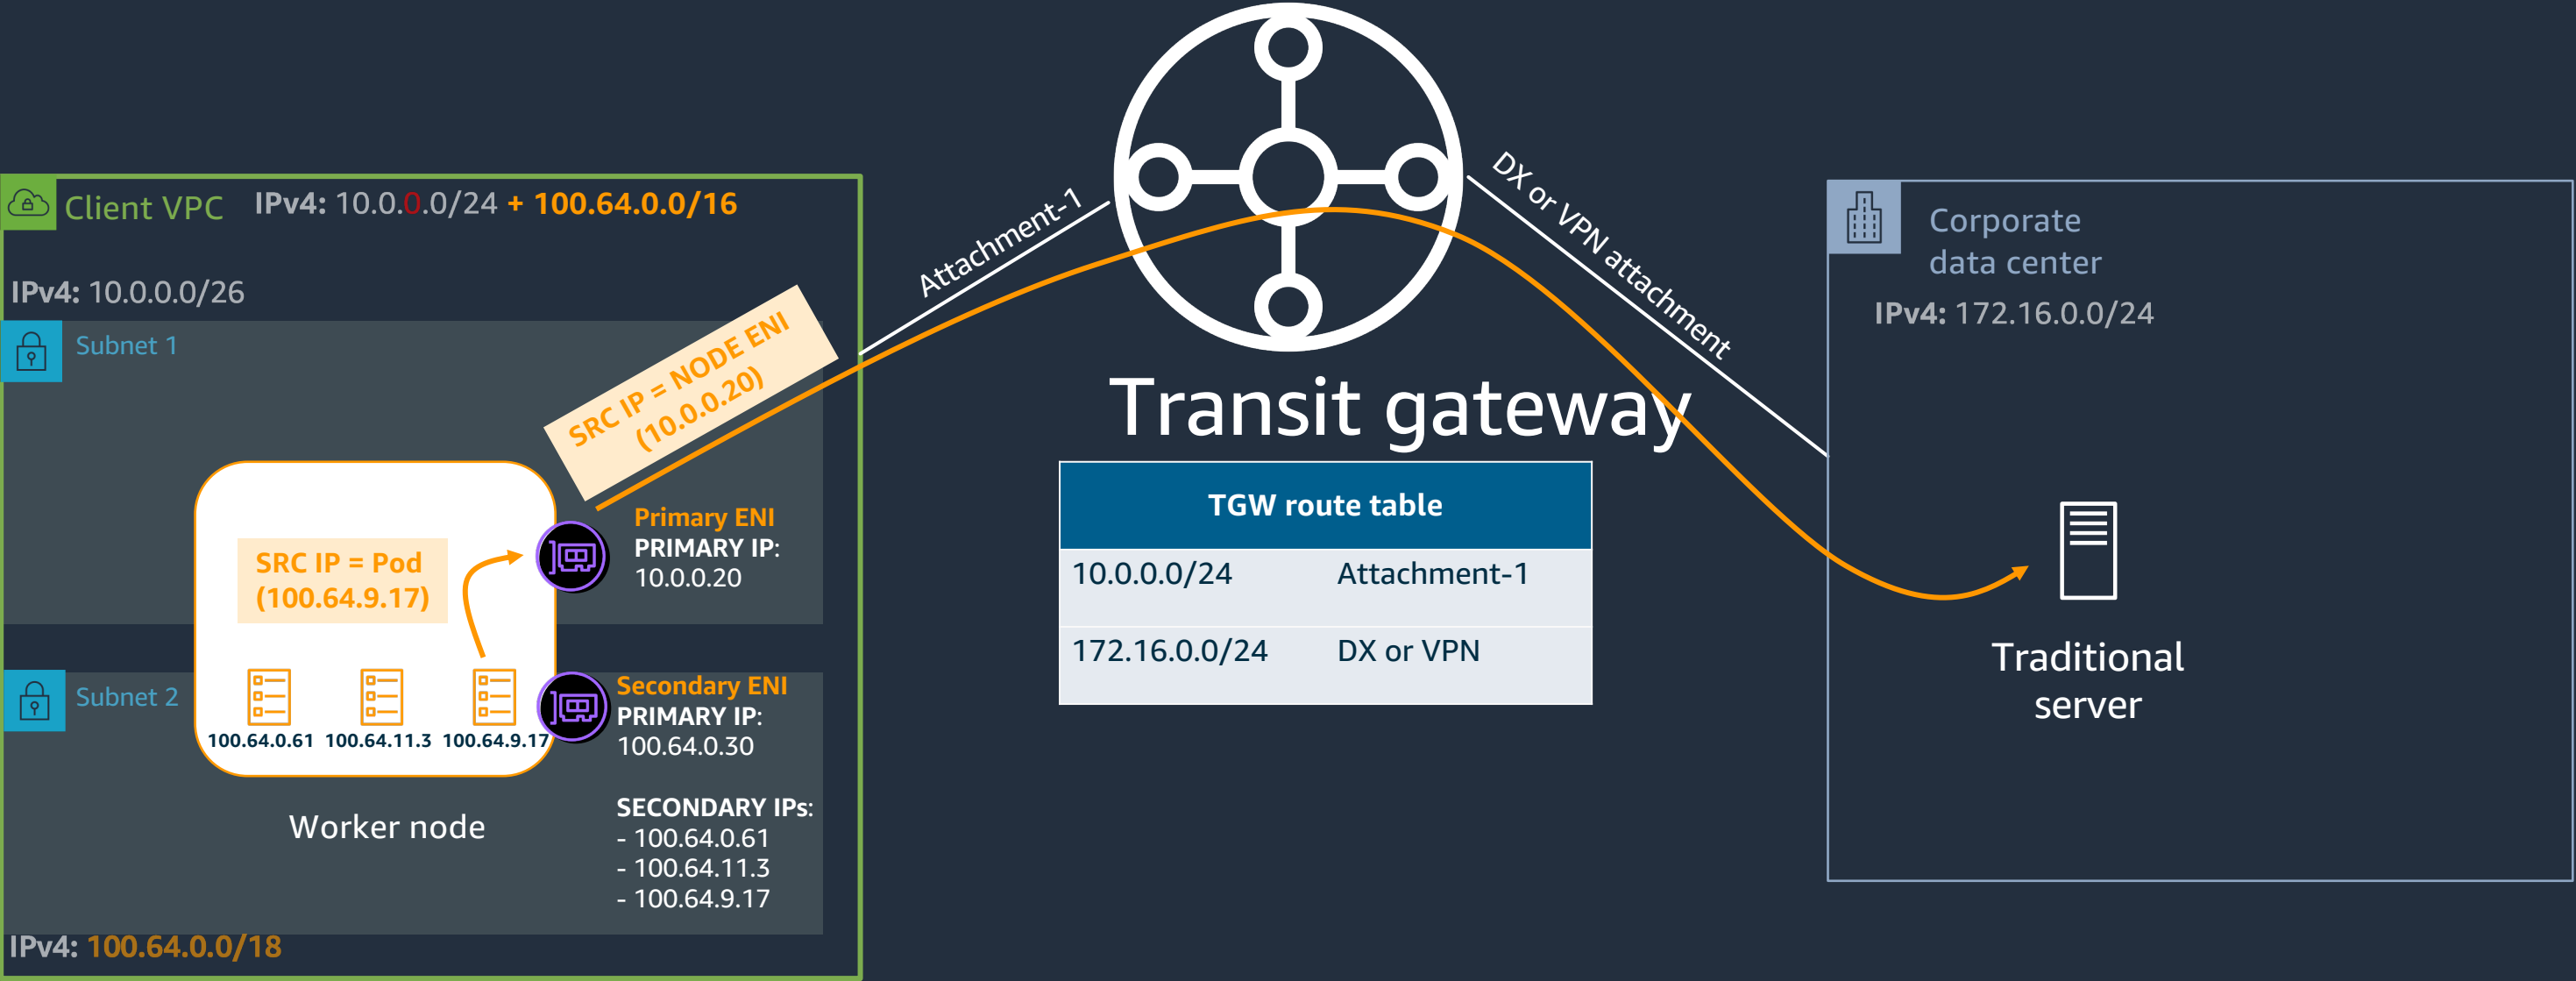
Custom networking + SNAT + transit gateway




Custom networking + SNAT + transit gateway



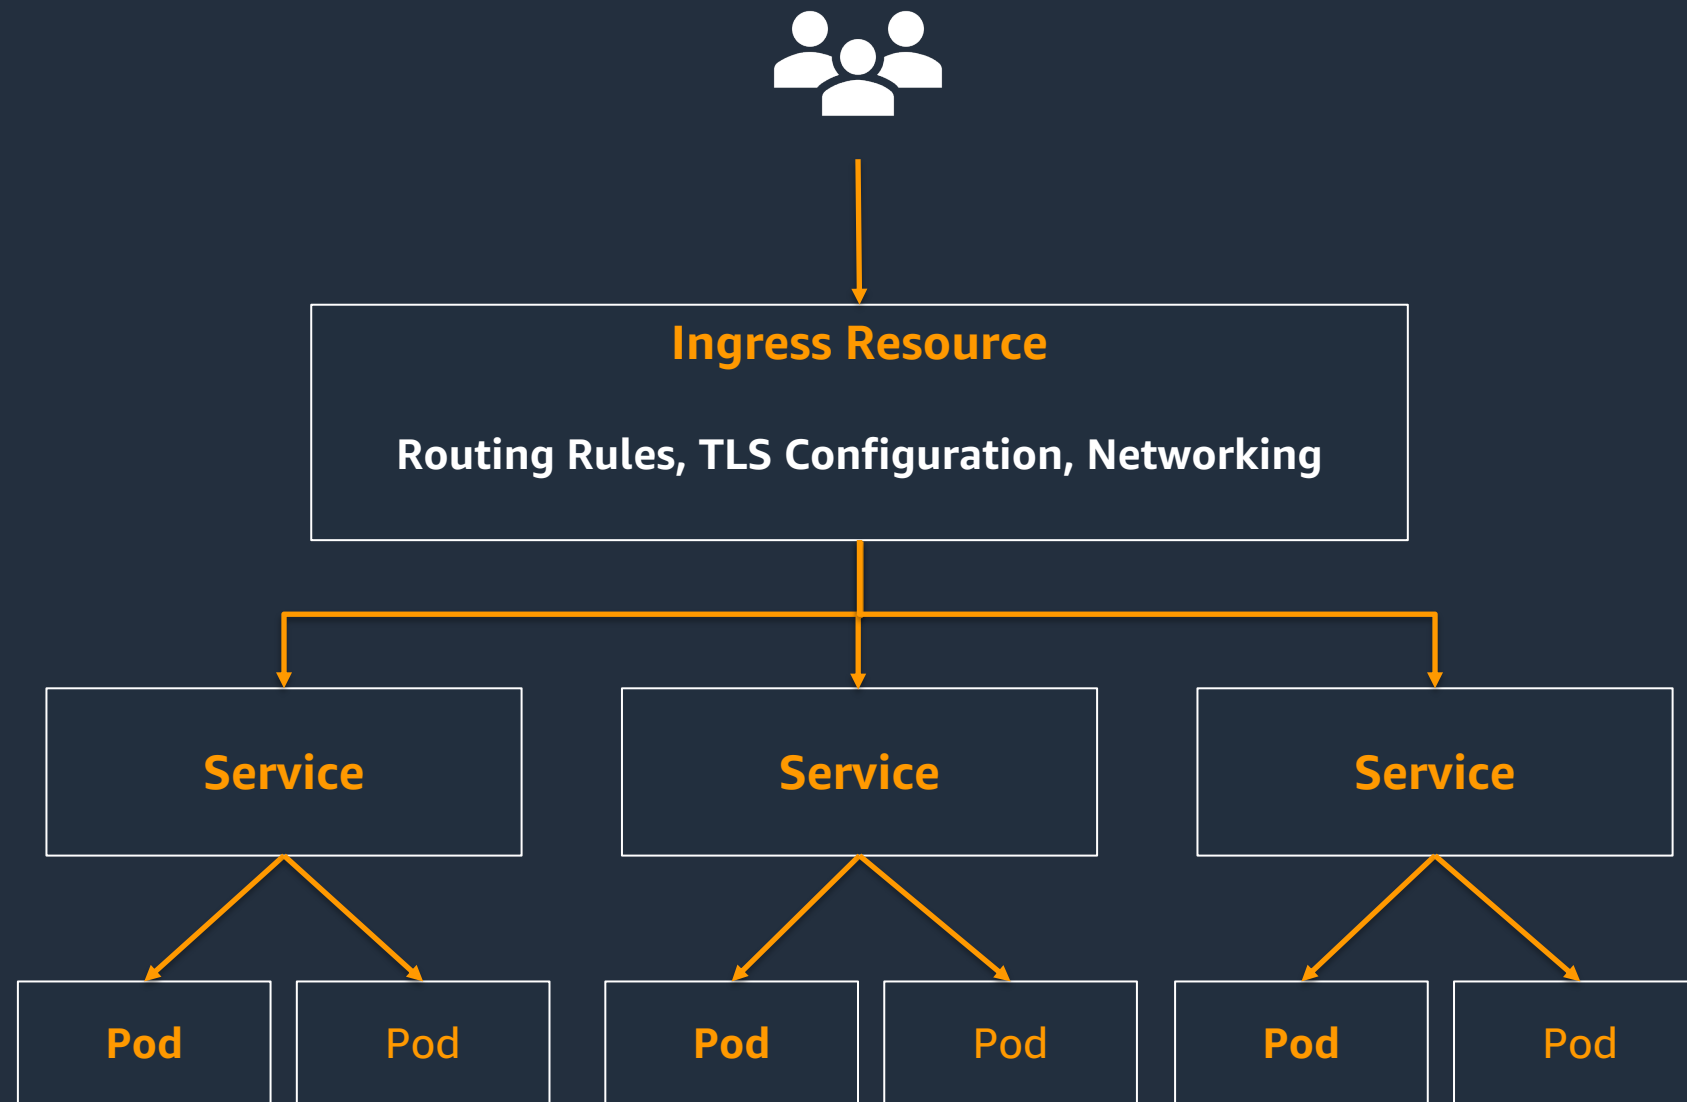
Custom networking + SNAT + transit gateway



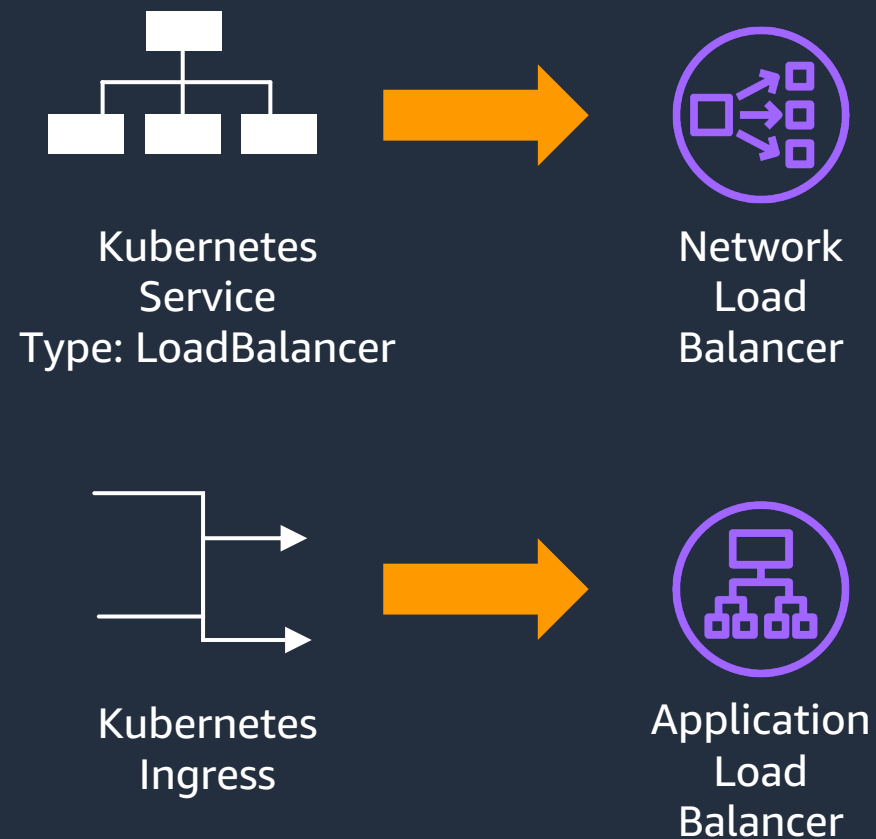
Ingress Controller



Kubernetes Ingress

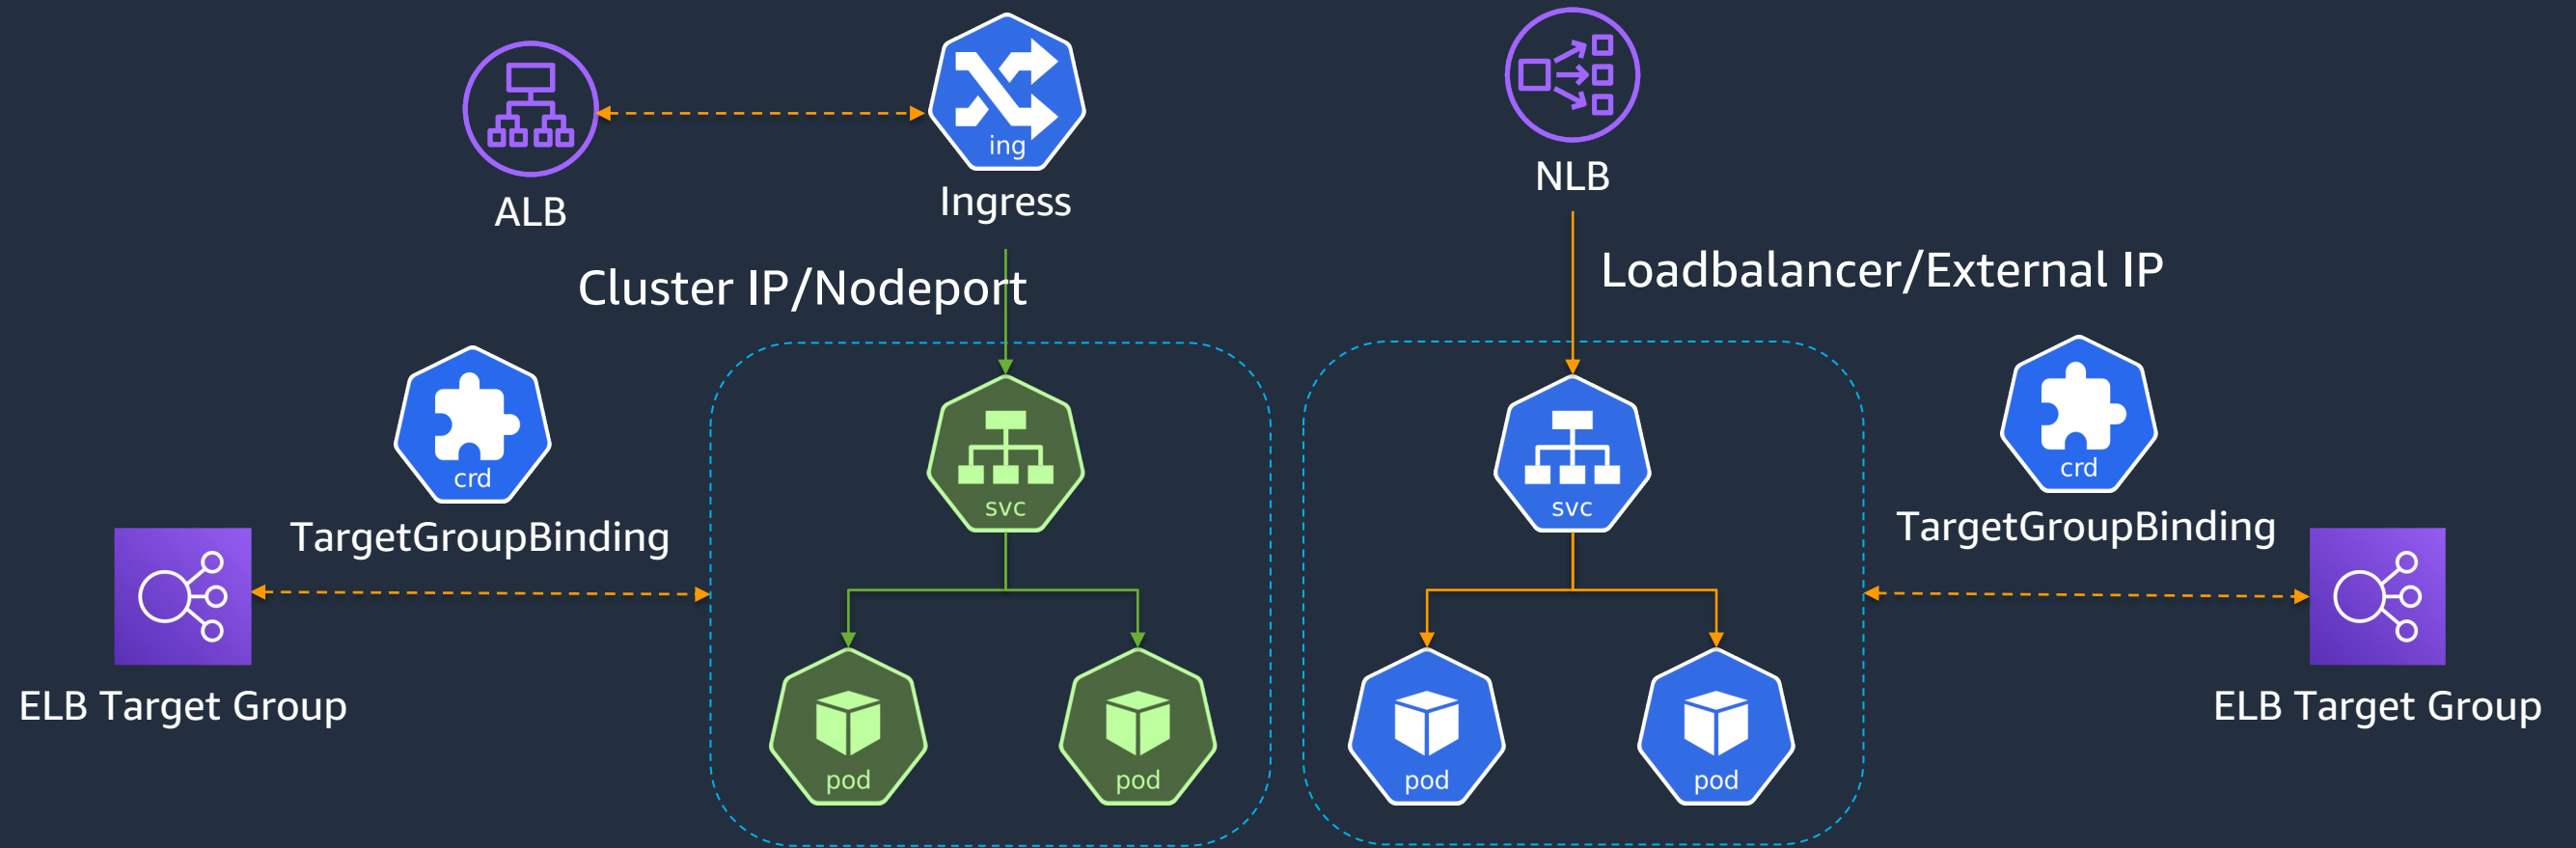



AWS Load Balancer Controller

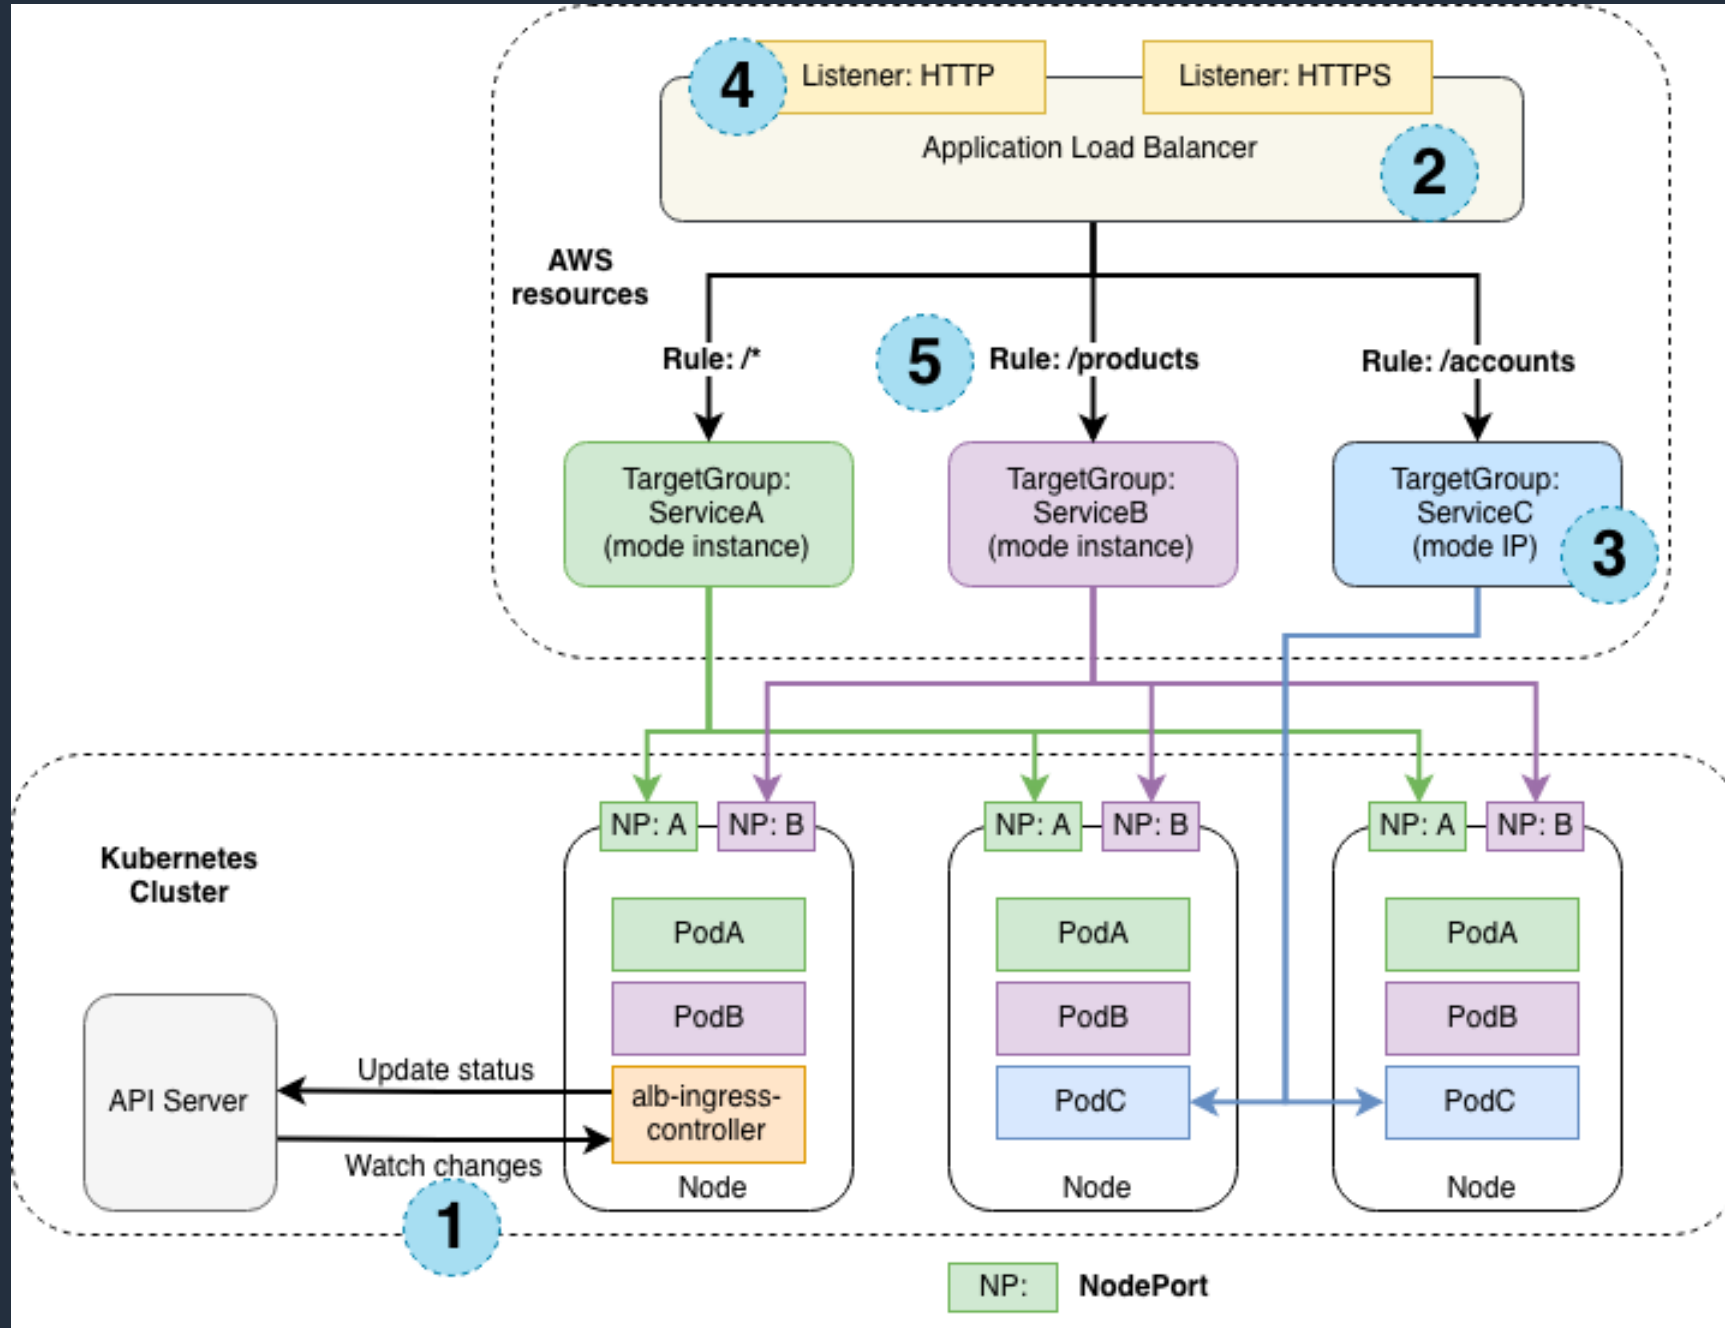


- Manages AWS Elastic Load Balancing for an Amazon EKS or a Kubernetes cluster
- Uses standard Kubernetes resources
 - v1: Service (Type: LoadBalancer)
 - networking.k8s.io/v1: Ingress
 - networking.k8s.io/v1: IngressClass
- Custom resources
 - elbv2.k8s.aws/v1beta1: TargetGroupBinding
- Supports latest ELB features


AWS Load Balancer Controller




Ingress with AWS CNI



1. Instance Mode



2. IP Mode

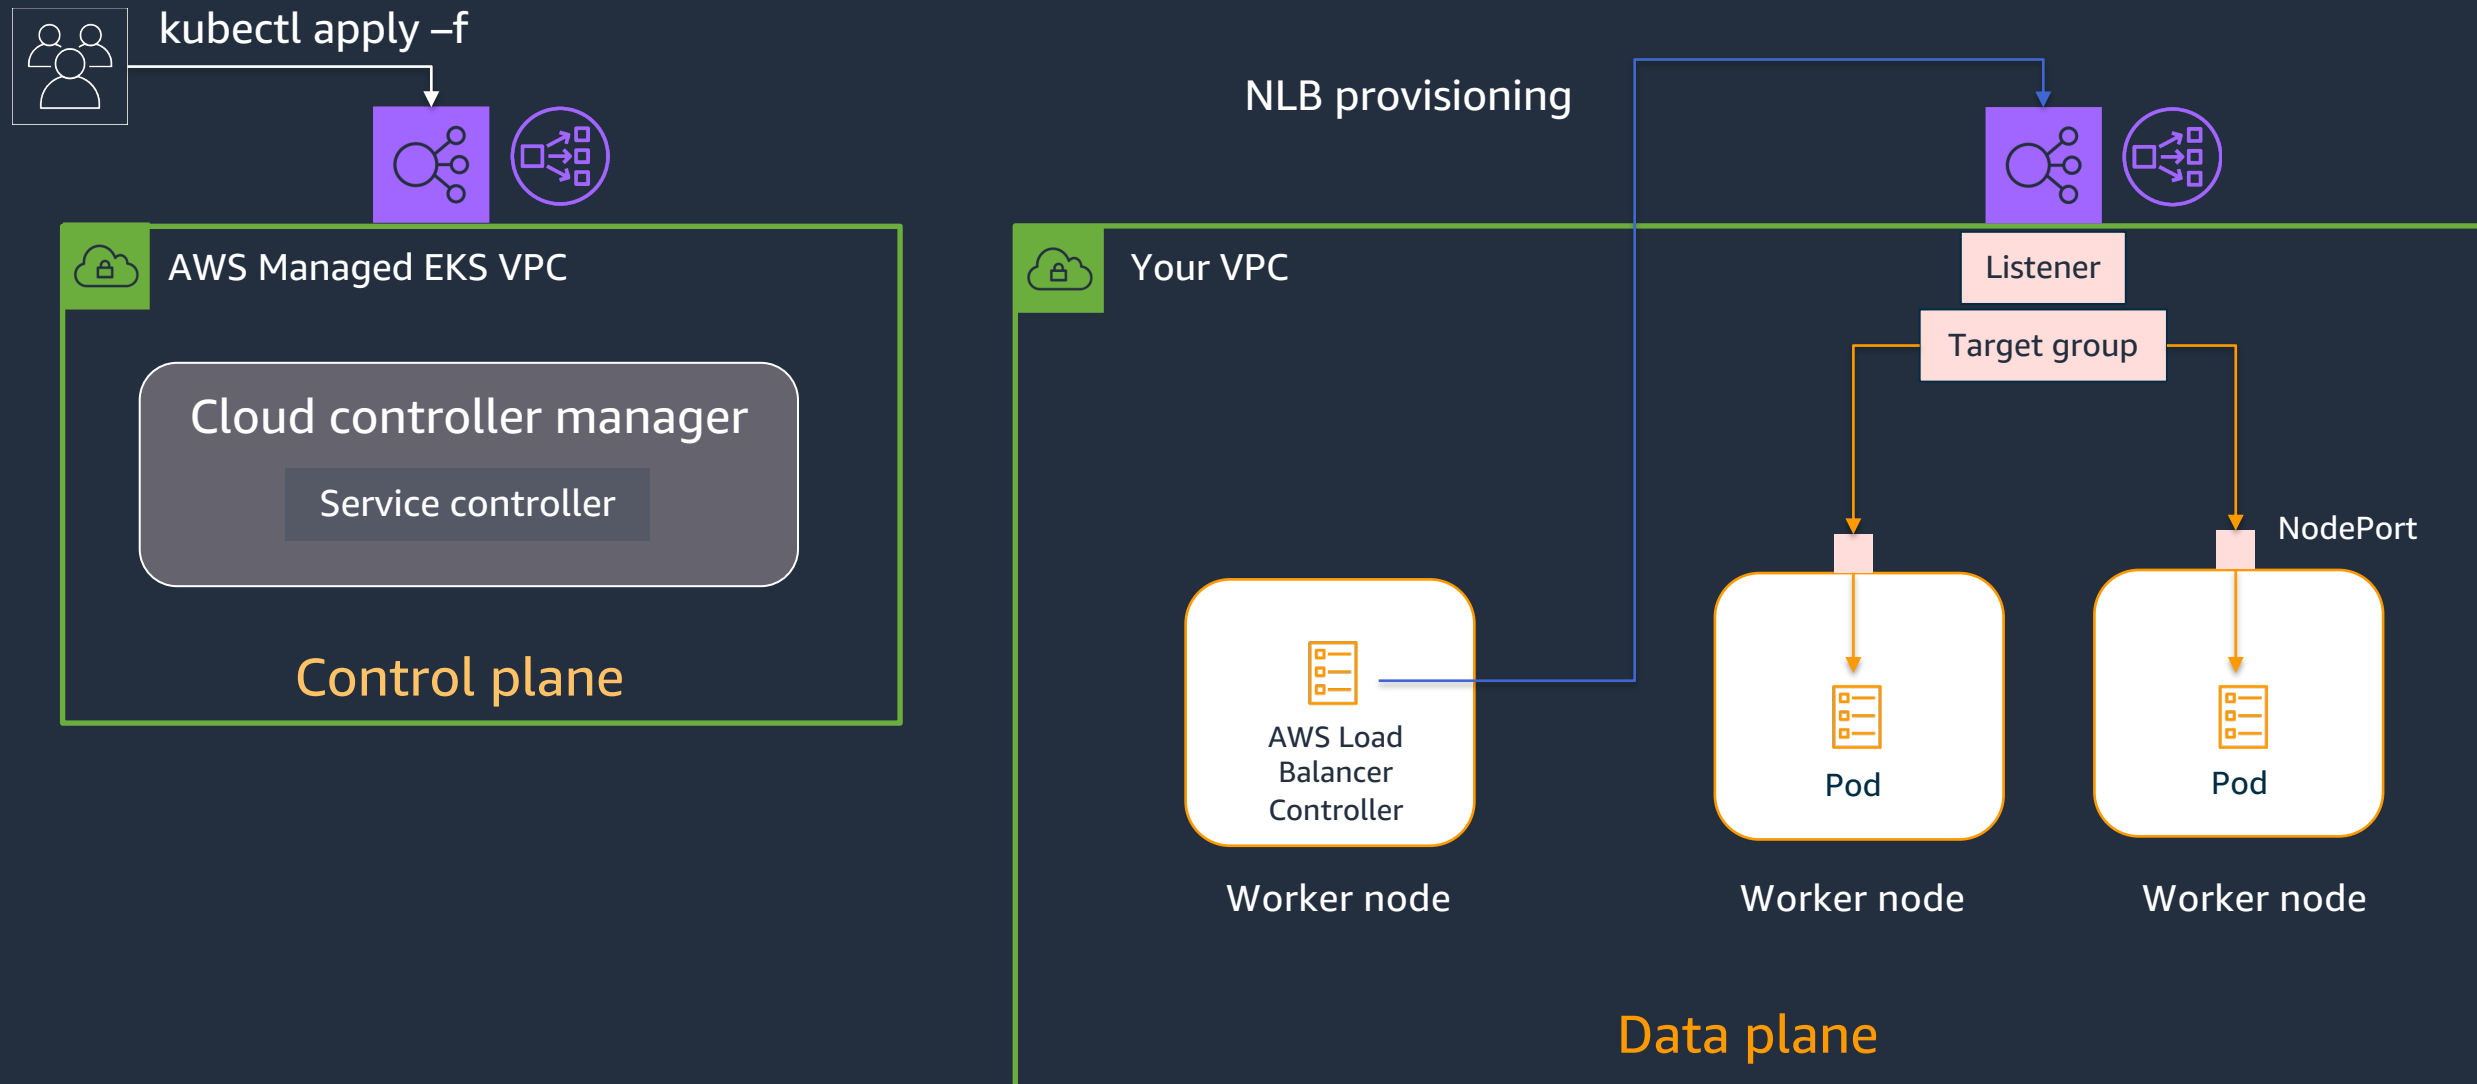


Network Load Balancer provisioning in EKS

```
kind: Service
apiVersion: v1
metadata:
  name: nginx-service
  namespace: ingress-nginx
  annotations:
    service.beta.kubernetes.io/aws-load-balancer-backend-protocol: tcp
    service.beta.kubernetes.io/aws-load-balancer-cross-zone-load-balancing-enabled: "true"
    service.beta.kubernetes.io/aws-load-balancer-type: "external"
    service.beta.kubernetes.io/aws-load-balancer-nlb-target-type: "instance"
    service.beta.kubernetes.io/aws-load-balancer-scheme: internet-facing
spec:
  type: LoadBalancer
  selector:
    app: nginx
  ports:
    - name: http
      protocol: TCP
      port: 8080
      targetPort: 80
```

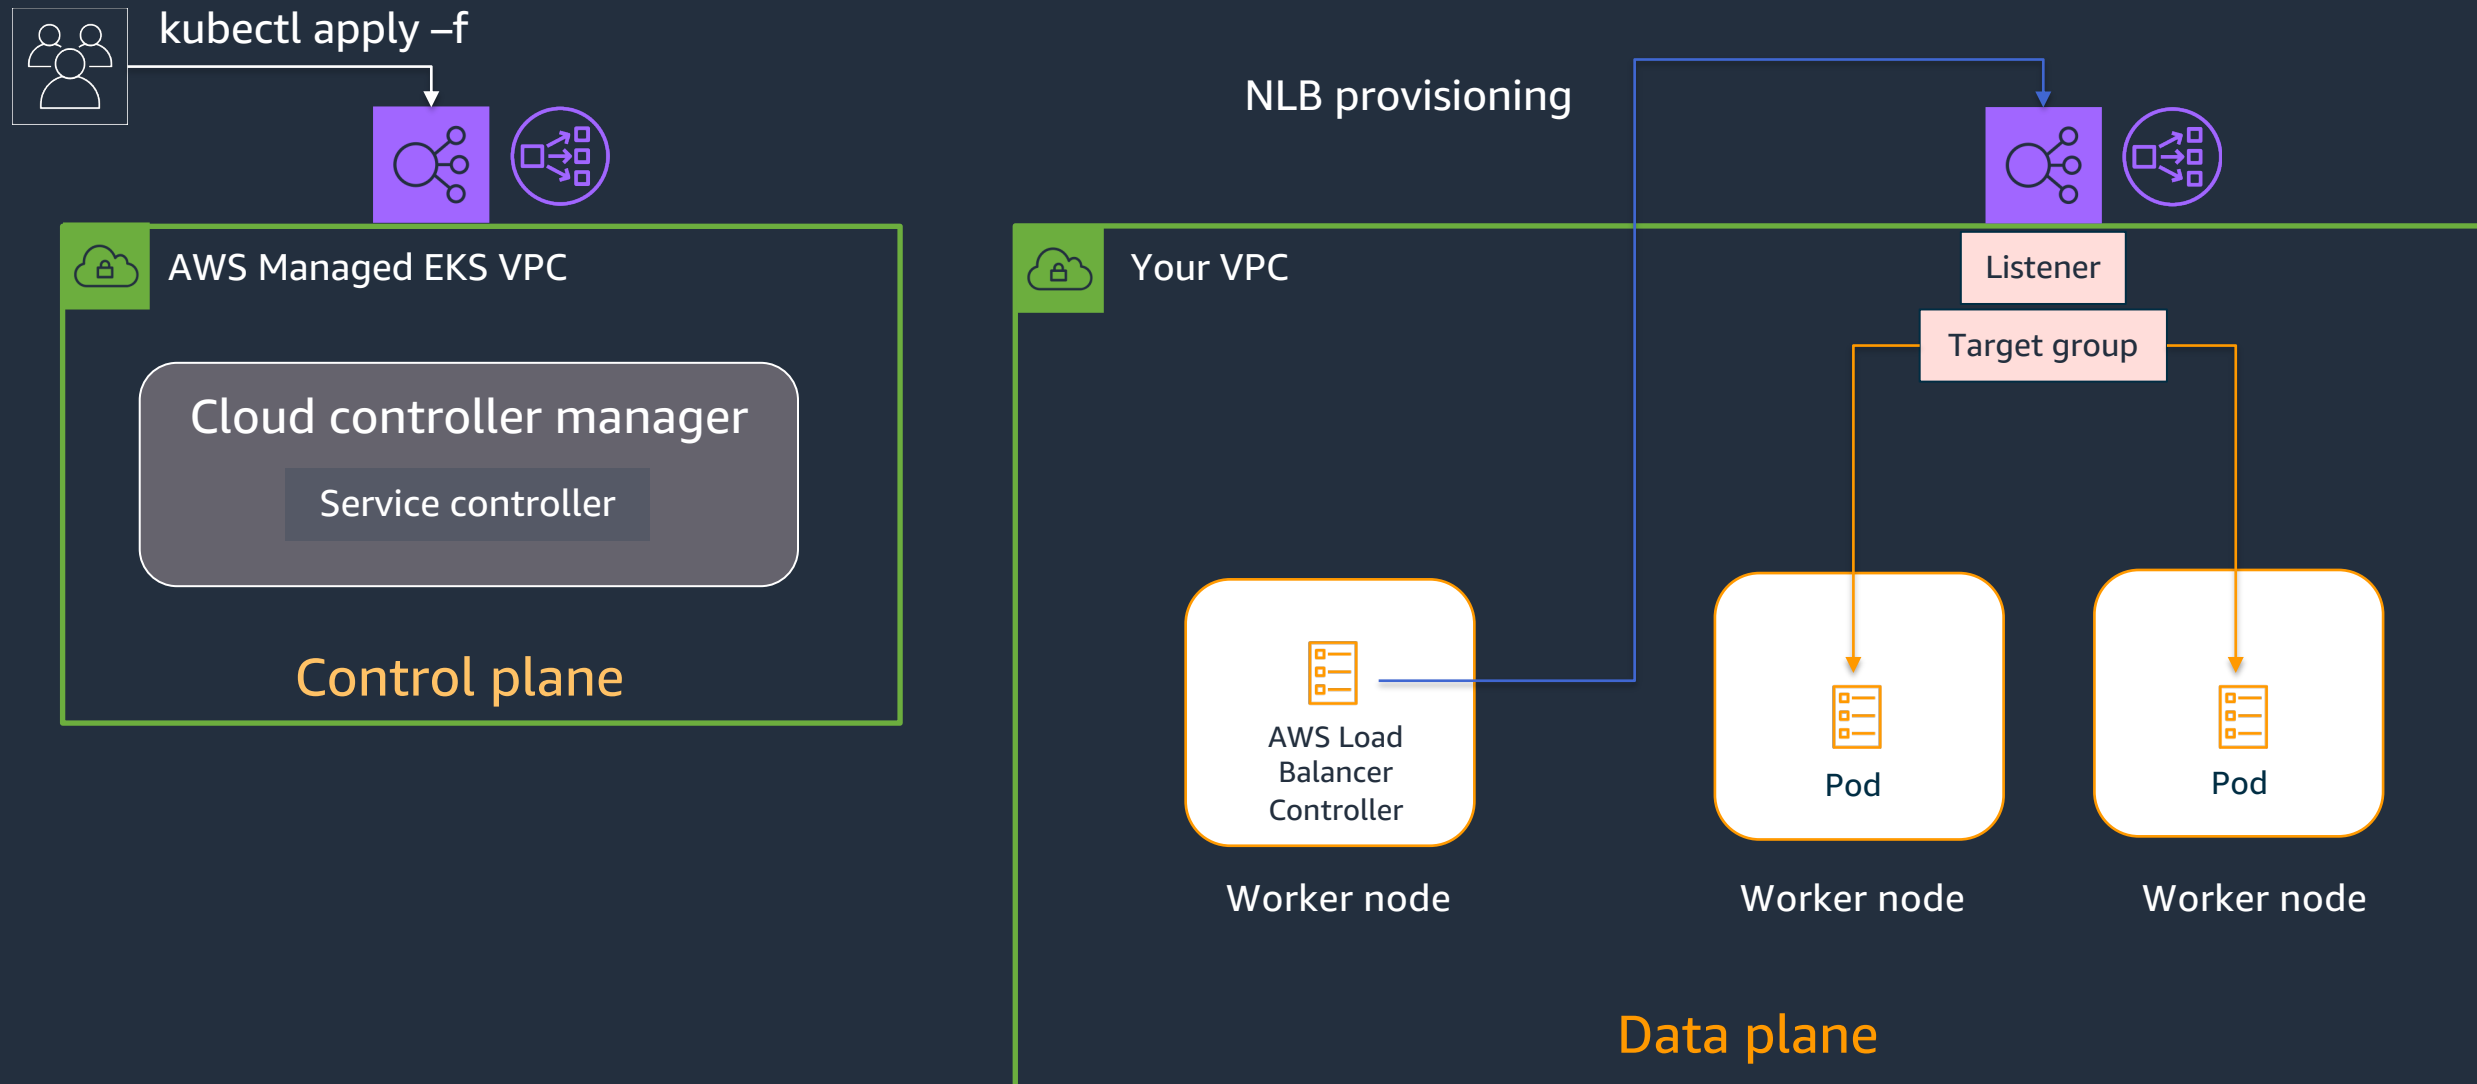
Using AWS Load Balancer Controller
requires version 2.2.0 or later

Network Load Balancer provisioning in EKS



[service.beta.kubernetes.io/aws-load-balancer-nlb-target-type: instance](https://service.beta.kubernetes.io/aws-load-balancer-nlb-target-type:instance)

Network Load Balancer provisioning in EKS



[service.beta.kubernetes.io/aws-load-balancer-nlb-target-type: ip](https://service.beta.kubernetes.io/aws-load-balancer-nlb-target-type:ip)

Application Load Balancer provisioning in EKS

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: echoserver-ingress
  namespace: echoserver
  annotations:
    alb.ingress.kubernetes.io/security-groups: sg-010fc3455c73f0a58, sg-049e999c68a291976
    alb.ingress.kubernetes.io/target-type: ip
    alb.ingress.kubernetes.io/scheme: internet-facing

    kubernetes.io/ingress.class: alb
spec:

  rules:
    - host: echoserver.example.com
      http:
        paths:
          - path: /*
            backend:
              serviceName: echoserver-svc
              servicePort: 80
```

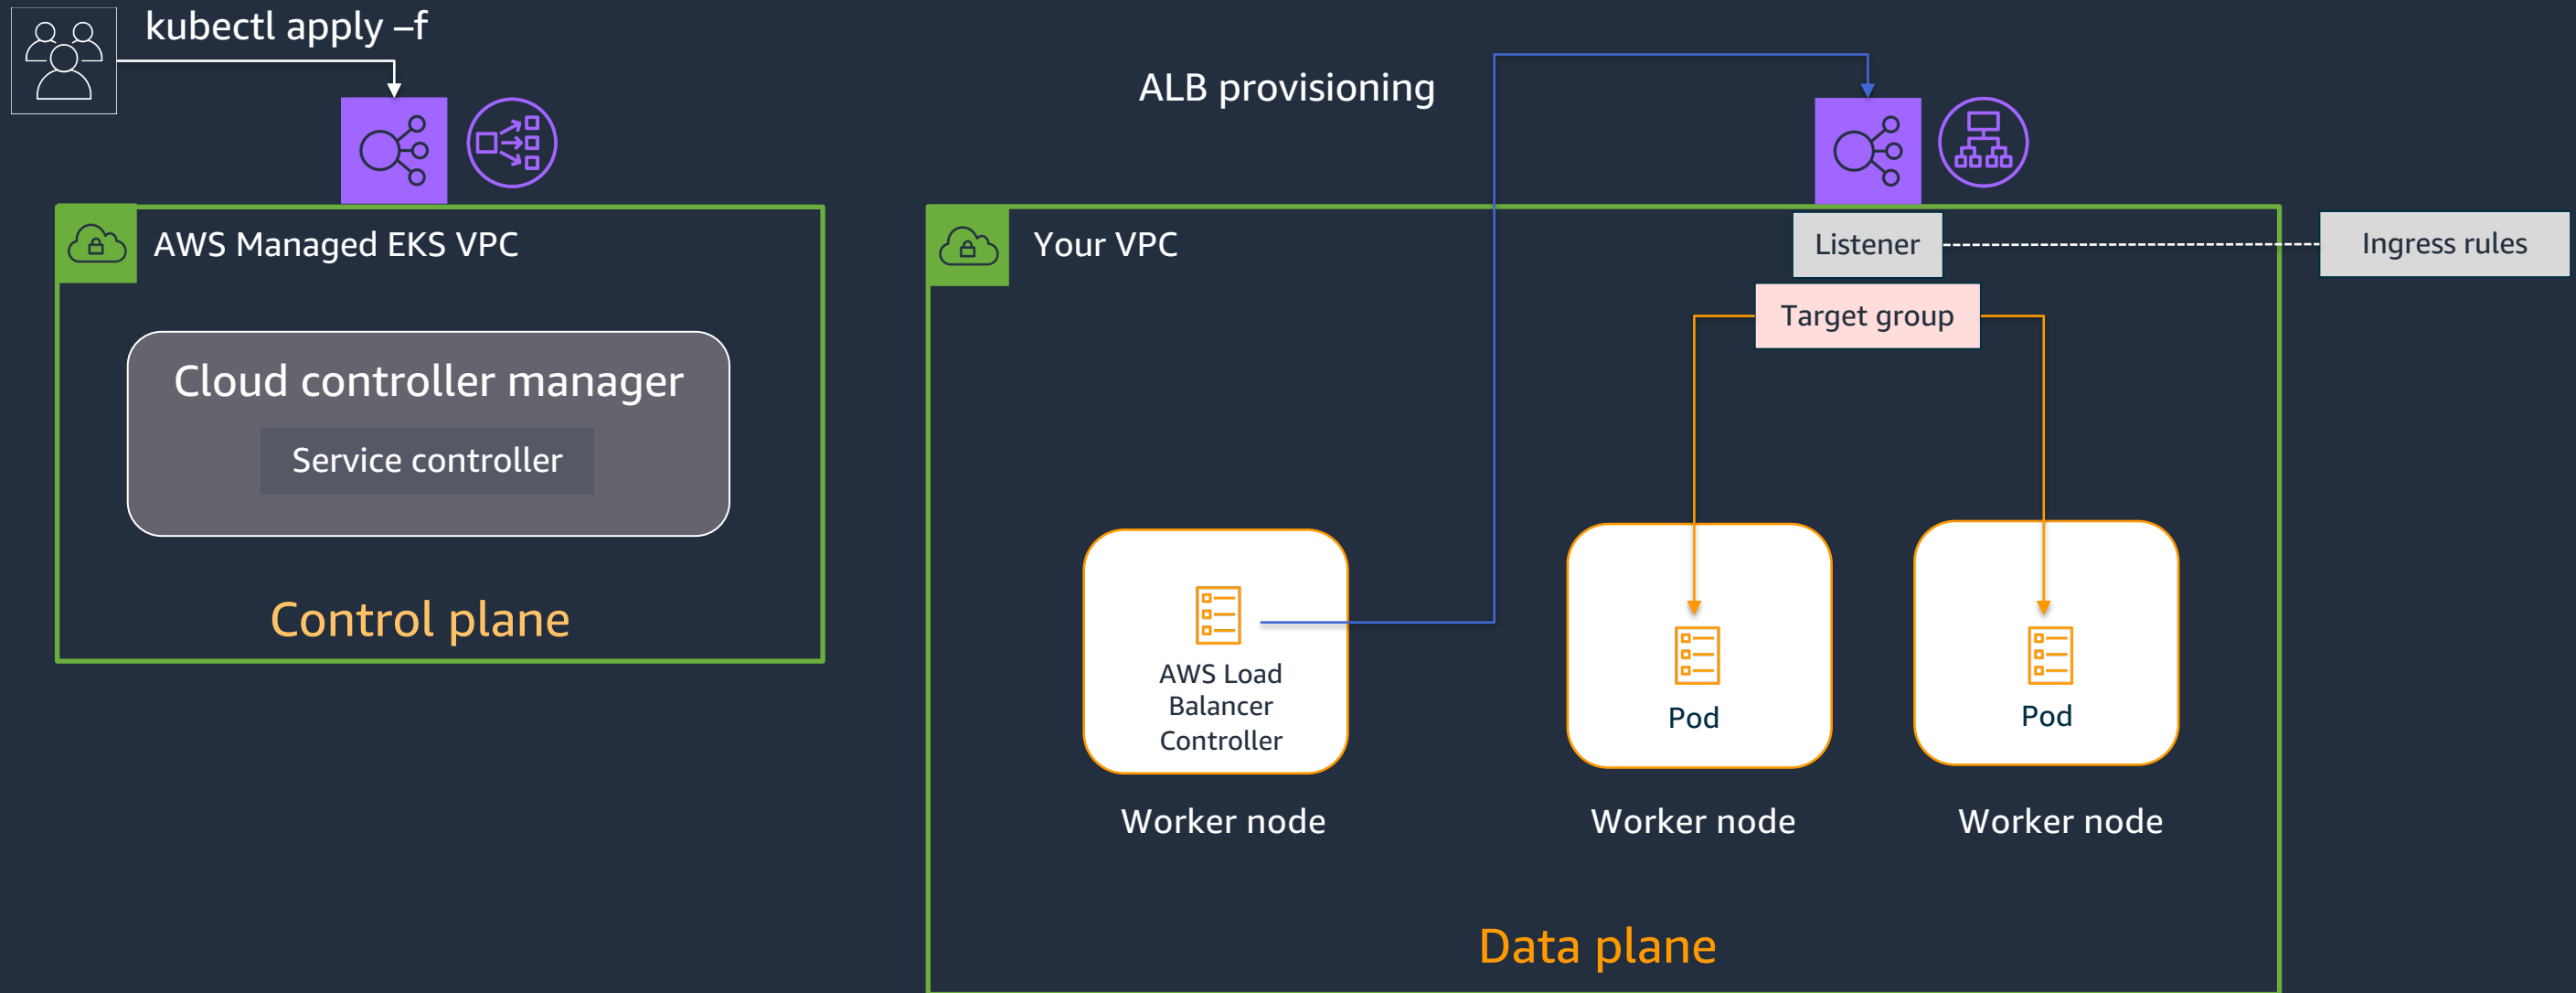
Application Load Balancer provisioning in EKS

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: echoserver-ingress
  namespace: echoserver
  annotations:
    alb.ingress.kubernetes.io/security-groups: sg-010fc3455c73f0a58, sg-049e999c68a291976
    alb.ingress.kubernetes.io/target-type: ip
    alb.ingress.kubernetes.io/scheme: internet-facing

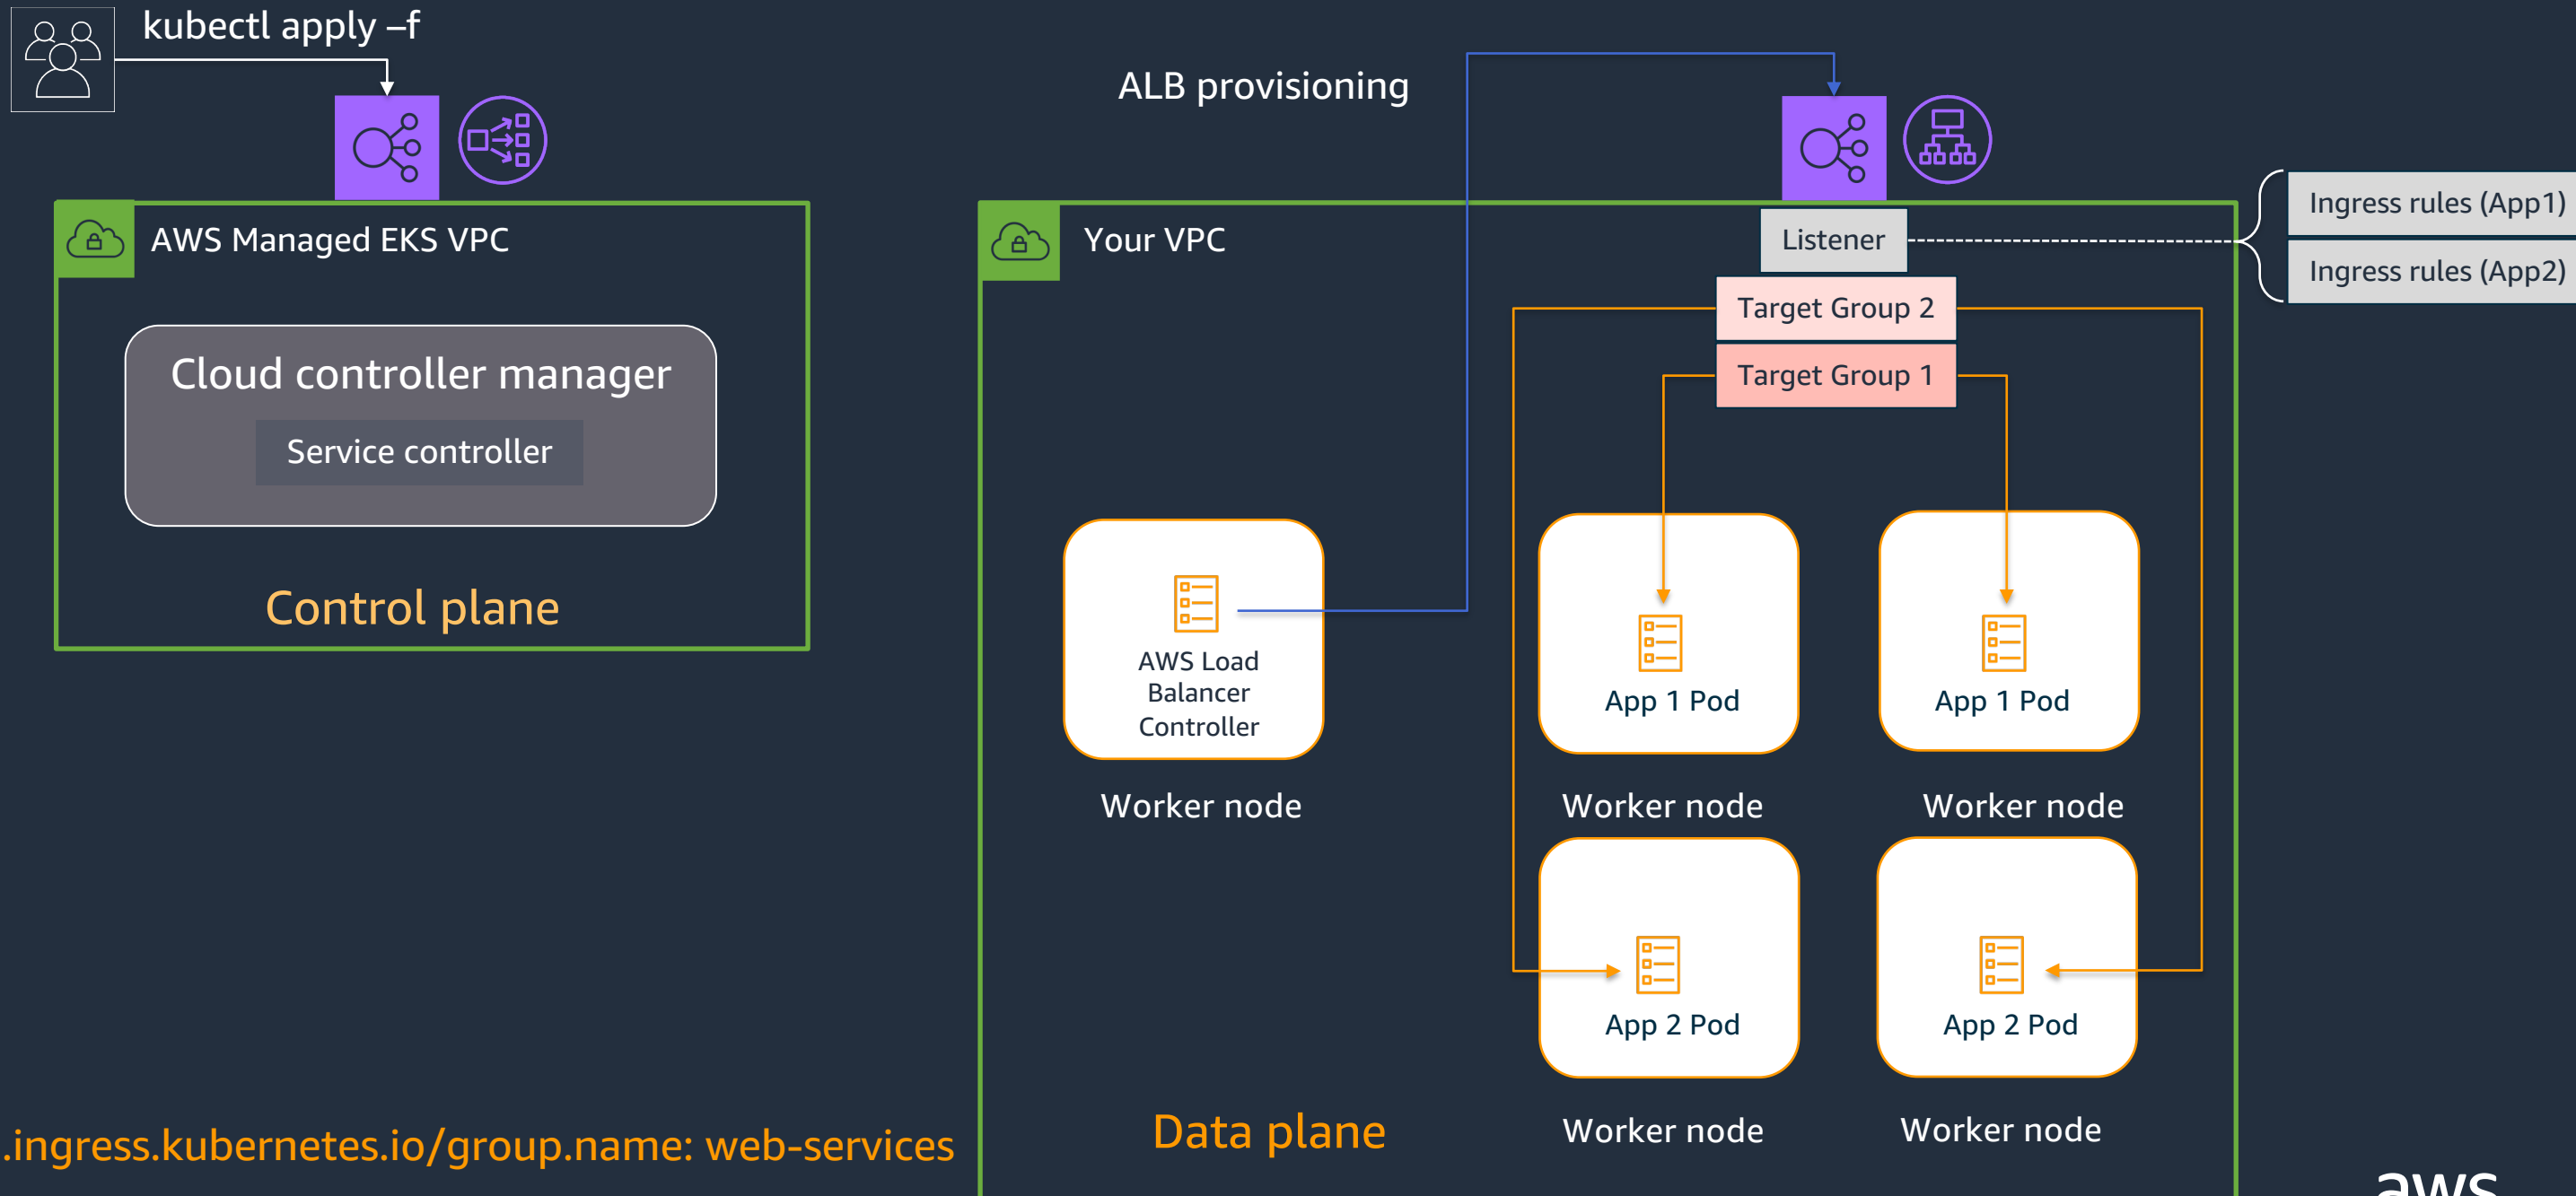
spec:
  ingressClassName: alb-ingress-class
  rules:
    - host: echoserver.example.com
      http:
        paths:
          - path: /*
            backend:
              serviceName: echoserver-svc
              servicePort: 80
```

```
apiVersion: networking.k8s.io/v1beta1
kind: IngressClass
metadata:
  name: alb-ingress-class
spec:
  controller: ingress.k8s.aws/alb
```

Application Load Balancer provisioning in EKS

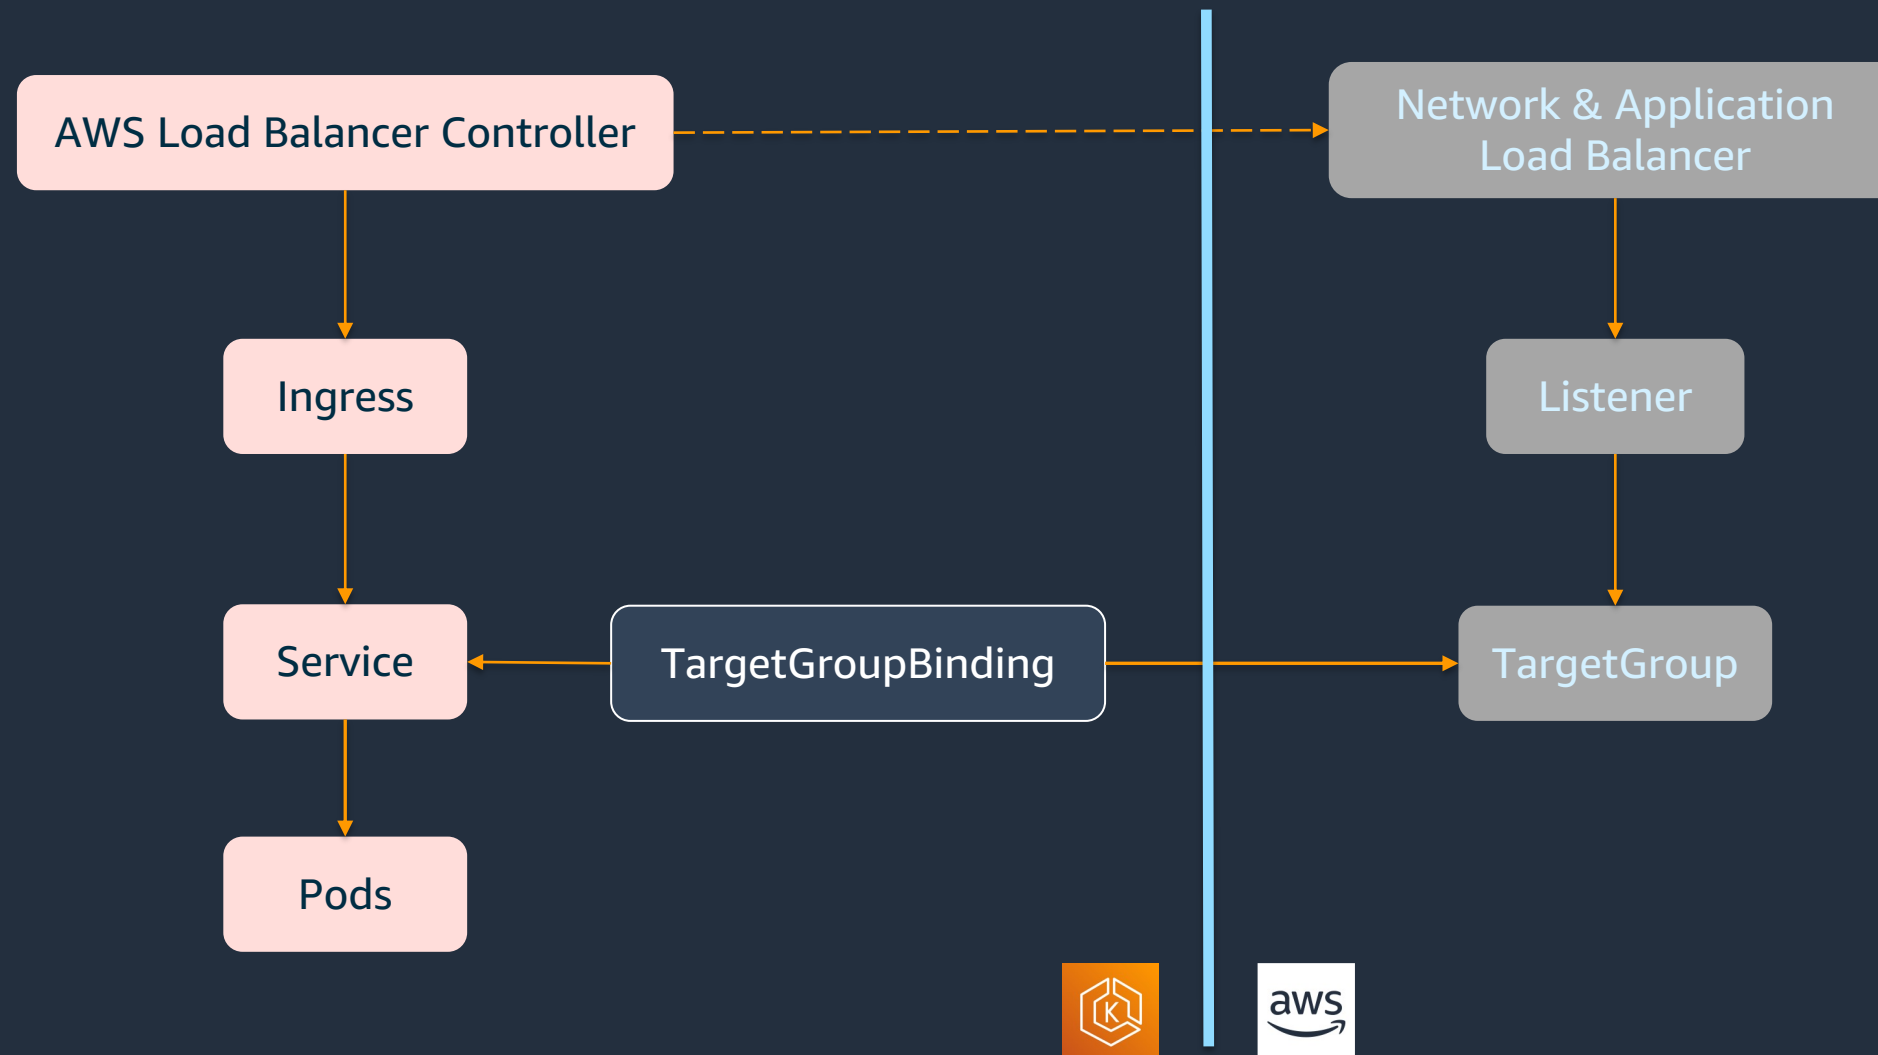


Application Load Balancer provisioning in EKS



alb.ingress.kubernetes.io/group.name: web-services

Binding AWS target groups and Kubernetes services



Workshop Links

- <https://www.eksworkshop.com/docs/networking/prefix/>
- <https://www.eksworkshop.com/docs/networking/custom-networking/>
- <https://www.eksworkshop.com/docs/networking/security-groups-for-pods/>
- <https://www.eksworkshop.com/docs/fundamentals/exposing/>

Questions